

Learning with AI

A Framework for Students, Instructors, and Universities

*From Graduate Mathematics to Undergraduate Education,
Departmental Practice, and Institutional Policy*

James M. Hyman

Tulane University
New Orleans, Louisiana

May 25, 2026

Preliminary version — for review only

Abstract

This book presents a framework for AI-assisted learning at colleges and universities. It is written for three audiences: students learning to use AI without surrendering their own thinking, instructors designing courses that remain intellectually honest under AI access, and academic departments and institutions developing coherent policy and practice. The framework was developed for graduate mathematics and generalized to undergraduate education and to departmental and institutional governance. The examples are especially detailed for STEM and research-intensive courses, but the principles also apply to writing, humanities, social sciences, professional programs, and creative work.

The framework rests on three commitments. First, courses must distinguish what students can do independently from what they can do with AI assistance, and assess each in proportion to the goals of the course. Second, AI use should be visible, verified, and disclosed rather than hidden or assumed. Third, students, instructors, and institutions need shared vocabulary if AI policy is to function in practice.

The book is organized in three Parts. Part I, the Student Guide, gives undergraduate and graduate students a practical compact for using AI as tutor, critic, and study partner without giving up the thinking each course is meant to build. Part II, the Instructor Guide, gives faculty assignment-design tools, assessment models for small and large classes, and rubrics for AI-aware grading. Part III, the Departmental and Institutional Guide, gives departments and universities shared policy language, privacy safeguards, equity considerations, and adoption workflows that work at the level where most decisions are actually made. Three Parts share a single vocabulary: Core Competence, AI-assisted Practice, and Trustworthiness; the Learning Spiral; and the NAI, AIT, AIC, and AIS assignment categories.

Disclosure of AI assistance

This book was prepared with the assistance of generative AI tools, used in the manner that the book recommends for student and faculty academic work. AI tools supported drafting, revision, structural critique, and verification, with the author retaining authorship and responsibility throughout. All arguments, recommendations, citations, and editorial choices are the author's. The work was developed under the same AIC-category practices described in this book: AI may assist scholarly work, but the human author verifies the content and takes responsibility for it.

How to read this book

This book is one bound volume in three Parts. Each Part is written for a different reader, and each Part can be read on its own.

Students should begin with Part I, returning to specific sections as their courses introduce AI policies. Faculty should read Part II, drawing on Part I to anticipate what their students

will read and on Part III to anticipate what their institution will require. Department chairs, teaching-center staff, and administrators should read Part III, drawing on Parts I and II to understand the student-facing and instructor-facing practices the institutional framework supports.

All three Parts use the same vocabulary. The Foreword introduces this vocabulary briefly. The full development of each term appears in the Part where the reader most needs it: Core Competence, AI-assisted Practice, and Trustworthiness in Part I; assignment categories and assessment models in Part II; institutional policy and adoption practice in Part III.

A reader looking for a single chapter to scan should consider the CAT framework section in Part I, the four assignment categories in Part II, or the CAT section in Part III, depending on role.

Readers who want a brief account of the learning-science principles this book operationalizes, including the calibration problem that AI use sharpens, should also see Section 5 (How learning happens, and how AI changes it) early in Part I.

About the companion repository

This book has a companion repository at github.com/machyman/hyman2026learning, with a rendered site at machyman.github.io/hyman2026learning. The repository hosts free, open-source materials that work alongside the book. Two of them are primary tools that students paste into AI applications: the Study Partner Protocol (SPP), a prompt that turns the assistant into a study partner that helps the student learn rather than producing finished answers, and Define Personal Preferences (DPP), a walkthrough that helps a student configure their AI tool's persistent settings the first time they use it. The repository also contains material for instructors (sample syllabus paragraphs, AI disclosure policies, an adoption quickstart) and standalone artifacts for students (a guide to sharing AI conversations with a teacher, disclosure templates, and the pre-submission checklist that appears in Part I).

The book and the repository are intended to be used together. The book teaches the framework; the repository hosts the operational artifacts. References from the book to specific companion materials appear throughout, in the form `companion/spp/`, `students/checklist`, and similar paths. The book remains under SIAM copyright; the materials in the repository are released under the MIT license and can be freely used, modified, adapted, translated, and redistributed with attribution. The repository will continue to evolve after publication; corrections to the book itself, if any are needed, appear in the repository's `ERRATA.md` file.

Acknowledgments

This book benefited from many conversations with friends and colleagues whose critical feedback and discussion on AI-assisted learning shaped the framework presented here. The

author also thanks the broader SIAM mathematical-sciences community for the institutional context that informs Part III’s policy recommendations, and the graduate students who tested early versions of the framework in their own work and whose questions sharpened the assessment practices in Parts I and II.

Particular thanks to those who provided detailed manuscript readings: Joshua Agbomola, Yuwei Bao, Scott McKinley.

The book was developed using the same disclosure-and-verification practices it recommends. Generative AI assistance supported drafting, structural critique, citation verification, and revision across multiple working sessions. The author drafted the original arguments and the framework’s conceptual architecture; AI tools assisted with prose refinement, organizational checks, bibliography verification, and the production of figures and dialogues. Every claim, every citation, every example, and every recommendation was reviewed and validated by the author. Errors that remain are the author’s. The author also thanks Anthropic for the development of the Claude models used during preparation; the relationship is that of a tool user, not a sponsorship.

A note on the disclosure that follows. This book argues that AI-assisted scholarly work should be transparent about the role AI played. The acknowledgment above is intended as one model of what such transparency looks like in practice: enough detail for a reader to understand what AI did and what the author did, without the lengthy methodological exposition that would belong in a separate paper. Readers who want a fuller account of AI-assisted writing standards should consult Appendix C.

Contents

Disclosure of AI assistance	1
How to read this book	1
About the companion repository	2
Acknowledgments	2
Foreword: From Graduate Mathematics to a University Framework	15
I Learning with AI — A Student Guide	22
1 Why this guide exists	22
2 Setting up your AI for learning	24
2.1 Why this step matters: a before-and-after	25
2.2 The prompt is only the visible part	26
3 Learning outcomes	28
4 The Student’s Trilemma	28
4.1 Three student failure modes	29
4.2 How the trilemma manifests differently for different students	30
4.3 What makes the trilemma hard to see	31
4.4 Why the Student’s Trilemma matters for the rest of Part I	31
5 How learning happens, and how AI changes it	32
5.1 The illusion of competence	33
5.2 Calibration checks in the moment	34
6 The basic compact	35

7	The CAT framework	37
7.1	Core Competence	37
7.2	AI-assisted Practice	37
7.3	Trustworthiness	38
7.4	Relation to international AI competency frameworks	38
8	The Learning Spiral	39
8.1	Try first	41
8.2	Ask for help, not replacement	41
8.3	Work independently	42
8.4	Verify	43
8.5	Reflect and disclose	43
9	Four assignment categories	44
9.1	NAI: No AI	45
9.2	AIT: AI as Tutor	45
9.3	AIC: AI as Collaborator	45
9.4	AIS: AI as Object of Study	46
10	Good uses of AI	46
10.1	Use AI to clarify confusing ideas	46
10.2	Use AI to practice retrieval	46
10.3	Use AI to test your explanations	47
10.4	Use AI to find weak points	47
10.5	Use AI to compare approaches	48
10.6	Use AI to prepare for direct questioning	48
11	Annotated dialogues: strong AI use	49
12	Dangerous uses of AI	53
12.1	Do not outsource first contact with the problem	53

12.2 Do not submit work you cannot explain	54
12.3 Do not trust generated citations	54
12.4 Do not let AI change your meaning	54
12.5 Do not trust code just because it runs	54
12.6 Do not use AI to evade the purpose of an assignment	55
13 How to verify AI-assisted work	55
13.1 Verification in STEM courses	56
13.2 Verification in writing courses	57
13.3 Verification in humanities courses	57
13.4 Verification in social science courses	57
13.5 Verification in professional and applied courses	58
13.6 Verification in graduate research	58
14 Disclosure	59
14.1 A short disclosure template	59
14.2 A fuller disclosure template	59
14.3 What good disclosure does	59
14.4 Thread sharing	60
15 Privacy and data protection	60
16 Small classes and the Oral-Defense Model	61
17 Large classes and the Evidence-of-Thinking Model	62
18 Asking better questions in mathematics	62
19 A practical study routine	64
20 Student checklist	65
21 Discipline examples	65

21.1 STEM example	65
21.2 Programming example	66
21.3 Writing example	66
21.4 History example	67
21.5 Social science example	67
22 When not to use AI	67
22.1 The choice not to use AI at all	68
23 What to do when you are unsure	68
24 Evidence base and limits of the guide	69
25 A final word	70
26 Reference materials	70
26.1 Template: AI account configuration	70
26.2 Template: AI use log	73
26.3 Template: verification checklist	73
26.4 Template: prompt library for learning	73
26.5 Template: one-page student compact	75
26.6 Short glossary	75
Publication note	77
II Teaching and Assessing with AI — An Instructor Guide	77
27 Purpose and scope	78
28 Instructor outcomes	79
28.1 Relation to international AI competency frameworks for teachers	79
29 The instructor-side calibration problem	80

30 The Instructor’s Trilemma	82
30.1 Three instructor failure modes	83
30.2 How the trilemma manifests differently by faculty type	84
30.3 Why the Instructor’s Trilemma matters for the rest of this Part	86
31 Three design axes: class size, level, and discipline	86
32 The CAT framework	88
32.1 Core Competence	88
32.2 AI-assisted Practice	89
32.3 Trustworthiness	89
33 Four assignment categories	89
33.1 NAI : No AI	90
33.2 AIT : AI as Tutor	90
33.3 AIC : AI as Collaborator	91
33.4 AIS : AI as Object of Study	91
33.5 A syllabus table for assignment categories	91
33.6 Category mix shifts with course level	92
34 Designing assignments that preserve thinking	92
34.1 Plan-Use-Verify in practice	94
34.2 Design features that improve assessment validity	95
34.3 Design features that weaken assessment validity	96
35 The Oral-Defense Model	97
35.1 Recommended structure	97
35.2 Examples across fields	98
35.3 Equity in oral assessment	98
35.4 Suggested weighting for small classes	99

36 The Evidence-of-Thinking Model	99
36.1 Scalable evidence sources	99
36.2 Suggested weighting for large classes	100
36.3 Random oral audits	100
36.4 TA-supported mini-conferences	101
37 The Hybrid Model	101
37.1 Recommended structure	101
37.2 Sampling strategy for live audits	102
37.3 Worked example: a hybrid course across the semester	103
37.4 Suggested weighting for hybrid classes	104
37.5 Two empirical anchors	104
38 Feedback and grading with AI	105
38.1 Safe uses of AI for feedback	105
38.2 High-risk uses of AI for grading	106
39 Annotated dialogues: weak AI use as teaching cases	106
40 Academic integrity without detector dependence	108
40.1 A fair response sequence	110
40.2 A constructive role for detectors	110
41 Privacy, access, and equity	110
41.1 Privacy	111
41.2 Access	111
41.3 Equity in assessment design	111
41.4 Multilingual students and accessibility	112
42 Discipline overlays	112
42.1 STEM courses	113

42.2 Writing and communication courses	113
42.3 Humanities courses	114
42.4 Social science courses	114
42.5 Professional and applied courses	114
42.6 Creative and design courses	115
43 Rubrics	116
43.1 A simple four-level rubric	117
44 TA calibration and consistency	117
45 Implementation workflow for a course	117
46 Evidence base and limits	118
47 Reference materials	118
47.1 Sample syllabus language	119
47.2 Assignment label template	120
47.3 Student disclosure block	120
47.4 Oral defense rubric	121
47.5 Random oral audit script	121
47.6 AI critique assignment template	121
47.7 Course design checklist	122
47.8 Instructor self-check	122
47.9 Short glossary	122
Publication note	123
 III A Framework for AI-Assisted Learning — A Departmental and Institutional Guide	 123
48 Purpose and scope	124

49 Institutional outcomes	125
50 The Institution’s Trilemma	126
50.1 Three institutional failure modes	127
50.2 How the trilemma manifests differently by institutional type	128
50.3 Why the Institution’s Trilemma matters for the rest of this Part	129
51 The institutional problem	129
52 Guiding principles	131
53 The CAT framework	131
53.1 Alignment with UNESCO learning outcomes	132
54 Shared vocabulary	132
55 Course-level AI-use categories	133
55.1 NAI : No AI	133
55.2 AIT : AI as Tutor	134
55.3 AIC : AI as Collaborator	134
55.4 AIS : AI as Object of Study	134
56 Assessment models by course size	134
56.1 Oral-Defense Model	134
56.2 Evidence-of-Thinking Model	135
56.3 Hybrid Model	135
57 Minimum independent-competence requirement	135
58 Disclosure standard	136
59 Privacy, data governance, and tool approval	137
60 Equity, access, and accessibility	137

61 Academic integrity procedure	138
62 Faculty use of AI	139
63 Required syllabus elements	140
64 Departmental responsibilities	140
65 Teaching-center responsibilities	141
66 Training plan	141
66.1 Student training	141
66.2 Faculty training	142
66.3 TA training	142
66.4 Administrative training	142
67 Implementation plan	142
67.1 Phase 1: Pilot	143
67.2 Phase 2: Department adoption	143
67.3 Phase 3: University adoption	143
67.4 Phase 4: Review and revision	143
67.5 Cost and resource considerations	143
68 Faculty change management	144
68.1 The change-management problem	144
68.2 Stages of faculty adoption	145
68.3 Barriers and how to address them	146
68.4 When the standard playbook is not enough	146
68.5 What works	147
68.6 What does not work	148
68.7 Department leadership	148
68.8 Faculty support resources	149

69 From department to institution: how departmental practice can shape institutional policy	150
69.1 The pattern in practice	150
69.2 What departmental faculty can do to support the bridge	150
69.3 What departments should ask of their institution	151
69.4 The honest limit	151
70 Looking forward	151
70.1 What is likely to change	152
70.2 What is likely to stay the same	152
70.3 What this means for the framework	153
70.4 Modesty about prediction	153
71 Evaluation metrics	154
72 Risk register	154
73 Model institutional policy statement	155
74 Evidence base and limits	156
75 Reference materials	156
75.1 Course classification worksheet	157
75.2 Syllabus template	157
75.3 Disclosure block	158
75.4 Privacy checklist	158
75.5 Academic-integrity workflow	158
75.6 Department adoption checklist	159
75.7 Program evaluation rubric	159
75.8 One-page adoption summary	160
Acknowledgment of source framework	160

Publication note	160
A Appendix A: Graduate Mathematical Intelligence Profile	161
B Appendix B: AI-Integrated Graduate Mathematics Assessment Framework	168
C Appendix C: Best Practices for Ethical AI-Assisted Scientific Writing	174
D Appendix D: Math-Specific Quick Reference	180

Foreword: From Graduate Mathematics to a University Framework

*AI changed how proofs are written.
It did not change what it means to understand one.*

This book began as a problem in graduate mathematics teaching.

In late 2022 and through 2023, students arrived at office hours with proofs they had not written and could not defend. Some were honest about asking a chatbot for help. Others were not. The proofs were uneven: occasionally correct, often plausible but flawed in ways that mattered, and sometimes just wrong in subtle ways the student could not detect. The students themselves were often unsure how much of the work was theirs. They wanted to learn. They also wanted the work done.

The standard responses to this situation were both unsatisfying. A blanket prohibition on AI tools was unenforceable in any course where work happened outside the classroom, and it left students unprepared for the professional and research environments they would enter. A blanket permission was no better: it conceded the field. A polished proof submitted by a student who could not reconstruct it on a blackboard was not evidence of mathematical understanding. It was evidence that the student had used a tool.

The framework presented in this book grew out of an attempt to find a third path for graduate mathematics, one that recognized AI as a permanent feature of mathematical practice while preserving the habits of mind that doctoral training is meant to build. That framework is the *Graduate Mathematical Intelligence Profile*, or GMIP. Its full technical statement appears as Appendix A of this book. The companion document, the *AI-Integrated Graduate Mathematics Assessment Framework*, appears as Appendix B.

Why this is not the first time

The framework presented here grew out of a particular crisis in graduate mathematics, but the underlying question is older. AI is not the first technology to frighten people by redefining skill. The pattern repeats often enough that it is worth naming before the rest of the book builds on it.

When photography became widely available in the nineteenth century, painters had reason to worry. A camera could capture a likeness quickly and cheaply, and the demand for realistic visual copying did decline. But painting did not disappear. Painters moved more deeply into color, abstraction, emotion, and personal vision, into territory a camera could not follow. The camera did not destroy the artist. It changed what the artist was needed for.

When calculators became common in classrooms, educators worried that students would stop learning arithmetic. The concern was not wrong: a student who reaches for a calculator before developing number sense may never fully acquire it. But calculators did not end

mathematics. They shifted the level at which mathematical work happens, from manual computation toward modeling, reasoning, and the judgment of when an answer is plausible. That last skill is what a calculator cannot do, and it is exactly where the mathematician earns the title.

Power tools tell a different version of the same story. A power saw does not make someone a carpenter. It lets a skilled carpenter work faster, and it also lets an unskilled person make mistakes faster. The tool has force but no judgment. As tools grow more powerful, the human responsibility to direct and evaluate them grows with that capability.

The history is not uniformly positive. Routine work disappears, and people whose livelihoods depend on that work are genuinely displaced. The point is not that everything always works out. The point is that across these transitions, the skill of the practitioner did not disappear. It relocated. The work moved up a level, from execution to judgment, from production to evaluation, from doing to deciding what should be done and recognizing when it has been.

The historical argument is summarized here. It is developed at greater length, with citations to research on AI's effects on productivity and learning, in a companion paper (Hyman, 2026a).

This book reads AI in that history. The framework is not built on the assumption that AI will replace teaching or learning. It is built on the assumption that AI shifts what teaching and learning are for, and that universities, like painters and mathematicians and carpenters before them, need to identify what the new level of work requires before they can teach it.

The two-transcript idea

The central move in GMIP is simple. A graduate student's mathematical record is split into two parallel transcripts.

The first is the *Mathematics Transcript*, which records what the student can do without AI assistance, or with tightly bounded help. It includes blackboard proofs, oral examinations, no-AI written exams, and the kinds of probing conversations where understanding either is or is not present. The Mathematics Transcript is graded against the standard the discipline has always used: can the student reason, prove, and explain?

The second is the *AI-Integrated Transcript*, which records what the student can do when they use AI tools responsibly. It includes take-home assignments completed under open-tool conditions with a process log, AI-critique tasks where the student audits a flawed argument, comparative work where multiple tools are triangulated, and research-style tasks where AI helps explore conjectures or organize a literature review. The AI-Integrated Transcript is graded against a different standard: can the student verify, disclose, judge, and produce mathematical work of higher quality with these tools than they could produce without them?

Neither transcript is sufficient on its own. A student who scores well on the AI-Integrated Transcript but cannot reason without tools has not become a mathematician. A student who scores well on the Mathematics Transcript but cannot use modern tools responsibly will be

at a disadvantage in research and professional practice. The framework requires both, and a department adopting it must decide how to weight them (figure 1).




Figure 1: **The two-transcript design.** Each graduate student’s mathematical record is split into two parallel transcripts. The Mathematics Transcript captures what the student can do independently; the AI-Integrated Transcript captures what the student can do with AI tools used responsibly. Neither alone is sufficient.

Generalizing to undergraduates and institutions

The two-transcript idea travels well, but the language does not. “Mathematics Transcript” is the right name only for mathematics. The underlying distinction generalizes: every course has some component of independent competence that students must demonstrate, and some component of AI-assisted practice that prepares them for work that uses tools. (In the CAT framework introduced in Part I, independent competence is named *Core Competence* and AI-assisted practice is named *AI-assisted Practice*; the phrases refer to the same ideas.)

The vocabulary used throughout this book reflects that generalization. The framework’s three axes are renamed for the broader audience.

Original (graduate mathematics)	Generalization (this book)
Mathematical Understanding (M)	Core Competence (C)
AI-Augmented Mathematical Skill (A)	AI-assisted Practice (A)
Epistemic and Ethical Responsibility (E)	Trustworthiness (T)

The renaming is more than cosmetic. “Mathematical Understanding” is too narrow for a writing course; “Core Competence” is the right concept across disciplines. “Epistemic

Responsibility” is correct for a research setting but heavy for an undergraduate first-year syllabus; “Trustworthiness” carries the same meaning in plainer language. The substance of the axes does not change (figure 2). A reader who is most comfortable with the original M/A/E terminology will find it preserved in Appendices A and B; a reader new to the framework will find CAT in the main text.

The mapping is not perfect at the boundaries. Some disciplines emphasize one axis more than others. A foundational calculus course may weight Core Competence heavily; a senior capstone in computer science may weight AI-assisted Practice heavily; a professional ethics course may weight Trustworthiness heavily. The framework gives instructors a vocabulary for those choices, not a recipe.




Figure 2: **CAT as a generalization of M/A/E.** The graduate-mathematics framework on the left uses domain-specific language. The CAT framework on the right uses cross-disciplinary language while preserving the underlying axes. Appendix A retains the original terminology.

Who this book is for

This book is written for three audiences who need to communicate with one another.

Students, both undergraduate and graduate, need a clear standard for what counts as their own work and how to demonstrate it. Part I addresses students directly. The graduate-mathematics roots of the framework appear most explicitly in the verification chapters and in the discipline overlay sections, but the practical compact applies in any course where AI is allowed.

Instructors need ways to design courses that remain assessable when polished AI output is freely available, and ways to grade AI-assisted work that reward thinking rather than fluency. Part II translates the framework into assignment categories, assessment models for small and large classes, and rubrics. Graduate instructors will find the GMIP appendix the most useful technical reference; undergraduate instructors will find Part II sufficient on its own.

Departments and institutions need shared policy language, privacy safeguards, equity considerations, and a fair process for suspected misuse. Part III addresses these directly. The graduate-school origins of the framework matter here too: graduate programs are often where institutional AI policy is first stress-tested, and where the consequences of getting it wrong

show up earliest.

A note on multiple audiences. The framework’s design choices vary along three axes—class size, course level, and discipline—and any specific course’s AI policy is a function of all three. Different parts of this book address different combinations. Graduate students working in research-style settings will find the proof-construction examples, audit-and-repair tasks, and verification practices most directly applicable. Introductory undergraduates in any discipline will find the homework-to-quiz patterns, in-class verification tasks, and category-mix discussions more directly useful. Instructors of small advanced courses and large introductory courses will recognize their settings in different parts of Part II. The framework supports all of these because the three axes are largely independent: an Oral-Defense Model can be used in a small introductory class and in a small graduate seminar, with the assignment-category mix shifting with course level rather than with class size. Part II section 31 makes this design structure explicit for instructors.

What to expect

The framework is presented through three Parts, but the reader will encounter several elements that may be unfamiliar in this kind of book. Part I and Part II include eight annotated student–AI dialogues showing strong, weak, and borderline AI use across calculus, real analysis, percolation theory, programming, history, and writing. The dialogues are not transcripts of real students; they are pedagogical illustrations of behaviors the framework’s distinctions are designed to surface. Part III includes a chapter on faculty change management that addresses the practical question of how a university actually moves from current practice to AI-aware teaching, and a section on what is likely to change as AI capabilities continue to evolve. Both chapters take the position that the framework’s structural commitments are durable while specific implementations require regular refresh.

What this book is not

A few clarifications about scope.

This is not a tool guide. The specific AI products available at the time of writing will not be the products available a year from now. The framework is intentionally tool-agnostic. When the book discusses prompts, examples, or model behaviors, the discussion is illustrative rather than prescriptive. A student or instructor who wants to know how to use a specific tool well should consult that tool’s documentation; the book treats the tool category, not the brand.

This is not a policy mandate. Universities, departments, and individual instructors retain the right to set rules appropriate to their courses, programs, and institutional contexts. The framework is offered as a vocabulary and a set of design choices that have been tested in practice. Local adaptation is expected.

This is not a discipline-balanced treatment. The book’s worked examples and discipline

overlays draw most heavily from mathematics, computational science, and adjacent STEM fields, reflecting the framework's origin in graduate mathematics and the author's primary teaching contexts. Universities and departments in the humanities, qualitative social sciences, and arts will find the framework's structure (CAT, four assignment categories, three assessment models, three design axes, verification-and-disclosure standards) directly applicable, but the specific examples will need adaptation to those disciplines' particular reading, writing, interpretive, and creative practices. The book's discipline overlays in Part II make a partial start on this adaptation; deeper humanities-specific treatments are a natural next step that this book deliberately leaves to disciplinary experts.

This is not an AI-skeptic manifesto. The framework treats AI as a powerful and consequential technology that students will use throughout their careers. Refusing to engage with it would not protect learning; it would simply leave students unprepared. Independent research on AI tutoring suggests that careful pedagogical design can make these tools genuinely useful for learning (Kestin et al., 2025; Mollick and Mollick, 2023), and the early position literature already mapped the opportunity-and-risk landscape that has guided most subsequent work (Kasneci et al., 2023); the framework draws on this body of evidence.

This is not an AI-utopian manifesto either. The same research base shows that unrestricted AI use can produce short-term performance gains while weakening durable understanding (Bastani et al., 2025), that AI systems can produce confident errors and invented sources (Bender et al., 2021), that detector-based enforcement is unreliable and inequitable (Liang et al., 2023), and that institutional AI adoption has so far outpaced policy maturity (American Association of University Professors, 2025b). The framework takes all of these seriously. The Bastani et al. field experiment with nearly a thousand high-school mathematics students is worth quoting precisely: practice-problem performance rose by 48% under an unrestricted ChatGPT-style interface and by 127% under a guardrailed AI-tutor variant, but on the closed-book exam that followed, students who had practiced with the unrestricted interface scored 17% *worse* than students who had never had AI access, while students who used the guardrailed tutor scored about the same as the no-AI group. Equally important, the same study found that students did not perceive any reduction in their learning from copying solutions: subjective self-assessment systematically failed to track the objective deficit. The framework's emphasis on objective Core Competence assessment rests on this metacognitive gap. The Bastani study examined a single secondary-mathematics intervention period at a Turkish high school; replication across other subjects, age ranges, and tool generations is still accumulating, and the specific effect sizes should be read as evidence for the qualitative pattern rather than as parameters that will reproduce exactly in every setting.

The book aspires to a position that has gained ground in recent treatments of AI in education: AI is best understood as a co-intelligence that augments human work without replacing the human who is responsible for that work (Mollick, 2024). Whether that augmentation strengthens or weakens learning depends on the design of the course around it. That is what this book tries to help with.

Acknowledgments and AI assistance

The framework has benefited from conversations with colleagues at Tulane and elsewhere, from drafts shared with graduate students who tested it in their own work, and from the considerable existing literature on academic integrity (McCabe and Treviño, 1993; Bertram Gallant, 2008) and AI-assisted scholarship (Miao and Holmes, 2023; OECD, 2026; U.S. Department of Education, Office of Educational Technology, 2023).

This book was written with the assistance of generative AI tools, used in the manner that the book recommends. I owe the reader a brief note on what that meant in practice.

There were nights when a paragraph I had drafted felt clumsier than the one the AI returned. The polished version made me sound a little more confident than I felt about that paragraph's argument, a little more current than I am with the AI literature, and a little more sure-footed than the underlying thinking yet was. The temptation to let the polished version stand was real, and not only once.

Resisting the temptation was the work. Where I have used AI assistance, I have tried to use it in the way I recommend students and colleagues use it: as a way to clarify thinking that was already mine, not as a substitute for thinking I had not yet done. The book's awkward sentences, the places where the argument moves haltingly, and the moments where I have chosen to sound like myself rather than like a more articulate version of myself are choices.

The author retains full responsibility for the arguments, the recommendations, and the editorial choices.

Part I

Learning with AI — A Student Guide

*A polished answer is not understanding.
AI can produce one and leave the other behind.*

Part I of this book is the Student Guide. It is written for undergraduate and graduate students in any discipline. The examples are especially developed for STEM and research-intensive courses, but the principles also apply to writing, humanities, social sciences, professional programs, and creative work.

Part I is built around a simple standard: if you submit work, you should be able to explain it, verify it, and take responsibility for it. AI may help you reach that point. It cannot replace it.

1 Why this guide exists

Before any of this, before chatbots, I used to catch a copied proof by its notation. A student would hand in a beautifully written argument that used symbols I had never put on the board. They had pulled it out of a textbook and did not understand it well enough to translate the symbols into the ones I had used. The proof was correct. The understanding was not theirs.

I used to catch a copied homework set by its wrong answers. Two students would arrive at the same incorrect numerical result with the same off-by-one error. The conversation that followed was usually not about academic dishonesty. It was about the smaller, more painful question of whether the work had ever been theirs to do.

*Thirty years of teaching: the conversation
after the work turns out to be someone else's
is always the same conversation.*

The technology has changed. The ache has not.

AI tools are now part of ordinary academic life. Students use them to explain hard ideas, summarize readings, generate practice questions, revise drafts, write code, check grammar, and prepare for exams. Used well, these tools can make learning more active and more personal. Used poorly, they can hide confusion, weaken memory, invent sources, and produce work that a student cannot defend.

The question is not whether AI will appear in college work. It already has. The better question is how to use it without surrendering your own thinking.

This guide gives you a practical way to do that. It does not treat AI as forbidden, magical, or harmless. It treats AI as a powerful tool that requires judgment. You should learn how to ask better questions, check what you receive, protect private information, and disclose your use honestly. You should also know when not to use AI at all.

A useful rule runs through the whole guide:

AI may help you learn, but it cannot learn for you.

That rule matters because learning is more than producing a correct-looking answer. Learning means you can explain ideas, make choices, test claims, notice errors, and transfer what you know to a new situation. A polished AI response is not the same as understanding.

The distinction is older than AI. When calculators became common in classrooms, educators worried that students would stop learning arithmetic, and that concern was not wrong: a student who reaches for a calculator before developing number sense may never fully acquire it. But calculators did not end mathematics. They shifted the level at which mathematical work happens. What we value in mathematicians is the ability to know when an answer makes sense even before they verify it, and a calculator cannot do that part of the work. That gap, between an answer and the judgment that an answer is correct, is exactly where your work with AI lives. AI tools will sometimes produce confident errors or miss context that you can detect only if you have done the underlying work yourself.

This is also why your own sense of how much you have learned can be misleading. In a pre-registered trial with nearly a thousand high-school mathematics students, those who practiced with unrestricted AI did not notice that they had learned less than students who worked without AI—but the closed-book exam afterward made the gap visible (Bastani et al., 2025). The work felt successful because the AI produced correct-looking answers. This guide’s emphasis on verification, explanation, and honest self-checks is meant to bridge that gap between what a polished AI artifact looks like and what you actually understand.

This distinction is consistent with current guidance on AI in education. UNESCO’s student framework emphasizes responsible citizenship, critical judgment, and human-centered use of AI (Miao et al., 2024). OECD’s 2026 education outlook describes the need to make generative AI a support for learning rather than a shortcut around learning, and frames the central concern as the gap between the apparent quality of student work (its performance) and its actual learning content (OECD, 2026). Research on study habits also points in the same direction. Practice testing and spaced practice help students retain knowledge because they require active retrieval, not passive exposure (Dunlosky et al., 2013).

Some students arrive at this guide with strong reservations about engaging with AI at all. The concerns are real: data-center electricity and water demand, labor conditions in AI development, concentration of economic power, copyright and creative ownership. These are not deflectable worries, and a guide that ignored them would not deserve the trust it is asking for. The case for engaging with AI critically rather than refusing it absolutely is developed in a companion paper (Hyman, 2026b). That paper approaches the question

through a five-layer literacy framework (technical, environmental, ethical, practical, civic) and treats principled refusal as a defensible position rather than a problem to be talked around. This guide proceeds on the assumption that most readers are deciding how to engage with AI rather than whether to engage at all, but the judgment this guide builds is the same judgment principled refusal also requires.

This guide turns those principles into habits you can use in any course.

2 Setting up your AI for learning

Two things matter before you ask AI anything for coursework. The first is that you have configured your AI to behave the way this guide's scripts assume. The second is that you have made an honest attempt at the problem yourself. Both are easy to skip, both shape every AI interaction that follows, and both are subjects of this Part. This section addresses the first; the Learning Spiral developed in section 8 addresses the second.

Configuration matters because the AI defaults you encounter out of the box are not designed for learning. They are tuned to be impressive in a demo: verbose, over-helpful, and quick to deliver answers. By default, an AI will explain too much, offer solutions when hints would serve, and skip the step where it asks what you actually understand. When you open a fresh conversation, the AI does not know whether you are a kindergarten student or a Nobel laureate, whether this is your first calculus course or your third graduate seminar; it answers at a generic level pitched at no one in particular. You spend extra effort translating the response to your level, or worse, take an answer at face value that was calibrated to someone else.

The attempt-first precondition is the more familiar of the two and the more easily skipped in practice. A student who reaches for AI before trying produces a particular failure pattern: they recognize the AI's solution but cannot generate the analogous one on their own, and they often cannot tell the difference. The pattern has a name in this guide. We call it the **friction-wall** failure, after the moment of productive difficulty that AI use can short-circuit. Trying first does not mean grinding through to a complete solution before allowing yourself to ask; it means making enough of an attempt that you know where you are stuck and what kind of help would be useful. The Learning Spiral develops this discipline.

When using AI, the prompt matters, but the preparation matters more.

The remedy for the configuration gap is straightforward. Every major AI platform now lets you set persistent preferences (sometimes called custom instructions, system prompts, or memory) that apply to all your conversations. A short, well-chosen configuration makes the difference between an AI that gives you generic explanations and an AI that acts as a tutor pitched roughly to your level and your goals. Experienced academic users have converged on six categories worth specifying:

1. *Who you are and what you are studying* (year, department, current courses, prior course-

work).

2. *How you want explanations pitched* (step-by-step reasoning, ask before assuming prerequisites).
3. *How you want uncertainty handled* (flag low-confidence claims, do not invent citations).
4. *Writing style preferences* (English variant, native vs. non-native, tone).
5. *Scope of use and integrity boundary* (what you will and will not use AI for; assignment-level rules take precedence).
6. *Optional: technical-work accuracy guard* (label confident vs. tentative claims; explain non-obvious code lines).

The full configuration template, with copy-and-adapt language for each category and a brief discussion of why each matters, appears in Section 26.1 (*Template: AI account configuration*) at the end of Part I. Spend ten minutes setting these once. The payoff is every subsequent conversation pitched closer to your actual level and goals, with fewer wasted exchanges, and a written record of the use boundaries you have committed to. Students who would prefer to be walked through this setup interactively can use the Define Personal Preferences (DPP) walkthrough in the companion repository ([companion/dpp/](#)).

2.1 Why this step matters: a before-and-after

The difference between a default AI and a configured AI is worth seeing concretely. A calculus student studying for a quiz might ask:

Help me find the limit of $(\sin x)/x$ as $x \rightarrow 0$.

A default AI, untuned and over-helpful, will typically answer immediately with the full machinery: “The limit is 1. The most elegant proof uses the squeeze theorem: for small positive x , we have $\cos x \leq \sin(x)/x \leq 1 \dots$,” followed by L’Hôpital’s rule, the Taylor series argument, and one or two corollaries the student did not ask for. The student now has the answer and three proofs they did not request, and has done none of the thinking that would produce understanding.

The same student asking the same question of an AI configured according to the template at the end of this Part receives something closer to:

What have you tried so far? If you graph $y = \sin(x)/x$ near $x = 0$, what value does the curve appear to approach? Once you have made an observation, I can help you build the formal argument.

The configured response asked back instead of answering. It pointed toward a representation (graphing) that builds intuition before formalism. It promised help *after* the student had done some thinking, not instead of it.

The difference is not the AI’s intelligence. It is the same model in both cases. The difference is whether the AI has been told who is asking and what kind of help is wanted. Ten minutes of one-time configuration produces this shift across every subsequent conversation in your coursework—thousands of exchanges over a semester, each one pitched closer to your level and your learning goals than the default would have been.

A practical note: when an AI gives you a default-style verbose response despite your configuration, the configuration has either been forgotten across a session boundary, or the preferences were not yet active when the conversation started. Most platforms apply persistent preferences to all new conversations, but the application is not always perfect, and platform updates can reset settings. If you notice the verbose default behavior returning, check that your preferences are still active before assuming you have configured something wrong.

2.2 The prompt is only the visible part

The persistent preferences in the previous subsection set the stage. They establish who you are and how you want the AI to engage with you in general. Each individual conversation still benefits from a brief setup that names what you are trying to do, what you already know, where you are stuck, and what kind of help would be useful. A good prompt is not magic. It works best when the AI has been given the setting, the purpose, the audience, the constraints, and the standard of success before it begins. The quality of the response usually depends less on the cleverness of a single question than on the world built around it.

A weak setup says, “Explain this topic.”

A stronger setup says, “I am writing for first-year students who are new to mathematical modeling. They are nervous about equations but curious about real-world problems. Explain this idea using plain language, one concrete example, and no unnecessary jargon. Emphasize intuition before notation.”

The second version does not merely ask for information. It gives the AI a role, a reader, a purpose, and a standard of success. That changes everything that follows.

This is one of the more important lessons students can learn about AI use. The best users are not just better prompters. They are better designers of context. They say what they already know, where they are confused, what kind of answer would help, and what kind of answer would not. They give the AI enough structure to be useful, but not so much that the AI simply fills in a template.

The principle applies most directly to learning. A student who asks, “Do my homework problem,” may receive an answer and learn very little. A student who says, “Here is my attempt. I think the error is in the second step, but I am not sure why. Please ask me one question at a time rather than giving away the solution,” has built a much better learning environment. The prompt is one piece of the request. The deeper act is setting up the

conditions for good thinking.¹

A good prompt asks a question. A good setup teaches the AI what kind of answer matters.

A short analogy makes the point stick. Asking AI for help without setup is like walking into a tutoring session, dropping a textbook on the table, and saying, “Teach me.” You might get something useful. You will do much better if you say, “I understand the definition, I am lost in the example, and I want help finding the first step without being told the whole solution.” The better you describe the learning situation, the better the AI can help you learn.²




Figure 3: *

The prompt is one sentence. The setup is everything before it.

¹Composing prompts in writing is not the only option. Most AI platforms accept voice input or transcribed dictation, which can help students whose spoken expression is more fluent than their written expression — for example, students whose first language is not English. Speak the question, let the platform transcribe it, then refine the wording in writing before sending. The underlying skill (setting up the situation, naming what kind of help would help) does not change with the input modality.

²Students who want to develop this craft more deliberately can draw on several kinds of resources: platform-provider documentation for setting persistent preferences and constructing well-scoped prompts; university AI-use guidance maintained through teaching-and-learning centers, writing centers, or libraries; accessible books on AI literacy (*Co-Intelligence* (Mollick, 2024) matches the orientation this guide develops); and open educational resources. These four categories are intended to age well even as specific instances within them change; the guide deliberately does not endorse specific commercial products or prompt-engineering courses.

3 Learning outcomes

After using this guide, you should be able to do five things.

1. Distinguish help that supports learning from help that replaces learning.
2. Match your AI use to the assignment category chosen by your instructor.
3. Verify AI-assisted work using methods appropriate to your discipline.
4. Disclose AI use briefly, honestly, and usefully.
5. Explain and defend submitted work as your own intellectual responsibility.

Be a clock builder, not a time teller.

These outcomes are modest by design. The guide does not ask you to become an AI expert. It asks you to become a better learner in a setting where AI is available.

4 The Student's Trilemma

Before the principles, a student named Wendy. In May 2025, Wendy was interviewed for *New York Magazine* as part of a long piece of reportage on AI use in college.³ She had just used a chatbot to produce a five-page paper she would turn in within the hour. She told the reporter she still liked writing, and remembered her high-school English class with something like nostalgia. But she would rather get good grades. With ChatGPT, she said, you don't really have to think that much.

Wendy is not the villain of this book. She is its central reader. She has correctly read the incentive structure she lives in, and she is paying a cost she can already feel. Her predicament has a name. Call it the **Student's Trilemma**: a three-way pull, like a dilemma but with one more horn, among earning the grade, being honest about how the work happened, and demonstrating genuine understanding. Each pair pulls against the third. Almost every student now feels some version of it.

The trilemma is not just one student's predicament. Recent surveys document how widely it is felt. A 2025 survey by the United Kingdom's Higher Education Policy Institute reports that 92% of students now use AI in some form, up from 66% in 2024 (Higher Education Policy Institute (UK), 2025). A 2025 Inside Higher Ed Student Voice survey of more than a thousand students from 166 institutions found that students overwhelmingly want their institutions to address academic integrity proactively rather than punitively (Inside Higher Ed and Generation Lab, 2025). A study of disclosure practices at King's Business School found

³Walsh (2025).




Figure 4: **The Student's Trilemma.** Three goals in genuine tension. Each edge labels the failure mode that arises when the two adjacent goals are satisfied at the cost of the third. The text develops the modes. AI sharpens each edge; it did not create the trilemma.

that roughly three-quarters of students failed to declare AI usage even when instructed to do so, with the cited reasons including fear of negative perceptions and lower marks (Gonsalves, 2024). The trilemma names the structural predicament these numbers describe. AI use is widespread; honest disclosure is widely avoided; the gap between performance and underlying competence is widely felt but rarely acknowledged.

4.1 Three student failure modes

Polished hollow. The student earns the grade and discloses AI use honestly but cannot reproduce the work. The submitted artifact is fluent, organized, and correctly cited; the disclosure statement is candid; the underlying competence has not formed. The student does not necessarily realize this in the moment. They have read the AI's solution, found it convincing, and submitted it; what they have not done is reconstruct the steps themselves under conditions that would reveal the gap. The polished hollow is the failure mode that is invisible from outside and felt from inside. The student knows the work is hollow even when no one else does.

Strategic compliance. The student earns the grade and genuinely understands the material but hedges on how AI helped. They used AI for first drafts, for clarifying confusing examples, for stress-testing their own arguments, for catching errors. The learning happened. But the account they give in disclosure statements is shaped to minimize what AI contributed, partly because they expect the disclosure to be read with suspicion and partly because the institutional signals about acceptable AI use are inconsistent. The disclosure is technically true and substantively misleading. Strategic compliance is the most common failure mode because it is the locally rational response to inconsistent expectations.

Honest risk. The student is fully honest about both their AI use and the limits of their

understanding. They disclose what they used AI for, what they verified independently, where they remain uncertain. Their work is documented and defensible. And they pay a competitive price for the transparency: the gradebook treats their disclosed AI-assisted draft the same as a peer's undisclosed identical draft, or worse if the disclosure prompts detector-driven suspicion. The student who refuses to play the strategic-compliance game pays the cost of refusing.

The three failure modes are not symmetric in their consequences. Polished hollow accumulates over time, hollowing out skill the student will eventually need. Strategic compliance protects the grade but corrodes the student's relationship with honest self-account. Honest risk costs the grade but preserves both competence and integrity. Different students under different pressures will find different modes locally least bad. The framework's value is not in telling students which mode to inhabit; it is in naming the structure so the choice is visible.

4.2 How the trilemma manifests differently for different students

The trilemma is universal but its pull patterns are not. The three corners weigh differently depending on the student's situation. Four dimensions of variation matter most. These are pressures that compound for individual students rather than discrete categories.

Economic stakes around the grade. For a first-generation college student whose family is funding the degree as a generational investment, the grade carries weight that students from continuously-college-going families do not feel in the same way. Educational researchers have documented an AI digital divide along this dimension: first-generation students report lower awareness and use of generative AI tools than continuing-generation students, partly because of access constraints and partly because of weaker informal networks that pass on AI-use norms (WGU Labs, 2024). The same student who would benefit most from AI literacy faces the trilemma's strategic-compliance edge most sharply, because the cost of the honest-risk failure mode is higher when the grade is more load-bearing.

Consequence severity around honest disclosure. For an international student on an F-1 visa, an academic-integrity finding can trigger consequences far beyond the grade. Suspension or expulsion can terminate the student's SEVIS record and revoke their visa status, with departure or removal proceedings to follow. The honest-disclosure corner of the trilemma carries weight that domestic peers do not feel, and the documented false-positive rate of AI-detection tools on non-native English writing (section 14) compounds the pressure toward strategic compliance for students who can least afford an unfounded accusation.

Career stakes around demonstrated understanding. For a pre-professional student preparing for medical school, law school, or competitive engineering programs, the credential's signaling value depends heavily on demonstrated independent capability. Polished-hollow work that produces high grades but weak underlying competence carries a delayed cost when the graduate enters a professional environment that tests them directly. The pre-professional student feels the polished-hollow edge of the trilemma most sharply, because the credential is being underwritten by future capabilities they have not actually built.

Time available for honest risk. For working students, students with caregiving responsibilities, and other non-traditional students, the time required to make AI use sustainable (try first, verify independently, document the process, write transparent disclosures) competes directly with employment hours and family obligations. Strategic compliance is more often a structural condition than a chosen response for students whose lives leave no slack for the framework's full practice.

A single student often faces several of these pressures at once. A first-generation international pre-medical student who works thirty hours a week experiences all four. The framework's claim is not that this student should adopt some single right response, but that they should be able to see the structure they are responding to, so the choice is not invisible.

4.3 What makes the trilemma hard to see

The trilemma has a structural feature that makes it harder to address than most predicaments students face: **it is invisible to the people best positioned to remove it.** The polished hollow is designed to be invisible from outside; that is what makes it polished. The strategic-compliance hedge sounds plausible to instructors who are not asking detailed questions about it. The honest-risk price is paid silently, as a missing grade differential the student rarely raises because raising it would sound like a complaint.

Instructors designing courses and administrators setting policy often do not see the trilemma their students are navigating, partly because the failure modes are structured to be unobservable from the instructor's vantage and partly because instructors and administrators face their own version of the same predicament (developed in Parts II and III). The framework's value to instructors and administrators is partly that it makes the student-side trilemma visible from the outside. The framework's value to students is that it names what they have been feeling without a vocabulary for it.

One direct counter-argument is worth taking seriously. A student interviewed in the same article that introduced Wendy asked the obvious question: *If I am being asked to go from point A to point B, why wouldn't I use a car to get there?*⁴ The honest answer is that the comparison fails because the road is the destination. The point of the homework is the walking, and the value of the walking is what the rest of this guide develops. The same is true in a different register for an essay, a problem set, or a research paper. The task that looks like a destination is usually a journey whose value lies in being travelled.

4.4 Why the Student's Trilemma matters for the rest of Part I

The principles, practices, and tools developed in the rest of Part I are the operational response to the trilemma. The Learning Spiral in section 8 is designed to convert the honest-risk failure mode from a competitive disadvantage into a sustainable practice: the student who

⁴Quoted in Walsh (2025).

tries first, verifies, and works independently is producing the kind of evidence that makes honest disclosure defensible without paying the full price of polished-hollow protection. The CAT framework in section 7 aligns the three corners by separating what genuine competence requires (Core Competence), what responsible AI use looks like (AI-assisted Practice), and what trustworthy self-account involves (Trustworthiness). The translation test in section 6 is the diagnostic that catches polished-hollow work before it becomes a credentialing problem. The disclosure templates in section 14 are designed to make honest disclosure both possible and proportionate, reducing the structural pressure toward strategic compliance.

The same three-cornered framework applies in inverted forms to instructors and to institutions; Parts II and III develop the Instructor's Trilemma and the Institution's Trilemma respectively, with parallel failure modes and structural features. Students who recognize the trilemma in their own work will recognize the same structural pattern when their instructors and administrators describe their own predicaments. The framework is the book's load-bearing conceptual contribution. Wendy's predicament is not unique to students.

5 How learning happens, and how AI changes it

Wendy's predicament has the shape it does because of how durable learning actually happens. The practices in this guide are not pulled from intuition. They are operational forms of principles from the science of learning that have been studied for decades. This short section names the principles, so the rest of Part I reads as application rather than assertion. The principles have a longer evidence base than the AI debates that frame them.

Five findings recur in studies of how durable learning happens.

Practice testing (also called retrieval practice). Trying to recall information, or produce a solution, or reconstruct an argument, builds memory more strongly than re-reading or re-watching it (Roediger and Karpicke, 2006). The act of retrieval is itself a learning event, not just a check on what was already learned. Most assessments measure production. Most studying measures recognition. The gap between them is one place where AI use is risky.

Distributed practice (also called spacing). Material reviewed at intervals, with forgetting and re-retrieval between sessions, is retained more durably than the same material reviewed in one block. The interval is the work, not the wasted time. A short return to a problem one to three days later usually does more for memory than a long block on the day of first encounter.

Interleaving. Mixed practice across problem types builds the discrimination skills that single-method practice does not (Rohrer and Taylor, 2007). Students who practice ten problems on the same method become fluent in that method. Students who practice ten problems mixing several methods become fluent at choosing which method applies. In mathematics, the second skill is often the one that matters on exams.

Generation effect. Information that you produce is remembered better than information

you read. Closing the AI window and rewriting the key step from memory is not a redundant exercise. It is the move that converts recognition into production.

Desirable difficulties. Conditions that slow initial performance, including productive struggle, retrieval failure, and effort to reconstruct, often produce more durable long-term learning (Bjork and Bjork, 2011). Easy is not the same as effective. Fluent reading of a polished AI explanation is the easiest study experience available; it is rarely the most effective.

These five findings show up at specific places in the Learning Spiral introduced later in Part I (section 8). *Try first* uses retrieval practice and desirable difficulties. *Verify* adds another retrieval pass and a chance for spaced and interleaved follow-up. *Work independently* and the recommendation to redo key steps from a blank page rely on the generation effect. *Reflect and disclose* closes the pass with elaboration and metacognition. Each step has a learning purpose, not just a procedural one.

5.1 The illusion of competence

A persistent finding in the science of learning is that students often think they understand material that they cannot actually reproduce, apply, or defend under questioning. The technical name for this gap is **calibration error**: the distance between what you think you know and what you can demonstrate. Calibration error has been studied for decades. AI does not create the problem, but it widens the gap in ways earlier study tools did not.

The widening shows up in four places.

The recognition–production gap. Reading a fluent AI explanation produces a strong feeling of understanding. That feeling is reliable for one thing: it tells you that you can recognize the explanation as correct when you see it again. It does not tell you that you can produce the explanation, or the underlying argument, on your own. Recognition memory and production memory are separate systems. AI delivers recognition. Assessments measure production.

I read an AI proof, nodded along, then couldn't reproduce step two. —student

The smoother the prose, the more carefully you should check the facts.

The mastery illusion on transfer. After AI-assisted homework, students often predict they can solve similar problems independently on a quiz. The prediction is systematically wrong. Fluency on assisted work is a poor predictor of unassisted performance, especially when the new problem requires a small variation of method or a transfer to a new context.

The effort illusion. Time spent with AI feels productive. The hours register as study hours. But productive learning hours are hours in which you did the cognitive work: retrieval,

reasoning, reconstruction, error repair. If AI did that work and you watched, the time is study time only in name.

The authorship illusion. In extended AI-assisted writing or proof work, the line between “I had this idea” and “AI suggested it” can become difficult to track. The sensation that the work is yours grows even when most of the structure came from somewhere else.

These illusions are not failures of effort or honesty. They are predictable consequences of how memory, metacognition, and fluency interact (Bjork and Bjork, 2011; Dunlosky and Metcalfe, 2009). The defense against them is not to avoid AI. It is to build in checks that reveal what you actually can do without it.

Most of the practices in this book are calibration checks. *Try first* gets a baseline before AI flattens the picture. *Verify* asks you to find errors in AI output, which requires that you understand the work. *Blank-page redo* tests whether you can produce, not just recognize. *Reflection and disclosure* track what you did versus what AI did. *No-AI quizzes* and *transfer questions* surface the gap between assisted and unassisted performance.

The instructor-side problem is parallel and is taken up at the start of Part II. Together, these checks, student-side and instructor-side, are what allow learning to remain visible in an environment where final products no longer reliably show how they were made.

5.2 Calibration checks in the moment

The four illusions develop slowly and become visible only when you try to do something new. Most students do not catch them in time, because the moment of catching is also the moment of submitting. The remedy is to build the catch into the work itself. Five short habits, run in the moment, do most of the work.

Predict before reading. Before you read AI’s next sentence, write down what step you expect next, or what conclusion you think it is heading toward. Then read. The gap between your prediction and AI’s continuation is your calibration error, made visible at the granularity of one sentence. If the gap is large and surprising, you have learned something. If the gap is large and you find AI’s reasoning hard to follow, you have learned that you are a step behind and need to slow down. The whole exercise takes ten seconds and converts passive reading into active hypothesis testing.

Forced unassisted gap. After AI helps you with one problem, close AI and try the next problem of the same type without it. The drop in fluency between the assisted problem and the unassisted one is a direct measurement of how much you actually learned. A small drop means the assisted work transferred. A large drop means you produced an answer without learning. Most students do not run this check because the assisted problem feels finished and the unassisted one feels like extra work; that feeling is the effort illusion at work, and is what makes the check valuable.

Blank-page rehearsal. At the end of each AI-assisted session, close AI, close your notes, and rewrite the key derivation, paragraph, or argument from a blank page. This is the strongest

single calibration check available, because reconstruction from nothing tests production memory rather than recognition memory. The gap between your draft and what you can rebuild is the gap between recognizing the work and owning it. The student checklist in section 20 ends with this item for the same reason.

Self-quiz on the next day. Within twenty-four hours of an AI-assisted session, set yourself one question about what you worked on the day before, and answer it without the materials in front of you. This is retrieval practice on a one-day delay; it is also a calibration check, because the question reveals whether yesterday’s understanding survived overnight. If it did not, you know to return to the topic before the next class or assignment.

Explanation aloud. After any AI-assisted exchange, say out loud, in your own words, what you just learned. The act of vocalizing forces production. If you find yourself stalling, hedging, or reaching for AI’s exact phrasing, you have caught a recognition–production gap in the act. The fix is to return to the underlying material, not to re-read AI’s response.

These five moves are not a separate study program. They are micro-checks distributed across the Learning Spiral introduced later in Part I (section 8), run as the work happens rather than at the end. *Predict before reading* sits at the boundary between *Ask* and the next AI turn. *Forced unassisted gap* sits at the transition between *Verify* and the next assignment. *Blank-page rehearsal* runs at the close of *Reflect*. *Self-quiz* runs the day after, on the next pass. *Explanation aloud* runs continuously, in any silence the work allows. Built into the work, calibration becomes a habit rather than a final exam.

6 The basic compact

AI-assisted learning works best when students and teachers share a clear compact.

You may use AI when your instructor allows it. You may ask it to explain ideas, question you, critique a draft, help you study, compare approaches, or point out possible weaknesses. You may use it as a tool for practice and improvement.

But when you submit work, the work is yours. That means you should be able to explain it without saying, “The AI told me.” You should know which sources support your claims. You should understand the code you submit. You should be able to describe how a paragraph changed during revision. You should know what you accepted, what you rejected, and why.

This is not only an academic integrity rule. It is a learning rule. If you cannot explain the work, the work has not yet become part of your understanding.

One concrete diagnostic helps. Call it the **translation test**: take a piece of work and re-present it in a different form. The broadest form of the test asks whether you can say the same thing three ways: visually (a picture, diagram, or graph), mathematically (with notation, equations, or formal structure), and in plain language (the kind of sentence you would use to explain the idea to a curious friend at dinner). A student who can move fluently between all three modes owns the underlying idea; a student who can perform only one likely

has surface recognition, not structural understanding.

A specific case of the test is the notation change: rewrite the proof or argument using different variable names and symbols. Figure 5 shows this form. The notation change is a useful diagnostic in mathematics and in any field that uses formal notation, but it is the narrowest version of the test. A student who can rewrite $f(x)$ as $g(t)$ has shown they recognize the structure under one specific kind of relabeling. They have not necessarily shown they could draw the same idea, or explain what is happening to someone who has never seen the notation. The three-modality form catches both kinds of gap.

For non-mathematical work, the test usually runs on the plain-language axis. Paraphrase the passage in your own words, give the same example in a different form, or re-derive the result by a different route. The diagnostic is the same: fluent translation across modes indicates ownership; difficulty indicates recognition. The guide returns to this diagnostic throughout.

Original notation	Translated notation
$M = \sup_{x \in [0,1]} f(x)$	$B = \sup_{t \in [0,1]} g(t)$
$\exists x_{n_k} \rightarrow x^* \in [0, 1]$ (Bolzano–Weierstrass)	$\exists t_{n_k} \rightarrow t^* \in [0, 1]$ (Bolzano–Weierstrass)
$f(x^*) = M$ by continuity	$g(t^*) = B$ by continuity

If you can write the right column fluently, the proof is yours.

Figure 5: **The translation test in practice.** The right column is the same argument with every name changed: $M \rightarrow B$, $x \rightarrow t$, $f \rightarrow g$, $x^* \rightarrow t^*$. Fluency in writing the right column confirms structural understanding rather than surface familiarity with the original notation.

The translation test catches more than single-chat overreliance. It also catches AI work produced outside any single, identifiable session, including work whose surface was polished in one chat while the actual content was drafted in another. The test asks what the student can do regardless of which chat produced what. The operational version of this concern is treated in section 14.4; the translation test is the structural safeguard that does not depend on which chat is shared.

The compact has five parts.

Attempt. Try the task before asking AI to do anything important.

Use. Ask AI for support that keeps you active.

Verify. Check important claims, sources, computations, and interpretations.

Disclose. State how AI was used when the assignment requires it.

Own. Submit only work you can explain and defend.

Different courses will apply this compact in different ways. A professor may ban AI on an exam, allow it for practice, require it for a critique assignment, or ask you to compare your answer with an AI answer. The exact rule belongs to the course. The responsibility belongs to you.

With AI tools: verify, verify, verify, then trust.

7 The CAT framework

This guide uses the CAT framework: **Core Competence**, **AI-assisted Practice**, and **Trustworthiness**. These three dimensions help you understand what your instructor may be trying to assess. Graduate students whose programs use the original M/A/E framework (Mathematical Understanding, AI-Augmented Mathematical Skill, Epistemic Responsibility) will find the technical version in Appendix A; the substance is the same.

7.1 Core Competence

Core Competence is what you must be able to do without AI, or with only tightly limited help. Every course has a core. In chemistry, it may include explaining mechanisms and interpreting lab results. In writing, it may include building an argument and revising a draft with intention. In history, it may include reading sources carefully. In mathematics, it may include solving problems and explaining definitions. In programming, it may include tracing code and debugging simple errors.

AI does not remove the need for these abilities. It makes them more important. If a tool can generate a fluent answer, your value as a learner comes from knowing whether the answer is right, relevant, honest, and useful.

You should expect some assignments to test Core Competence directly. These may include no-AI quizzes, in-class writing, lab practicals, oral explanations, short exams, source analyses, board work, or live coding checks. These assessments are not a rejection of AI. They protect the part of learning that must remain yours.

7.2 AI-assisted Practice

AI-assisted Practice is your ability to use AI well. This is also a real skill. It includes asking useful questions, judging responses, comparing explanations, finding weak points, generating practice problems, revising carefully, testing code, checking assumptions, and using feedback without becoming dependent on it.

Good AI use is not measured by how much AI contributed. It is measured by whether AI helped you think better.

For example, a weak use of AI is asking for a complete solution and copying it. A stronger use is asking for a hint, trying the next step yourself, and then checking whether your reasoning holds. A weak use of AI is asking it to write a paper. A stronger use is asking it to find the least convincing paragraph in your own draft and explain why a reader might not believe it.

7.3 Trustworthiness

Trustworthiness is your ability to verify, disclose, critique, and take ownership of AI-assisted work. This is the ethical center of the guide.

Trustworthiness asks practical questions.

- Did I check whether the AI response is true?
- Did I verify sources rather than trusting generated references?
- Did I protect private or restricted information?
- Did I disclose AI use when required?
- Can I explain what I submitted?

Your own notes on ethical AI-assisted scientific writing state the same principle in a research setting: AI assistance becomes acceptable only when the human author verifies the content and accepts responsibility for it. The guide applies that principle to student learning. AI may assist, but the final responsibility is yours.

7.4 Relation to international AI competency frameworks

The CAT framework is narrower in scope than the UNESCO AI competency framework for students (Miao et al., 2024), which articulates twelve competency blocks across four aspects (Human-centred mindset, Ethics of AI, AI techniques and applications, AI system design) and three progression levels (Understand, Apply, Create). UNESCO addresses what every citizen should know about AI in society. CAT addresses something narrower: how a student demonstrates that submitted academic work reflects their own understanding when AI tools are available. The two frameworks complement each other; CAT is best understood as a focused subset of UNESCO’s broader competency map applied specifically to the assessment of academic work.

Table 1 sketches the correspondence. Core Competence overlaps most strongly with the discipline-specific knowledge a student demonstrates within UNESCO’s *AI techniques and applications* aspect at the *Apply* level, but extends beyond AI competence into the discipline’s own foundational skills. AI-assisted Practice corresponds most closely to UNESCO’s *Apply*

level across all four aspects. Trustworthiness maps onto UNESCO’s *Human-centred mindset* (specifically Human accountability) and *Ethics of AI* (specifically Safe and responsible use). The mapping is not one-to-one because the frameworks were designed for different purposes, but a student or program that has met UNESCO’s competency expectations will already have done much of what CAT asks for.

Table 1: CAT framework and the UNESCO AI competency framework for students (Miao et al., 2024): a sketch of correspondence. The mapping is partial because the frameworks have different scopes; UNESCO covers AI literacy and AI design more broadly, while CAT focuses on the assessment of academic work.

CAT component	Closest UNESCO competency blocks
Core Competence	Discipline-specific foundations; partial overlap with AI foundations and Application skills (<i>Understand</i> and <i>Apply</i> levels of <i>AI techniques and applications</i>).
AI-assisted Practice	The <i>Apply</i> level across all four UNESCO aspects, especially Application skills under <i>AI techniques and applications</i> .
Trustworthiness	Human accountability (<i>Apply</i> of <i>Human-centred mindset</i>); Safe and responsible use (<i>Apply</i> of <i>Ethics of AI</i>); Embodied ethics (<i>Understand</i> of <i>Ethics of AI</i>).

The framework’s positioning relative to UNESCO matters for adoption. A program that already aligns its curriculum with UNESCO’s framework can adopt CAT as the assessment-level operationalization of UNESCO’s expectations for AI-assisted academic work, without abandoning the broader UNESCO structure.

8 The Learning Spiral

The Learning Spiral is the main habit this guide asks you to practice. It has five steps: try first, ask for help, work independently, verify, and reflect (figure 6). The Spiral assumes you have already configured your AI account so it knows who you are and how you want it to respond. If you have not done that yet, do it now using the starter template in section 2; the Spiral’s calibration depends on that configuration. The five steps recur on every AI-assisted task, and each pass through them should deepen what you understand rather than return you to the same place—this is why the framework is called a spiral rather than a loop.

The same five-step spiral, read as a rubric, becomes what AI-assisted work should show. See section 32. —ed.




Figure 6: **The Learning Spiral.** The five-step practice students run on every AI-assisted assignment. Trying first protects independent thinking; asking for help rather than for a solution preserves learning; working independently after the help confirms understanding; verification catches errors; reflection and disclosure close the pass and prepare the next, deeper, attempt. The figure shows the five-step structure; the framework is a spiral rather than a closed cycle because each pass should build on the last.

Step	What it means
Try first	Make an honest attempt before asking AI for help.
Ask for help	Ask for hints, explanations, questions, feedback, or critique.
Work independently	Return to the task and produce your own answer, draft, code, or analysis.
Verify	Check important claims, sources, computations, and assumptions.
Reflect and disclose	Record what AI contributed, what you changed, and what you can now explain.

The loop runs at multiple time scales. Within a single task, the five steps are a sequence. Across an assignment, a unit, or a semester, the same loop repeats: what you learned from verification in one pass becomes part of the try-first in the next; what you reflected on becomes the setup for the next round of questions. The local cycle on each task supports a longer cycle on the work as a whole.

The five steps name skills, not behaviors that come automatically. A first-generation student new to a discipline may not know what trying first looks like for a proof, what verifying means for a citation in a specific field, or what counts as working independently in a programming class versus a literature class. Each step is a skill to be built, often in partnership with an instructor, advisor, or the study practices the rest of this Part describes. Readers who recognize themselves in this paragraph should know that the framework was designed with their starting point in mind, and that the companion materials in `companion/spp/` explicitly support students who are new to any of this.

8.1 Try first

Before using AI, make a first attempt. Write down what you think the problem asks. Draft a paragraph. Sketch the argument. Try the calculation. Read the primary source. Mark the confusing sentence. Run the code. List the assumptions.

Before you type a prompt, also write down the exact step where your reasoning failed or the specific concept that feels unclear. Naming the point of failure is itself a learning move; it forces you to organize what you already know and identify the boundary of your current understanding.

This first attempt does not need to be good. Its purpose is to give your mind something to work with. Once you have tried, AI can help you compare, clarify, and repair. Without that first attempt, AI is more likely to become a substitute for thinking.

A counterintuitive finding in the learning literature is that effortful generation, including unsuccessful generation, often produces more durable memory than studying correct examples alone. Even when your first attempt is wrong, the act of attempting before consulting AI strengthens what you will learn from AI's response.

A recent neuroimaging study at MIT, currently circulating as a preprint pending peer review, examined what happens in students' brains during essay writing under three conditions: writing with no tools, writing with a search engine, and writing with an AI assistant. The reported finding is striking: within an hour of finishing, AI-assisted writers showed dramatically reduced ability to recall content from their own essays compared with those who wrote without AI assistance. The brain imaging suggested that students who started with the AI shifted their cognitive activity from generating content to supervising AI-generated content, with lower neural connectivity overall. Most relevant to this step: students who wrote alone first and then turned to AI preserved much higher levels of recall and engagement than students who started with AI and continued alone (Kosmyna et al., 2025). The specific numbers should be read as preliminary pending peer review, but the qualitative pattern is consistent with the rule above: the first attempt gives your mind something to work with, and that initial cognitive activation is what later AI assistance amplifies, not replaces.

As a practical threshold: for routine homework, give yourself at least ten to fifteen minutes of serious attempt before asking AI for a hint. For a proof or open-ended problem, spend enough time to write down the definitions, the goal, one possible strategy, and the point where you are stuck. Do not ask AI before you have something for it to respond to.

8.2 Ask for help, not replacement

A useful AI question keeps you active. It asks for support that leaves the central thinking with you.

Instead of asking:

Write my answer to this problem.

try asking:

Give me one hint about how to start, but do not solve the problem.

Instead of asking:

Write my essay on this topic.

try asking:

Here is my thesis and outline. What is the weakest part of the argument?

Instead of asking:

Fix my code.

try asking:

Ask me diagnostic questions that would help me locate the bug.

The best prompts are not commands to replace your work. They are requests for coaching. Students who want a tested prompt for converting the AI into a study partner rather than an answer machine can use the Study Partner Protocol (SPP) in the companion repository ([companion/spp/](#)).

8.3 Work independently

After receiving help, step away from the AI response and do the work yourself. Write the paragraph in your own words. Complete the proof or calculation. Revise the code. Rebuild the argument. Explain the source. Decide what belongs in the final submission.

This step is easy to skip because AI output often looks finished. Do not confuse finished-looking with learned.

A good test is simple: close the AI window and try to explain the idea aloud. If you cannot explain it, return to the material.

8.4 Verify

Verification is the habit that turns AI assistance into responsible academic work. It is also the most important calibration check in the Learning Spiral: finding errors in AI output requires that you understand the work well enough to recognize when something is wrong. AI systems can produce confident errors. They can invent sources, misread prompts, use the wrong method, make hidden assumptions, or change the meaning of your sentence.

Verification looks different in different fields.

In a STEM course, you may check units, test a limiting case, compare with a known result, reproduce a calculation, run code on a simple example, or check whether an answer is reasonable in magnitude.

In a writing course, you may check whether the revision still says what you meant. You may compare the new version to your original thesis. You may make sure the paragraph still sounds like you.

In a humanities course, you may return to the text, verify a quotation, check historical context, or compare an interpretation against the evidence.

In a social science course, you may check a claim against the data, examine the method, identify possible confounders, and ask whether another explanation fits the evidence.

The details vary. The standard does not: do not submit AI-assisted work until you have checked the parts that matter.

The most common failure: students verify that the AI answer looks right, not that they can explain why it is right.

The AI answer is a draft, not a witness.

Where possible, return to the same problem type, or a mixed set of related problem types, one to three days later without AI. This spaced and interleaved retrieval is what converts a single working session into durable understanding (Rohrer and Taylor, 2007). Verification confirms that you understood the work today; spaced and interleaved practice confirms that you will still understand it next week.

8.5 Reflect and disclose

Reflection helps you notice what you learned. Disclosure helps your instructor understand how the work was produced.

To make reflection concrete, after each AI-assisted session answer in your own words: Why does each step in this proof, solution, or argument work? What would change if one assumption were weakened or removed? These two questions take thirty seconds and turn the session into elaborative practice. A useful weekly habit is to schedule two or three short interleaved retrieval sessions using your hard-problems list or verification checklist; the pattern is short, mixed, and unassisted.

When disclosure is required, use clear and ordinary language. You do not need to apologize for allowed AI use. You do need to be honest.

A good disclosure block looks like this:

Tools used: I used an approved AI tool.

Tool role: I used it to generate practice questions and critique my draft introduction.

What I accepted or changed: I used two suggested questions, but rewrote the introduction myself.

Verification: I checked the factual claims against the assigned reading and removed one unsupported claim.

What I can explain: I can explain the final thesis, evidence, and revision choices without AI.

Some courses will require more detail. Some will require none for minor uses. Follow the assignment instructions. When an instructor asks to see the AI conversation itself, thread sharing (section 14.4) is the strongest disclosure available; for most AI-assisted work it makes the written block unnecessary.

9 Four assignment categories

Each assignment should tell you what kind of AI use is allowed. When the instructions are unclear, ask before using AI in a way that affects submitted work.

Instructors increasingly mark their assignments with one of four short labels that signal how AI fits the learning goal. The labels are color-coded; the colors carry meaning. Three of them sit on a traffic light:

- **NAI** (No AI): red. Stop. The assignment measures what you can do without assistance.
- **AIT** (AI Tutor): yellow. Proceed with care. AI may coach you toward understanding, but the submitted work must be yours.
- **AIC** (AI Collaborator): green. Go. AI may co-produce work with you, and the assignment treats the collaboration as part of the learning.

The fourth label is a different kind of sign:

- **AIS** (AI Study): blue. Off the traffic light entirely. The question is not how much AI you may use but what AI is doing—you study AI itself rather than using it as a tool.

NAI, **AIT**, and **AIC** vary along one axis: how much AI freedom the assignment allows. **AIS** is not a point on that axis. It is a different question altogether, which is why it gets a different color.

9.1 **NAI**: No AI

In an **NAI** assignment, AI use is not allowed. The purpose is to measure what you can do independently.

Examples include closed-book exams, short quizzes, in-class writing, oral explanations, lab practicals, live coding checks, language drills, or foundational calculations.

An **NAI** assignment is not a claim that AI is bad. It is a way to protect Core Competence. You need some abilities that do not depend on a tool being open in another window.

9.2 **AIT**: AI as Tutor

In an **AIT** assignment, AI may be used for study support, but not to produce submitted work.

You may ask AI to explain a concept, quiz you, give hints, make flashcards, create practice problems, or help you prepare for class. You should not ask it to write the final answer.

AIT use is common when the instructor wants you to learn from AI without letting AI shape the submitted product. The Study Partner Protocol (SPP) in the companion repository ([companion/spp/](#)) is designed for exactly this kind of AIT work. When **AIT** work produces a tutoring conversation worth showing your instructor, thread sharing (section 14.4) documents the interaction more cleanly than any written summary.

9.3 **AIC**: AI as Collaborator

In an **AIC** assignment, AI may help with brainstorming, drafting, debugging, revision, critique, outlining, coding, or comparison. You must still verify the result and disclose AI use when required.

AIC assignments are useful when the learning goal includes process, revision, research, modeling, design, computation, or professional practice. The instructor is not only asking what you know. They are also asking whether you can use tools responsibly. Because the AI's contribution is part of what is being assessed, thread sharing (section 14.4) is often the most natural way to disclose how the work was produced.

9.4 AIS: AI as Object of Study

In an **AIS** assignment, you study AI output directly. You may compare AI responses, find errors, test generated code, identify bias, repair a flawed explanation, evaluate sources, or analyze how AI changes work in a field.

AIS assignments are especially valuable because they train you to examine AI instead of merely using it.

10 Good uses of AI

Good AI use keeps you mentally active. It helps you see, test, practice, revise, or explain. It does not remove the main work of the course.

10.1 Use AI to clarify confusing ideas

AI can often provide a second explanation when a lecture, reading, or textbook passage does not make sense on the first pass.

Useful prompts include:

Explain this concept in plain language, then give a college-level version.

Give me a concrete example and a non-example.

Ask me three questions to check whether I understand this idea.

After reading the response, return to the course material. The AI explanation is a bridge, not the authority.

10.2 Use AI to practice retrieval

One of the strongest learning habits is trying to recall information without looking at the answer. Practice testing and spaced practice have broad support in learning research (Dunlosky et al., 2013). AI can help you create these practice opportunities.

A good routine is to ask AI for questions, close the AI response, answer on your own, and then check your work.

Create ten practice questions on this topic. Do not give the answers until I ask.

Quiz me one question at a time. Wait for my answer before giving feedback.

This matters because recognition is not the same as recall. Reading an answer and thinking “that makes sense” is weaker than producing the answer yourself.

10.3 Use AI to test your explanations

One of the most reliable ways to discover what you do not yet understand is to try teaching it. Explaining a concept or a solution in your own words, with the steps in your own order, forces you to produce the structure of an idea rather than recognize it. AI can play the role of the listener, the curious student, or the skeptical colleague, whatever audience you need to practice for.

I'll explain this concept to you. Tell me which parts are unclear and which parts are incomplete.

Act as someone outside this field. Listen to my explanation and ask the questions a non-expert would ask.

Here is my answer to this problem. Critique it as if I were teaching you the topic. Where would a student be confused?

The shift in roles matters. When you ask AI to solve a problem, AI does the producing and you evaluate. When you ask AI to evaluate your explanation, you do the producing and AI evaluates. The second form keeps the cognitive work where it needs to be. The same shift lets you rehearse for different audiences, which is one of the more useful side benefits of teaching yourself by teaching AI.

10.4 Use AI to find weak points

AI can help you find places where your work may be unclear or incomplete. This is often more valuable than asking it to improve the work directly.

Here is my argument. What would a skeptical reader challenge?

Find the first step in this reasoning that needs more justification.

What assumptions am I making that I have not stated?

When you use AI this way, you remain the author. AI becomes a critic.

10.5 Use AI to compare approaches

A good way to deepen understanding is to compare two methods, two interpretations, or two explanations.

Compare these two solution strategies. What does each one reveal?

Give me two possible interpretations of this passage, then explain what evidence would support each one.

Show me why this programming approach might be less stable or less efficient than another approach.

Comparison helps you notice structure. It can also reveal when an AI response is too confident.

10.6 Use AI to prepare for direct questioning

In small courses, your instructor may ask you to defend your work orally. In large courses, you may face a short audit, an explanation video, or an in-class follow-up question.

You can use AI to practice.

Act as my instructor. Ask me questions about this solution until you find what I do not understand.

Ask me to explain each step of this argument. Do not move on until my explanation is clear.

Give me a flawed version of this answer and ask me to find the error.

This is a strong use of AI because it prepares you to think without AI.

When it seems like magic, you need more information.

11 Annotated dialogues: strong AI use

The next pages show what good AI use looks like in practice. Each dialogue is a short exchange between a student and an AI tool. Marginal annotations show what the student did right or wrong at each turn. Read the dialogues alongside the principles in section 8 and section 10; the principles describe the habits, the dialogues show the habits at work.

These dialogues draw heavily from mathematics, programming, and humanities research. Students in social sciences, professional programs, and creative or design fields should translate the patterns to their own disciplines; the underlying compact (try first, ask, work, verify, reflect) is field-independent, and the moves illustrated below transfer with minor adaptation.

Dialogue 1: Undergraduate optimization problem (calculus)

Student. I have a calculus problem: a farmer wants to fence a rectangular field of area 400m^2 along a river, fencing only three sides. What dimensions minimize fencing cost? I am stuck setting up the problem.

→ Good move: stated the problem and named where they are stuck.

Student. Don't solve it. Give me one hint about the main obstacle so I can try first.

→ Good move: requested a hint, not a solution.

AI. *One hint: you have two unknowns (length and width) but the area constraint relates them. Use the constraint to write the fencing cost as a function of one variable.*

Student. OK. Let length parallel to river be L , width be W . Area constraint: $L \cdot W = 400$, so $L = 400/W$. Fencing is $L + 2W = 400/W + 2W$. I want to minimize $f(W) = 400/W + 2W$. I'll take the derivative.

→ Good move: did the substitution and set up the function before consulting AI again.

Student. $f'(W) = -400/W^2 + 2$. Setting to zero: $W^2 = 200$, so $W = \sqrt{200} \approx 14.14\text{m}$, and $L = 400/14.14 \approx 28.28\text{m}$.

Student. Can you check whether this is a minimum, not a maximum?

→ Good move: asked AI to confirm, not to redo the work.

AI. *Take the second derivative: $f''(W) = 800/W^3$. At $W = \sqrt{200}$, f'' is positive, so the critical point is a local minimum. Sanity check: as $W \rightarrow 0$ or $W \rightarrow \infty$, $f(W) \rightarrow \infty$, so the critical point is the global minimum on the relevant domain.*

Student. Got it. Final answer: $W = \sqrt{200}\text{m}$, $L = 2\sqrt{200}\text{m}$. I can explain why the second derivative test confirms it's a minimum, and I can state what would change if the constraint were different.

→ Good move: stated what they can explain. Disclosure block would note: AI used for one hint and one verification check; setup, derivative, and conclusion are the student's.

Dialogue 2: Graduate analysis (proof audit-and-repair)

Student. I have an attempted proof that every continuous function on $[0, 1]$ attains its supremum. I’m going to write it first, then ask you to audit it.

→ Good move: produced the artifact before AI involvement.

Student. My proof: “Let $M = \sup_{x \in [0, 1]} f(x)$. By definition of supremum, there exists a sequence $\{x_n\} \subset [0, 1]$ with $f(x_n) \rightarrow M$. Since $[0, 1]$ is bounded, the Bolzano–Weierstrass theorem gives a convergent subsequence $x_{n_k} \rightarrow x^*$. By continuity, $f(x_{n_k}) \rightarrow f(x^*)$. Therefore $f(x^*) = M$.”

Student. Audit this. Find any step that is incomplete or relies on an unstated hypothesis.

→ Good move: asked for an audit, not a rewrite (GMIP D2: proof audit; see Appendix A).

AI. *Two issues. First, you used Bolzano–Weierstrass, which requires the sequence to be bounded; you have $\{x_n\} \subset [0, 1]$, so this is fine, but you should state it. Second, you concluded $x^* \in [0, 1]$ implicitly, but $[0, 1]$ is closed, so the limit lies in $[0, 1]$. State this. Third (subtle): you wrote that sup attains a sequence approaching it. This requires the sup to be finite. You have f continuous on a compact set, but you haven’t established $M < \infty$. The argument circles: you’re using the result you want to prove.*

→ Note: AI flagged a real subtle gap. The student must verify whether this critique is correct.

Student. The third point is wrong, I think. The supremum of any set of real numbers is well-defined as either a real number or $+\infty$. If $M = +\infty$, then we can still find x_n with $f(x_n) \rightarrow \infty$, take a convergent subsequence $x_{n_k} \rightarrow x^* \in [0, 1]$, and continuity would force $f(x^*) = \infty$, contradiction. So M is finite, but I should make this explicit.

→ Excellent move: student pushed back on AI critique, reasoned through it, and identified what to repair (make finiteness explicit) without accepting the framing as written.

Student. Revised proof: “Let $M = \sup_{x \in [0, 1]} f(x) \in \mathbb{R} \cup \{+\infty\}$. We first show $M < \infty$: ...”. The repair clarifies the structure rather than contradicting my original argument.

→ Disclosure: AI used for proof audit; student verified each critique against textbook definitions and rejected one. Original argument and repair are the student’s.

Dialogue 3: Graduate research (literature review with verification)

Student. For my qualifying-exam topic on percolation theory, give me a list of foundational papers I should read.

AI. *Here are six often-cited foundational papers: Broadbent and Hammersley (1957) introducing percolation; Kesten (1980) on the critical probability for 2D bond percolation on $\mathbb{Z}^2$; Grimmett’s textbook Percolation (2nd ed., 1999); Smirnov (2001) on conformal invariance for site percolation on the triangular lattice; Lawler–Schramm–Werner (2001) on intersection exponents; and Aizenman–Newman (1986) on tree graph inequalities.*

Student. Before I cite any of these, I’m going to verify each one independently by checking MathSciNet (the AMS bibliographic database) or the journal record.

→ Good move: applied the verification standard from appendix C.

Student. [Several minutes later.] I verified five of the six. The Aizenman–Newman 1986 paper does exist, but the actual title and journal are different from what you described. The reference appears to be “Tree graph inequalities and critical behavior in percolation models,” *Journal of Statistical Physics* 36 (1984), pp. 107–143, not 1986. I’m using the verified version.

→ Good move: caught a subtle hallucination—real authors and topic, but wrong year. This is exactly the kind of error AI literature suggestions produce (Bender et al., 2021).

Student. Update: I should also flag for my advisor that I will note in the disclosure that the initial reference list came from AI but every entry was checked against MathSciNet before inclusion.

→ Good move: prepared an honest disclosure block consistent with appendix C.

Dialogue 4: Programming (debugging by diagnostic questions)

Student. My Python function is supposed to compute the moving average of a list with window size k , but it gives wrong values at the edges. Don’t fix it for me. Ask me diagnostic questions until I find the bug.

→ Good move: asked for diagnostic questions, not corrected code.

Student. Code: `def moving_avg(arr, k): return [sum(arr[i:i+k])/k for i in range(len(arr))]`

AI. *Three diagnostic questions. First: what does `arr[i:i+k]` return when `i+k > len(arr)`? Second: how many elements does `arr[i:i+k]` contain in that case? Third: are you dividing by the right number when the slice is shorter than `k`?*

Student. Oh. Python slices return a partial slice if the end goes past the array, but I’m still dividing by k . So at the edges, I’m dividing a shorter sum by the wrong denominator. The fix is either to skip those edge cases or to divide by `len(arr[i:i+k])`.

→ Good move: identified the bug from the diagnostic questions, articulated two valid fixes.

Dialogue 5: Humanities (background research with primary sources)

Student. I’m writing a paper on agricultural reform in late 18th-century England. Give me background context on the Enclosure Movement; I will then return to my assigned primary sources.

AI. *The Enclosure Movement converted common-use lands to private ownership. Parliamentary enclosures peaked between roughly 1760 and 1830. Landlords benefited from the legal*

restructuring; many small tenants and commoners lost access to grazing and gathering rights. Critics like William Cobbett saw it as displacement; defenders argued it improved agricultural productivity.

Student. This is helpful as orientation. Now I'm going back to the primary sources assigned for the seminar: a parliamentary act from 1773 and an excerpt from Cobbett's *Rural Rides*.

→ Good move: treated AI as orientation, not as the source.

Student. [After reading.] My initial claim was going to be "Enclosure caused widespread rural displacement." But the parliamentary act I read is narrower: it permits enclosure of specific commons in one parish, with compensation procedures and dispute mechanisms. Cobbett's account, written 50 years later, is the more rhetorical critique. I need to narrow my claim.

→ Good move: revised the claim based on what the primary sources actually said, not on AI summary.

Student. Revised thesis: "In the parish of [X], the 1773 enclosure act provides a structured but contested transfer of rights, with formal procedures that nonetheless produced individual hardship cases visible decades later in critical literature."

→ Disclosure: AI used for general background; thesis and argument grounded in assigned primary sources.

These five dialogues illustrate one consistent pattern. The student does the substantive work first, uses AI for narrowly defined help, verifies what AI provides, and can defend the final result. AI is present in each dialogue, but it does not displace the student's thinking.

A borderline case

Not all AI use is clearly strong or clearly weak. Some uses sit in a gray zone where the student must exercise judgment about what to keep. Dialogue 6 shows one such case.

Dialogue 6: Borderline use (writing, certainty drift)

Student. Here's a paragraph from my history paper. Tighten the prose; keep my meaning.

Student. [Student's original]: "The evidence in the parish records suggests that enclosure produced uneven outcomes for small tenants. Some lost grazing access; others received small compensation. The records do not show a single uniform pattern."

AI. [AI revision]: "*The parish records demonstrate that enclosure produced sharply unequal outcomes for small tenants. Many lost grazing access while a few received compensation. The records reveal no uniform pattern.*"

→ Watch the meaning drift. “Suggests” became “demonstrates.” “Some / others” became “many / a few.” “Do not show” became “reveal no.” Each change tightens the prose, and each one inflates a hedged claim into a confident one.

Student. Hmm. The AI version reads better, but it’s making a stronger claim than my evidence supports. I have records from one parish, not a body of evidence about “the parish records” in general.

→ Good move: caught the certainty drift. This is the verification step from appendix C: AI revision can subtly change the certainty level of central claims.

Student. I’ll keep the AI’s sentence rhythm but restore my hedging: “The parish records I examined suggest that enclosure produced uneven outcomes for small tenants in this parish. Some lost grazing access; a few received compensation. The records show no uniform pattern across the cases I reviewed.”

→ Borderline use that landed in the right place. The student kept the prose improvement but rejected the certainty inflation.

This case sits in a gray zone because the AI did exactly what was asked: it tightened the prose. The certainty drift was a side effect, not a failure of the prompt. The student is responsible for catching it. A student who accepts the AI revision unread would have published a stronger claim than their evidence supports, with no malice and no policy violation, but with reduced scholarly accuracy. The verification habit is what catches these cases.

12 Dangerous uses of AI

Dangerous AI use often feels helpful at first. It saves time, removes friction, and produces a polished answer. The problem appears later, when you cannot remember, explain, defend, or transfer what you submitted. Cognitive science offers one explanation for why this happens: fast, intuitive thinking is comfortable with surface fluency and tends to skip the slower, effortful checking that catches errors (Kahneman, 2011). AI output that reads smoothly engages the fast system; verification requires the slow one.

The easiest person for you to fool is yourself.

12.1 Do not outsource first contact with the problem

If AI sees the problem before you do, it can shape your thinking too early. You may stop noticing what confused you. You may accept the first framing it gives you. You may never develop your own plan. UNESCO’s guidance on generative AI in education characterizes

the risk this way: when a student starts an assignment with an AI-generated outline rather than developing their own, the practice can amount to writing without thinking (Miao and Holmes, 2023). Verification cannot substitute for that lost thinking, because the only thing being verified is the AI's plan, not the student's understanding.

For serious assignments, begin with your own attempt. Even a rough attempt helps.

12.2 Do not submit work you cannot explain

This is the clearest boundary.

If you cannot explain a paragraph, proof, calculation, source choice, design decision, or line of code, it should not be in your submitted work. You may use AI to help you understand it. You may not use AI to hide the fact that you do not understand it.

12.3 Do not trust generated citations

AI systems can invent sources or misstate what real sources say. A citation that looks real may still be wrong, and guidance on generative AI in education treats source verification as a central part of responsible use (Miao and Holmes, 2023).

Never cite a source you have not checked. For research assignments, you should find the source yourself, read the relevant part, and confirm that it supports the claim.

12.4 Do not let AI change your meaning

AI revision can improve clarity, but it can also change what you meant. It may make a claim stronger than your evidence allows. It may remove a useful qualification. It may replace your voice with a generic academic tone.

After using AI to revise, compare the new version with your original intent. Ask yourself: Is this still my claim? Is the level of certainty right? Would I defend this sentence if asked?

12.5 Do not trust code just because it runs

Code can run and still be wrong. It may solve the wrong problem, use the wrong formula, mishandle edge cases, hide numerical instability, or produce plausible-looking output from flawed logic.

If AI helps write code, you should test it on simple cases where you know the answer. You should read the code line by line. You should know what each major step does.

12.6 Do not use AI to evade the purpose of an assignment

Every assignment has a learning purpose. Sometimes the purpose is to practice a skill slowly. Sometimes it is to struggle with a hard text. Sometimes it is to learn a method well enough to use it later.

AI can defeat that purpose if you use it too early or too broadly. Before using AI, ask: What is this assignment trying to teach me? Will this use help me learn that, or help me avoid it?

13 How to verify AI-assisted work

Verification is not one more task added to the assignment. It is part of responsible learning. The cognitive habit underneath verification is what learning science calls self-explanation: stating in your own words how a new idea relates to what you already know, and why a claim is true. Self-explanation has measurable benefits for retention and transfer (Dunlosky et al., 2013), and it does not depend on AI being involved at all—which is precisely why it is the right habit for verifying AI-assisted work. Figure 7 summarizes how verification looks across disciplines. Graduate students preparing manuscripts for publication will find a fuller verification protocol in Appendix C, which addresses research-grade work specifically.

Sage advice from Dick Hamming, whom I knew during his visits to Los Alamos: do the back-of-the-envelope calculation before you build the cathedral. The principle scales. In an AI-assisted setting, a rough check on the answer’s shape and order of magnitude is worth more than a precise validation of the wrong thing. Verification starts with what you would expect to see, not with what the AI gave you.

Use the following questions before submitting AI-assisted work.

Meaning. Does the final work say what I mean?

Accuracy. Are the facts, claims, computations, and sources correct?

Method. Is the method appropriate for the question?

Evidence. Does each major claim have support?

Assumptions. Have I identified the assumptions being made?

Ownership. Can I explain this without reading from the AI response?

The strongest single ownership check is the **translation test** introduced in section 6. Close the AI window, take a key claim or step, and re-present it in a different form—different notation, a different example, a visual representation if the original was symbolic, or a plain-language explanation if the original was formal. If you can move fluently to another mode, the work has become yours. If you cannot, return to the underlying material before submitting.




Figure 7: **Verification across disciplines.** The verification habit is constant; the specific checks vary by domain. Every discipline has a small set of routine checks that catch most AI errors. Students should learn the verification toolkit appropriate to each course they take.

13.1 Verification in STEM courses

In STEM courses, verification often means testing the work against structure.

You might check units, signs, limiting cases, special cases, order of magnitude, known formulas, benchmark examples, code output, or experimental data. You might solve a simpler version of the problem. You might compare two independent derivations. You might ask whether the result behaves sensibly when an input becomes very large, very small, or zero.

For code, useful checks include small test cases, boundary cases, reproducibility, comments explaining the logic, and comparison with a known implementation.

For laboratory work, useful checks include instrument limits, measurement error, procedural consistency, data provenance, and whether the conclusion follows from the actual results.

For mathematics and quantitative work, a quick verification kit is:

- Check a trivial or special case (for example, $x = 0$, $x = 1$, or a known solution).
- Check a limiting case (large or small parameter, asymptotics).
- Check units, dimensions, and signs for each term.
- Compare to a known formula or theorem statement.

- Run one independent derivation or route (for example, differentiate your integral result).

These five checks rarely take more than a few minutes. If a result fails one of them, the underlying work needs another look before submission.

13.2 Verification in writing courses

In writing courses, verification often means checking meaning, evidence, and voice.

After AI-assisted revision, compare the draft to your original purpose. Did the thesis change? Did the AI add a claim you did not intend? Did it make your tone too formal, too vague, or too certain? Did it remove an important nuance?

A useful practice is to write a short author's note after revision:

In this revision, I changed the structure by moving the evidence earlier. I rejected the AI suggestion to make the claim stronger because the source did not support that level of certainty.

That note helps you remain the author.

13.3 Verification in humanities courses

In humanities courses, verification often means returning to the text or source.

AI can summarize, but it can also flatten ambiguity. It may miss irony, context, genre, translation issues, or the difference between a primary source and a later interpretation. If you use AI to explore a text, return to the passage itself. Ask whether the evidence supports the interpretation.

Do not let AI replace close reading.

13.4 Verification in social science courses

In social science courses, verification often means examining evidence, method, and causal claims.

Ask whether the data support the conclusion. Ask whether there are confounders. Ask whether a claim is descriptive, causal, predictive, or normative. Ask whether the sample is appropriate. Ask whether the method can answer the question being asked.

AI may help you generate alternative explanations. It cannot decide which explanation the evidence supports unless you check the evidence yourself.

13.5 Verification in professional and applied courses

In professional courses, verification often includes ethics, privacy, standards, and consequences.

If you are working with cases, clients, patients, designs, policies, or institutional data, ask whether AI use is permitted. Check whether the work includes sensitive information. Confirm that recommendations follow professional standards and course expectations.

A fluent answer can still be unsafe, unfair, or irresponsible.

13.6 Verification in graduate research

Graduate research raises the verification bar. The standards of independent competence and tool-augmented practice that apply in coursework apply with extra force in research, where errors propagate through manuscripts, citations, and into the literature itself. Graduate students who use AI in their research should treat verification as part of the research workflow rather than a final check before submission. The full protocol appears in Appendix C; this subsection summarizes the core practices.

A graduate researcher should be able to defend every result, citation, computation, and interpretation in their work without reference to AI. The defense need not be unaided in real time, but the underlying understanding must be the researcher’s own. If a reviewer asks “why did you choose this approach,” the answer should reflect scientific judgment, not tool-mediated procedure. Three workflow practices support this standard.

Independent attempt before tool use. Begin every nontrivial research task with the researcher’s own attempt: a proof outline, a model setup, a literature shortlist, a code sketch. The independent attempt is what AI assistance refines. If AI sees the problem first, the researcher’s framing risks becoming AI’s framing; a published paper that originated this way often shows it in subtle inconsistencies that reviewers find before authors do.

Citation discipline. Every cited source must exist, must be read, and must support the claim being made. AI tools can produce citations that look correct but reference the wrong year, the wrong journal, or the wrong author combination, a failure mode documented across early research on AI-assisted writing (Bender et al., 2021). Verification against MathSciNet, Web of Science, the publisher’s own record, or the journal page itself is non-negotiable. Dialogue 3 in section 11 illustrates exactly this kind of catch.

Audit-and-repair. For mathematical and computational work, AI-produced derivations and code should be treated as drafts to audit, not as results to accept. The audit identifies where each step uses each hypothesis, what edge cases are excluded, and where the derivation could fail. The graduate-mathematics version of this practice is GMIP D2 (proof construction and proof repair), described in detail in Appendix A; analogous practices exist in computational science (input validation, boundary tests, comparison against analytical solutions) and in empirical work (sensitivity to parameter choices, robustness across specifications).

These practices are not separate from research. They are research, in an environment where AI-fluent text can pass shallow inspection. Researchers who internalize them publish more carefully, defend their work more easily under questioning, and contribute to a literature that remains trustworthy as the volume of AI-assisted scholarship grows.

14 Disclosure

Disclosure is not a confession. It is a record of process.

When an instructor permits AI use and asks for disclosure, your goal is to be factual and clear. State what tool you used, how you used it, what it contributed, what you changed, and how you verified the result.

14.1 A short disclosure template

Use this when AI played a small role.

AI disclosure: I used an approved AI tool to generate practice questions, check grammar, or suggest revision questions. I reviewed the output and made the final choices myself.

14.2 A fuller disclosure template

Use this when AI shaped the process more substantially.

Tools used: Name the tools used, if any.

Tool role: Describe what the tools helped with.

Student role: Describe what you wrote, revised, checked, or decided.

Accepted, changed, or rejected: Describe the AI contributions you used, revised, or set aside.

Independent verification: Describe how you checked important claims, sources, data, code, or reasoning.

What I can explain without AI: Name the main parts of the work you can now explain on your own.

14.3 What good disclosure does

Good disclosure helps your instructor see your thinking. It also protects you from two mistakes: pretending AI was not used when it was, and giving AI more credit than it deserves.

The level of disclosure should match the assignment. Minor grammar help may need only a brief note. A project that used AI for brainstorming, code, revision, and source checking needs more detail.

14.4 Thread sharing

When an instructor asks to see the AI conversation that produced your work, sharing the chat link or transcript is now a routine option. Most major platforms (ChatGPT, Claude, Gemini, Copilot, and others) support a share function that produces a read-only link to the conversation. The companion repository's `students/sharing-ai-conversations.md` guide describes the operational steps for each major platform.

Thread sharing is the strongest form of disclosure available to a student. A written disclosure template, however careful, depends on your account of what the AI contributed. Thread sharing lets your instructor see what was actually asked and what was actually returned. The gap between your account and the underlying interaction closes. For an instructor's purposes, this is sometimes the difference between trusting the work and having to ask further questions.

Two honest limits deserve naming. First, thread sharing presumes the shared thread is the one that actually produced the work. A student who runs a polished AI study session for show while drafting the actual work in a separate, unshared chat has not solved the disclosure problem; they have only relocated it. The translation test in section 6 is the structural safeguard that closes this gap, because the test asks what the student can do regardless of which thread produced what. Second, thread sharing must respect privacy: the considerations described in section 15 apply to anything visible in a shared thread, including identifying information about classmates, instructors, or third parties. Edit or redact before sharing if any such material appears, and ask your instructor for an alternative if redaction would distort the record.

15 Privacy and data protection

Before uploading material to an AI system, ask whether you have the right to share it. Many AI tools are run by outside companies. Some retain data. Some may use inputs to improve services. Some institutional tools may have stronger privacy protections, but you should not assume that every tool is safe for every kind of information.

Unless your instructor or institution has clearly approved the tool for that use, do not upload:

- private information about classmates;
- unpublished research data without permission;
- patient, client, or human-subject information;

- restricted exam questions or solutions;
- confidential institutional documents;
- private correspondence;
- recommendation letters or evaluations;
- copyrighted course material if your instructor has not permitted it;
- personal data that could identify someone else.

If you are unsure, do not upload the material. Ask your instructor how to proceed.

Privacy is not only a legal concern. It is part of academic trust. Ethical guidance for AI and data use in teaching stresses that learners and educators should handle data systems with care, transparency, and attention to risk (European Commission, Directorate-General for Education, Youth, Sport and Culture, 2022).

16 Small classes and the Oral-Defense Model

Small classes often allow more flexible AI use because the instructor can ask students directly about their work. Examples include seminars, capstones, advanced labs, writing workshops, and many graduate courses. The defining feature is class size, not course level: a small introductory seminar uses the same Oral-Defense pattern as a small graduate seminar, even though the assignment-category mix and AI autonomy granted are likely to differ.

In these courses, you may be asked to defend AI-assisted work in conversation. That defense may look different across fields.

STEM. Explain a derivation, model, computation, experiment, or code segment.

Writing. Explain your thesis, structure, evidence, and revision choices.

Humanities. Defend an interpretation by returning to the text or source.

Social sciences. Explain your method, evidence, assumptions, and limits.

Professional programs. Justify a recommendation, privacy choice, or ethical judgment.

Creative work. Explain process, iteration, constraint, and authorship.

The main standard is the same:

If you submitted it, you should be able to explain it.

You can use AI to prepare for this kind of questioning. Ask it to play the role of a skeptical instructor. Ask it to question your assumptions. Ask it to find the weakest part of your explanation. Then practice answering without looking at the AI response.

17 Large classes and the Evidence-of-Thinking Model

Large classes cannot rely on one-on-one oral exams for every student. Instead, they often use scalable evidence of thinking. Most large classes are introductory or service courses simply because that is where enrollment concentrates, but the model is class-size-driven, not level-driven: a large advanced course uses the same evidence-of-thinking patterns.

You may be asked to complete short no-AI quizzes, write in class, submit a draft history, explain one step from your homework, critique a flawed AI answer, provide a process log, verify sources, record a short explanation, or participate in a random oral audit.

These practices are not meant to assume misconduct. They are meant to protect learning. In a world where AI can produce polished answers, instructors need ways to see the work behind the answer.

A common pattern is the homework-to-quiz check. You complete AI-allowed homework, then answer a short no-AI question about the same idea in class. This checks whether the homework became part of your understanding.

Another pattern is the audit-and-repair task. You receive an answer that may contain a subtle error. Your job is to find the error, explain why it matters, and repair it. This is one of the most useful AI-era skills because it rewards judgment over fluency.

Many courses fall between the small-class and large-class extremes — too large for individual oral defenses on every assignment, but small enough for some direct contact, or arranged so that lab time supports oral defense even at medium enrollment. Part II (section 37) describes the resulting Hybrid Model that combines elements of both. As a student in a hybrid course, your experience may include occasional oral conferences (often at the start of the term for calibration, then again at major project milestones) alongside continuous process documentation through Plan-Use-Verify artifacts on regular assignments. The expectations are the same: explain what you submitted, verify what you used, own the work.

18 Asking better questions in mathematics

Section 8.2 introduced the general principle: a useful AI question keeps the central thinking with you. This section extends that principle to mathematics specifically, where AI use splits cleanly into two patterns. For computational work (calculus, numerical methods, modeling), strong prompts ask AI to check, debug, and probe rather than to solve. For proof-based work (real analysis, abstract algebra, topology), strong prompts ask AI to critique, generate examples, and identify gaps rather than to write the proof.

For computational mathematics:

Weak prompt:

Solve this problem.

Stronger prompt:

Ask me one question that would help me decide which method applies.

Weak prompt:

Explain compactness.

Stronger prompt:

Give me one example, one non-example, and one common false belief about compactness.

For proof-based mathematics, the same pattern looks like this. The strong prompts preserve the student's role as the proof-builder.

Weak prompt:

Prove this theorem for me.

Stronger prompt:

Here is my proof attempt. Identify the first point where the logic may fail, but do not rewrite the proof.

Weak prompt:

Check my proof.

Stronger prompt:

Identify whether each hypothesis in the theorem is actually used in the proof. If any hypothesis is unused, ask whether the theorem can be strengthened by removing or weakening that assumption.

These prompts train the use of AI as a mathematical coach rather than a solution engine.

There is a deeper pattern worth naming. Cognitive psychologists have spent decades evaluating which study techniques actually improve learning. Practice testing (quizzing yourself), distributed practice (spreading study over time), elaborative interrogation (asking why a fact is true), and self-explanation (stating in your own words how a new idea relates to

what you already know) all show measurable benefit. Highlighting, rereading, and producing summaries of unread material show much weaker benefit (Dunlosky et al., 2013).

AI tools are study-technique amplifiers. Whatever technique you ask AI to support, AI will support it efficiently. Asked to summarize a reading you have not yet engaged with, AI amplifies a technique that has limited learning value. Asked to highlight what is important, AI amplifies a technique that has limited learning value. Asked to test you on material you have studied, AI amplifies a technique that has strong learning value. Asked to ask you why something is true, AI amplifies a technique that has strong learning value. The prompt determines which amplification you receive. This is why weak prompts tend to ask AI to do the low-utility work for you, while stronger prompts tend to keep the high-utility work yours.

19 A practical study routine

Here is a weekly routine for using AI without weakening your learning.

Before class, skim the assigned material and write down two questions. If you use AI, ask it to explain background concepts, not to replace the reading.

After class, close your notes and write a short summary from memory. Then use AI to quiz you. Ask for one question at a time. Answer before seeing feedback.

When starting homework, attempt each problem first. If you are stuck, ask for a hint or a simpler related example. After finishing, ask AI to check for possible gaps, but verify the final answer yourself.

When writing, draft first. Then ask AI to identify unclear points, missing evidence, or possible objections. Revise in your own words.

Before an exam or oral defense, ask AI to question you. Practice explaining without looking at notes. Focus on the places where you hesitate.

A simple way to make these habits stick is to convert them into when-then plans. Some examples:

- When I receive AI feedback on a proof or paragraph, then I close the AI window before writing my final version.
- When I finish AI-assisted homework, then I work one nearby problem without AI before submitting.
- When I cannot explain a step in my own words, then I do not submit the work yet.

The when-then form turns a vague intention into a trigger-action plan and tends to improve follow-through compared to general resolutions.

This routine uses AI as a scaffold. It does not let AI carry the weight.

20 Student checklist

Use this checklist before submitting AI-assisted work.

- ☐ I know the AI policy for this assignment.
- ☐ I made my own first attempt before using AI for substantial help.
- ☐ I used AI for support, not replacement.
- ☐ I checked important claims, sources, calculations, or interpretations.
- ☐ I removed any AI-generated material I could not explain.
- ☐ I protected private, restricted, and confidential information.
- ☐ I disclosed AI use if the assignment required it.
- ☐ I can explain the final work in my own words.
- ☐ I can close AI and my notes and redo the key argument, calculation, proof, or paragraph from a blank page.
- ☐ I know what I learned from the process.

The blank-page item is the strongest single calibration check on the list. Explanation while looking at finished work is easy to confuse with understanding; reconstruction from a blank page is harder to fake.

21 Discipline examples

The verification subsections in Section 13 listed verification techniques per discipline. This section walks through five short vignettes showing the whole Learning Spiral in action across fields: try first, ask for help, work independently, verify, reflect and disclose. Each vignette is a compressed version of an honest week in a course. Your instructor's policy always comes first.

21.1 STEM example

A physics student starts a problem on energy conservation. *Try first*: she sketches a free-body diagram, writes down the conservation equations she remembers, and gets stuck on which surface friction does work. *Ask*: she asks AI for one hint about how to identify the correct system boundary, not for the solution. *Work*: with the hint, she completes the algebra and arrives at a numerical answer. *Verify*: she checks units, tests the limiting case where friction is zero against the textbook formula, and confirms the answer's order of magnitude. *Reflect*

and disclose: she records that AI helped with the system-boundary identification; the rest is hers, and she can explain the steps without notes.

This is responsible use because AI helped her enter the problem without replacing the reasoning.

Within mathematics, the practical use of AI differs between proof-based courses (for example, real analysis, abstract algebra, topology) and computational courses (for example, calculus, numerical methods, modeling). Proof-based work centers on choosing the right argument and verifying that each hypothesis is used; AI is most helpful as a critic of proof attempts and a generator of examples and non-examples. Computational work centers on choosing the right method and verifying the result; AI is most helpful as a debugger of code and a checker of intermediate calculations. The two patterns share the same loop, but the verification moves and the prompt patterns are tuned differently.

21.2 Programming example

A student writes a data-cleaning function. *Try first*: he runs the code and observes that the output looks wrong on rows with missing values. *Ask*: he asks AI to suggest diagnostic tests rather than rewriting the function. *Work*: he builds three small test cases the AI suggested, plus one he thought of himself. *Verify*: the tests reveal that missing values are passed through as zeros rather than handled separately; he revises the function and reruns the tests until they pass. *Reflect and disclose*: his submission includes the tests as part of the work, with a note that AI suggested the diagnostic-test approach; the analysis of test outcomes and the fix are his.

This is responsible use because the student used AI to improve debugging, not to hide ignorance.

21.3 Writing example

A student is drafting an essay on labor history. *Try first*: she writes a full draft from her notes and the assigned readings, identifying the paragraph she's least confident about. *Ask*: she asks AI to identify where the argument loses focus, without rewriting it. *Work*: AI flags her third paragraph; she returns to the source she cited there and finds her evidence is thinner than her claim. *Verify*: she rewrites the paragraph in her own words, narrowing the claim to what the evidence supports; she rejects an AI suggestion that would have inflated her wording back to the original level of certainty. *Reflect and disclose*: her disclosure block notes that AI helped with diagnostic critique on one paragraph; the thesis, evidence, and final wording are hers.

This is responsible use because the student remains the writer and lets the evidence control the claim.

21.4 History example

A student is writing a paper on a primary source. *Try first*: he reads the assigned passage carefully and drafts a one-paragraph interpretation. *Ask*: he asks AI for general historical context about the period, not for the interpretation itself. *Work*: with the context, he returns to the primary source and tests whether his original interpretation still fits, narrowing one claim that turned out to be too broad. *Verify*: he checks AI's contextual claims against the assigned readings and removes one detail that AI mentioned but the readings did not support. *Reflect and disclose*: his disclosure says AI provided period background; the textual interpretation, the narrowing, and the verification are his.

This is responsible use because the primary source remains the authority.

21.5 Social science example

A student is analyzing a study about education and income. *Try first*: she works through the dataset codebook and lists the variables she thinks matter. *Ask*: she asks AI to suggest possible confounders she may have missed. *Work*: AI suggests six; she classifies them into those measured in the dataset and those that are speculative. *Verify*: for each measured confounder, she runs a quick check on whether including it changes the coefficient of interest; for the speculative ones, she notes them as limitations. *Reflect and disclose*: her writeup separates evidence-based statements from acknowledged speculation, with a disclosure block noting AI's role in confounder generation.

This is responsible use because AI helped expand the student's thinking, but the evidence controlled the conclusion.

22 When not to use AI

Sometimes the best use of AI is not to use it.

Do not use AI when the assignment forbids it. Do not use AI when the task is designed to build fluency through practice. Do not use AI when it would expose private information. Do not use AI when you are too tired to check the result. Do not use AI when using it would make you skip the part of the work you most need to learn.

This last case is common. If you are weak at algebra, do not use AI to bypass algebra practice. If you struggle to write introductions, do not ask AI to write every introduction. If you are learning to read difficult texts, do not let AI summarize every text before you have faced it yourself.

Use AI where it strengthens your learning. Avoid it where it removes the struggle that the course is trying to teach you how to handle.

If a question arises about whether work is yours, process evidence matters. AI detection tools are not perfect, and research has found that GPT detectors can misclassify non-native English writing as AI-generated (Liang et al., 2023). Your best protection is honest disclosure, draft history, verification notes, and the ability to explain your work.

22.1 The choice not to use AI at all

Some students choose not to use AI for any coursework, on grounds beyond what specific assignments require. The reasons vary: environmental concerns about the energy cost of large-model inference, labor concerns about how training data was sourced and how AI is reshaping professional work, ethical concerns about complicity in harms the technology produces. From some fields, especially in the humanities, the social sciences, and the arts, the position is sometimes stated more strongly: there is no ethical use of AI in their discipline at the present moment.

This guide is not written to talk those students out of their position. The framework still applies, but the application is different. Every assignment is implicitly **NAI** for a student who has chosen not to use AI; the Learning Spiral becomes the practice they were going to use anyway; the Student's Trilemma names the predicament they are still asked to navigate when peers are using AI freely. What the framework offers is vocabulary for the situations they will find themselves in: grading curves shaped by tools they do not use, group projects where teammates do, courses where the instructor assumes AI use without saying so. Naming the patterns helps a student who has chosen to abstain navigate a coursework environment being reshaped around them.

The institutional implications are developed in Part III.

23 What to do when you are unsure

AI policies vary by course. They may even vary by assignment within the same course. If you are unsure, ask your instructor before using AI in a way that affects submitted work.

A good question is specific:

For this assignment, may I use AI to generate practice questions or critique my draft if I disclose that use?

Another useful question is:

May I use AI to debug code if I include a short note explaining what I tested and what I changed?

Avoid vague questions such as “Can I use AI?” That question is often too broad. Ask about the role you want AI to play.

If you already used AI and are unsure whether the use was allowed, tell your instructor early. It is usually easier to correct a misunderstanding before submission than to explain it afterward. Part II describes how instructors design assignments and what they look for in disclosure; reading the relevant section there can help you frame the conversation.

24 Evidence base and limits of the guide

This guide draws on three kinds of evidence: learning science, emerging research on AI-supported tutoring, and institutional guidance on responsible AI use. The learning-science foundation is mature enough to support several strong claims. Active retrieval, spaced practice, self-explanation, and elaborative interrogation help students learn because they require the learner to do cognitive work (Dunlosky et al., 2013).

The AI-specific evidence is newer and still heterogeneous, but a clear pattern is emerging: AI helps most when it is embedded in active learning, verification, feedback, and independent reconstruction. Across multiple methodologies and disciplines, results converge on the same conclusion: the design of AI use determines whether learning improves, stays flat, or declines. A pre-registered randomized trial with nearly a thousand high-school mathematics students found that an unrestricted ChatGPT-style interface improved practice performance by 48% but reduced exam performance by 17% relative to no-AI controls, while a guardrailed AI tutor improved practice performance by 127% with no exam-performance loss (Bastani et al., 2025). A randomized controlled trial in an introductory physics course at Harvard found that a carefully designed AI tutor produced learning gains over double those of in-class active learning. The tutor was engineered around seven pedagogical practices: active engagement, cognitive-load management, growth mindset, scaffolding, feedback accuracy, timely targeted feedback, and self-pacing. The system prompt handled the first three; the platform’s structured walkthrough handled scaffolding; the remaining three are properties of the AI medium itself (Kestin et al., 2025). Recent meta-analyses confirm a moderate positive overall effect on academic achievement (Sun and Zhou, 2024), with mixed effects on higher-order thinking outcomes (Zhao et al., 2025) that depend largely on whether the use design preserves student cognitive engagement. UNESCO frames AI literacy as a shared responsibility for students and teachers, not only a student skill (Miao et al., 2024; Miao and Cukurova, 2024); the OECD’s 2026 Digital Education Outlook synthesizes the same conclusion at the international policy level (OECD, 2026).

The mechanisms by which unrestricted AI weakens learning are now reasonably well understood. Researchers describe them with overlapping names: cognitive offloading (skipping the diagnostic and evaluative steps of help-seeking), metacognitive laziness (reduced self-monitoring during AI-assisted tasks), and the crutch pattern (asking AI for solutions rather than for help). Neuroimaging evidence suggests reduced cognitive activation when students start with AI rather than with their own attempt (Kosmyna et al., 2025). The Bastani

study adds a metacognitive finding that matters here: students often do not perceive that their learning has been reduced, even when objective assessment shows it has. The guide's emphasis on trying first, verifying, reflecting, and explaining is designed to address each of these mechanisms in turn.

For that reason, this guide avoids a simple pro-AI or anti-AI position. It treats AI as useful when it keeps you active and risky when it lets you skip the work the course is meant to teach. The best test is practical: after using AI, can you explain more, verify more, and do more on your own than you could before?

25 A final word

AI can make schoolwork faster. That is not always good. Some parts of learning are slow because they require attention, memory, judgment, and practice. The goal is not to avoid all difficulty. The goal is to use help wisely.

A student who uses AI well does not hide behind it. They use it to ask better questions, test their thinking, find weak places, and improve their work. They know what the tool contributed. They know what they checked. They know what they still do not understand.

That is the standard this guide asks you to meet.

Use AI when it helps you think. Stop when it starts thinking for you.

26 Reference materials

The following templates are intended for student use during AI-assisted assignments. Instructors may adopt or adapt them. For institution-wide reference materials, see Part III.

26.1 Template: AI account configuration

Use this six-category template to set persistent preferences in your AI platform (sometimes called custom instructions, system prompts, or memory). Copy them, adapt the bracketed parts to your situation, and paste the whole thing into your platform's preferences panel. The motivation for configuring your account at all is given in section 2; the categories below are the operational template.

1. Who you are and what you are studying.

*I am a [year] [undergraduate/graduate] student in [department] at [university].
I am currently taking [course names this semester]. My academic background*

includes [prior coursework relevant to current work]. I am using AI to support my learning, not to substitute for it.

This single paragraph gives the AI enough context to calibrate explanations to your level and your courses. Update it each semester as your courses change.

2. How you want explanations pitched.

Explain reasoning step by step. When I ask you to solve a problem, walk me through the approach before giving the answer; do not just give me the final answer when I am working through something I should be learning. If a concept builds on prerequisites I have not told you about, ask me what I already know rather than assuming.

The last sentence is the most important. It teaches the AI to ask, which prevents the two most common failure modes: explanations pitched too high (where you cannot follow the reasoning) and explanations pitched too low (where the AI assumes ignorance you do not have). Both waste your time.

3. How you want uncertainty handled.

When you are uncertain about a fact, say so explicitly. Do not invent citations, sources, or specific numbers if you do not know them. If a question is at the edge of your knowledge, tell me, so I can verify in textbooks or other sources rather than trusting an answer that might be wrong.

AI systems can produce confident-sounding wrong answers, especially for specific facts (dates, citations, equations from particular papers, statistics). Asking the AI to flag its own uncertainty may help, but it does not remove the need to verify important claims yourself.

4. Writing style preferences.

I write in [American English / British English / English as a second language]. [If English is not your first language: please flag passages where my phrasing might read as unclear or non-idiomatic, and suggest alternatives.] I prefer prose that is [clear and direct / formal / conversational].

The English-as-a-second-language clause matters more than it may seem. Detector tools used by some journals and some courses systematically misread non-native-English writing as AI-generated, an issue covered in Part III and Appendix C. If English is not your first language, telling the AI so, and asking it to help with idiomatic clarity, both improves the writing and gives you a record of what you wrote and what the AI suggested, which protects you if your work is questioned.

5. Scope of use and integrity boundary.

I will use AI for [studying, brainstorming, debugging, getting unstuck, checking my own work]. I will not use AI to [write assignments I am supposed to do myself, generate code I do not understand, produce text I will submit as my own work without disclosure]. When my course or assignment policy says no AI, I will not use AI for that work, regardless of what is in this configuration.

This is partly a reminder to yourself. The configuration does not override your course's rules; the assignment categories described in section 9 explain when AI is permitted, encouraged, restricted, or forbidden, and the assignment-level rules always take precedence over your account-level configuration. Putting the boundary in writing where you see it every time you open an AI conversation helps keep the lines clear.

6. Optional: technical-work accuracy guard.

For mathematical, scientific, or technical answers: show your work. Label any claim you are confident about with [confident] and any claim you are less confident about with [tentative]. If I ask for code, explain what each non-obvious line does.

This is optional because it adds complexity to every answer. For students in mathematics, the sciences, and engineering, it is the highest-leverage single addition you can make. The labels train you to read AI output critically by making the AI's confidence level visible, and the show-your-work instruction makes the verification practices in section 8 easier because the work is already laid out.

A living document. The starter template above captures the categories that experienced academic users have found matter most. As you use AI more, you will discover specific failure modes in your own work and want to add specific instructions to prevent them. This is normal and expected. Treat your AI preferences as a living document. Periodically review what is there, prune what no longer applies, and add what you have learned the hard way to specify. After two or three semesters of use, you will have a personalized configuration that reflects your work style and your discipline's particular demands. The point of starting with the elements above is to give you the structure those experienced users converge on, so you can grow your own configuration on a sound foundation rather than discovering each lesson by repeated failure.

Where to learn the platform-specific details. The instructions for setting persistent preferences vary by AI platform, and the user interfaces change often enough that any specific click-by-click guide written today will be partly outdated within a year. For current instructions, consult the documentation page for whichever platform you use; major platforms (ChatGPT, Claude, Gemini, Copilot, and others) all maintain documentation on system prompts, custom instructions, or memory configuration. If your university's library, IT services, or teaching center has a guide for academic AI use at your institution, start there; institutional guides usually take account of any university-licensed AI tools and the specific privacy considerations for student use.

26.2 Template: AI use log

Use this template for assignments that require process documentation.

Question	Your response
Assignment	
AI tool used	
Date used	
Purpose of use	
Prompt or question asked	
Useful output received	
Output rejected or changed	
Verification method	
What I learned	
What I can explain without AI	

26.3 Template: verification checklist

- ☐ I can restate the task in my own words.
- ☐ I know which parts of the work were helped by AI.
- ☐ I checked each major claim against an appropriate source or method.
- ☐ I verified all quotations and citations.
- ☐ I checked calculations, code, or data where relevant.
- ☐ I looked for hidden assumptions.
- ☐ I removed or rewrote anything I could not explain.
- ☐ I can answer questions about the final work.

26.4 Template: prompt library for learning

For understanding

Explain this idea in three levels: simple intuition, course-level explanation, and one technical detail I should not miss.

Give me an example, a non-example, and a borderline case.

Ask me questions that reveal whether I understand this concept.

For problem solving

Give me one hint about the main obstacle. Do not solve the problem.

What assumptions should I check before choosing a method?

What is a simpler related problem I could solve first?

For writing

Here is my draft. Identify the paragraph where the argument is least clear. Do not rewrite it. Ask me revision questions.

What claim in this paragraph needs stronger evidence?

Does this revision change the level of certainty in my original claim?

For coding and computation

Suggest simple test cases for this code. Do not rewrite the function yet.

What edge cases might break this algorithm?

Explain what this error message means and ask me what I have already checked.

For exam preparation

Quiz me one question at a time. Wait for my answer before giving feedback.

Ask me to explain the main concept without using jargon.

Give me a flawed answer and ask me to find the mistake.

26.5 Template: one-page student compact

I understand that AI may be allowed in some parts of this course and not allowed in others. I will follow the policy for each assignment. When AI use is allowed, I will use it to support learning rather than replace my own work. I will verify important AI-assisted content, protect private information, disclose AI use when required, and submit only work that I can explain as my own.

26.6 Short glossary

AI-assisted Practice. The middle component of the CAT framework: the student's ability to use AI well—asking useful questions, judging responses, comparing explanations, finding weak points, generating practice problems, revising carefully, testing code, checking assumptions, and using feedback without becoming dependent on it.

AI-assisted work. Work shaped in some way by an AI tool, including brainstorming, drafting, revision, coding, explanation, summarization, or feedback.

Assignment categories. Four labels for AI use in coursework: **NAI** (No AI), **AIT** (AI as Tutor, study only), **AIC** (AI as Collaborator on submitted work, with verification and disclosure), **AIS** (AI as Study, the object of analysis itself). Three sit on a traffic light (**NAI** red, **AIT** yellow, **AIC** green); **AIS** is blue, off the traffic light, because the question is different.

Calibration error. The gap between what a student thinks they know and what they can actually demonstrate. AI use widens this gap if it allows recognition to substitute for production.

CAT framework. Three dimensions of student capability under AI: Core Competence (what you can do without AI), AI-assisted Practice (skill at using AI well), Trustworthiness (verification, disclosure, ownership).

Core Competence. The first component of the CAT framework: the knowledge and skills a student must be able to demonstrate without AI or with tightly limited assistance.

Desirable difficulties. Conditions that slow initial performance (productive struggle, retrieval failure, effortful reconstruction) but produce more durable learning.

Disclosure. A clear statement of how AI was used in producing an assignment.

Distributed practice. Studying material at intervals, with forgetting and re-retrieval between sessions, rather than in a single block. Also called spacing.

Evidence of thinking. Process artifacts that show how a student reasoned, revised, verified, or made choices.

Friction-wall failure. The failure pattern produced when a student reaches for AI before making a real attempt at the problem. They recognize the AI's solution but cannot

generate the analogous one on their own, and often cannot tell the difference. Named for the moment of productive difficulty that AI use can short-circuit.

Generated citation. A source suggested by AI. Generated citations must be checked before use.

Generation effect. The finding that information you produce is remembered better than information you only read.

Illusion of competence. The mistaken sense that you understand material you have only seen explained, including the recognition–production gap, the mastery illusion on transfer, the effort illusion, and the authorship illusion (see Section 5.1).

Interleaving. Mixed practice across problem types or topics, rather than blocked practice on one type at a time. Builds the discrimination skills that single-method practice does not.

Learning Spiral. The five-step habit this guide recommends: try first, ask for help, work independently, verify, reflect and disclose. Run on every AI-assisted task. Each pass through the five steps should deepen understanding from the previous pass—which is why the framework is named a spiral rather than a loop.

Oral defense. A conversation or examination in which a student explains submitted work and answers questions about it.

Polished hollow. One of the three failure modes of the Student’s Trilemma: the student earns the grade and discloses AI use honestly but cannot reproduce the work; the surface is right and the disclosure is clean, but the underlying competence has not formed.

Practice testing. Trying to recall information, produce a solution, or reconstruct an argument, rather than re-reading or re-watching. Also called retrieval practice.

Strategic compliance. One of the three failure modes of the Student’s Trilemma: the student earns the grade and genuinely understands but hedges on how AI helped; the learning happened but the account of it is shaped to protect the grade.

Student’s Trilemma. A three-way pull (like a dilemma but with one more horn) among earning the grade, being honest about how the work happened, and demonstrating genuine understanding. Each pair pulls against the third. AI sharpens each edge but did not create the trilemma. Three named failure modes are the polished hollow, strategic compliance, and honest risk.

Translation test. A diagnostic for whether a student owns their work: can they re-present it in a different form—different notation, different variables, different example, or a different representational mode (visual, mathematical, plain language)? If they can translate fluently, the work is theirs. If not, they have recognition, not production.

Trustworthiness. The third component of the CAT framework: the student’s verification, disclosure, and ownership of AI-assisted work. A student demonstrates trustworthiness

by checking what AI produced, telling readers honestly what AI did, and standing behind the result.

Try first. The first step of the Learning Spiral: make an honest attempt before asking AI for help.

Verification. The act of checking AI-assisted content against reliable evidence, methods, sources, or reasoning. The act that drives Trustworthiness in the CAT framework.

Publication note

This manuscript is prepared as part of a three-guide package for possible publication and institutional adoption. The source is intentionally conservative in formatting: standard L^AT_EX structures, stable section labels, explicit templates, and minimal custom commands. Final production should use the current SIAM book macro package specified by SIAM Publishing at the time of submission.

Part II

Teaching and Assessing with AI — An Instructor Guide

*The question is not whether a student used AI.
It is what the work still shows.*

Part II of this book is the Instructor Guide. It helps faculty, lecturers, teaching assistants, and course coordinators design courses that remain intellectually honest, fair, and useful when students have access to powerful AI tools. It is general in scope, with especially detailed examples for STEM and research-intensive courses.

The central premise is simple. AI-assisted education must assess two things separately: what students can do independently, and what students can do responsibly with AI. A course that ignores AI may lose assessment validity. A course that accepts AI output without safeguards may lose learning. The goal is a better middle path: teach students to use AI as a tool while preserving the habits of attention, judgment, verification, and explanation that higher education is meant to build.

27 Purpose and scope

Generative AI has changed the practical conditions of teaching. Students can now obtain fluent explanations, polished essays, working code, plausible citations, study questions, summaries, translations, and step-by-step solutions in seconds. These tools can help students learn. They can also make it difficult to know what a submitted artifact reveals about the student's understanding.

This guide offers a practical framework for that problem, consistent with broader guidance that treats generative AI as a tool requiring human oversight, transparency, and institutional judgment (Miao and Holmes, 2023). It does not ask instructors to choose between banning AI everywhere and allowing it everywhere. Both responses are too blunt for the range of courses taught at a university. A writing seminar, a chemistry lab, a design studio, a calculus course, a history survey, and a graduate methods course need different rules. Still, they share the same underlying responsibility: preserve meaningful learning while teaching students to use modern tools responsibly.

The guide has four aims.

1. Help instructors state clear AI policies for each assignment.
2. Help instructors design assessments that reveal student thinking.
3. Help instructors teach verification, disclosure, and responsible judgment.
4. Help departments develop consistent practices across small, medium, and large courses.

One variable runs through the rest of this Part more than any other: class size. A thirty-student seminar, a seventy-student lecture, and a three-hundred-student introductory course are completely different worlds for AI-era assessment design. Approaches that work in the seminar do not scale to the introductory course; approaches that scale may not preserve what makes the seminar valuable. The framework developed here is built around three assessment models (Oral-Defense, Evidence-of-Thinking, and Hybrid) each calibrated to a different range of class sizes, and around three design axes that include class size as the first (section 31). The class-size axis recurs throughout the Part wherever the framework's recommendations bend with it. Readers whose class size is their dominant constraint may want to scan section 31 and the three assessment models early.

The recommendations are intentionally modest. The evidence on generative AI in education is developing quickly, and specific tools will change. Recent studies suggest that AI tutoring can support learning when designed around sound pedagogy, while unguarded access can weaken later independent performance (Bastani et al., 2025; Kestin et al., 2025). The durable principles are less likely to change: students need core competence, active practice, verification habits, transparent process, and fair assessment (Dunlosky et al., 2013; OECD, 2026).

28 Instructor outcomes

After using this guide, an instructor should be able to design a course that states AI expectations clearly and assesses student learning fairly. In practice, this means being able to do the following.

1. Identify the Core Competence that students must demonstrate without AI or with tightly bounded assistance.
2. Design AI-assisted assignments that develop AI-assisted Practice rather than reward outsourcing.
3. Assess Trustworthiness through verification, disclosure, process evidence, and explanation.
4. Choose an Oral-Defense, Evidence-of-Thinking, or Hybrid model that matches class size and staffing.
5. Use AI to support teaching without delegating grades, integrity decisions, or student evaluation to an automated system.

These outcomes keep the guide focused on course design. They also make clear that AI policy is not separate from pedagogy. It is now part of pedagogy.

You are not there to fill a bucket. Your goal is to launch a career.

28.1 Relation to international AI competency frameworks for teachers

The instructor outcomes above are scoped to course design and assessment. They sit inside the broader AI competency framework for teachers articulated by UNESCO (Miao and Cukurova, 2024), which specifies fifteen competency blocks across five aspects (Human-centred mindset, Ethics of AI, AI foundations and applications, AI pedagogy, AI for professional development) and three progression levels (Acquire, Deepen, Create). Two aspects of UNESCO's teacher framework are particularly relevant to the instructor outcomes in this Part: *AI pedagogy* (using AI to facilitate learning) and *AI for professional development* (using AI to support the instructor's own professional growth). The first is the central focus of this Part; the second is addressed briefly in the discussion of faculty use later in the chapter.

Table 3 sketches the correspondence. The mapping is partial because the instructor outcomes in this guide focus on how to design assignments and assess learning when AI tools are available, while UNESCO's framework also covers a teacher's broader AI literacy and ethical engagement with AI in society. A faculty development program that prepares instructors for the outcomes of this Part will move them through UNESCO's *Acquire* and *Deepen* levels in

Table 3: Instructor outcomes in this Part and the UNESCO AI competency framework for teachers (Miao and Cukurova, 2024): a sketch of correspondence. The mapping is partial; UNESCO covers AI literacy, ethical engagement, and professional development more broadly than this Part addresses.

Instructor outcome (this Part)	Closest UNESCO competency block
Identify Core Competence to be demonstrated without AI	AI pedagogy: AI-pedagogy integration (<i>Deepen</i>); design assignments that target core skill domains.
Design AI-assisted assignments that develop AI-assisted Practice	AI pedagogy: AI-pedagogy integration and AI-enhanced pedagogical transformation (<i>Deepen</i> and <i>Create</i>).
Assess Trustworthiness through verification, disclosure, process evidence	Ethics of AI: Safe and responsible use (<i>Deepen</i>); Human-centred mindset: Human accountability (<i>Deepen</i>).
Choose an assessment model matched to class size and staffing	AI pedagogy: AI-assisted teaching (<i>Acquire</i>) and AI-pedagogy integration (<i>Deepen</i>).
Use AI to support teaching without delegating grades or integrity decisions	AI pedagogy: AI-assisted teaching (<i>Acquire</i>); AI for professional development: AI enabling lifelong professional learning (<i>Acquire</i>).

the AI pedagogy aspect; programs aiming for the *Create* level will need to extend beyond this guide.

A program that already aligns its faculty development with UNESCO’s framework can adopt the instructor outcomes here as the assessment-focused operationalization of UNESCO’s *Acquire* and *Deepen* levels in the AI pedagogy aspect, without abandoning the broader UNESCO structure.

29 The instructor-side calibration problem

The same diagnostic I once used to catch a copied proof, by its notation, is what helps a teacher distinguish owned understanding from polished output today. The technology is different. The work of telling them apart is not.

A teacher with seventeen years of teaching experience told the *Chronicle of Higher Education* that the work of distinguishing student writing from AI-assisted writing now consumes time that previously went elsewhere. “I’m grading fake papers instead of playing with my own

kids,” he said.⁵ The line is one teacher’s, but the predicament is widely shared. The labor cost of the calibration problem is real, and acknowledging it is the place to start.

Students sense the calibration problem from an unexpected angle. Focus groups with ninety-five undergraduates at the University of Pittsburgh in 2025 surfaced a quiet finding: some students prefer AI chatbots to faculty office hours because, as one student put it, “with ChatGPT you can ask as many questions as you want and it’s not going to judge you.”⁶ The line is not a verdict on any particular instructor. It is a description of how the instructor’s classroom now looks from the other side of the desk. A patient, infinite, non-judgmental tutor has appeared, and it is competing for the role that the instructor used to hold uncontested.

The translation test introduced for students in Part I (section 6) is the same diagnostic the instructor can use. A student who cannot translate their work into different notation, with different variable names, by a different example, has produced an artifact whose origin is uncertain regardless of its surface polish. Section 40 develops the test as a teaching tool.

When students submit AI-assisted work, instructors face a parallel problem to the one Part I named for students. The cues that previously signaled mastery, including clean prose, organized argument, articulate explanation, and polished proofs, no longer reliably distinguish students who understand the material from students who used AI to produce work that resembles understanding. A student can submit a coherent essay, a well-organized lab report, or a correctly proved theorem without having done the cognitive work that produces those things in a pre-AI environment.

*AI didn’t make the problem harder to solve.
It made the problem harder to see.*

This is a calibration problem for the instructor: the difficulty of distinguishing what a finished artifact reveals about *this student*, separate from what AI may have contributed. The three assessment models presented later in this Part (Oral-Defense, Evidence-of-Thinking, and Hybrid) are designed to address it. They restore the signal that AI has weakened.

Reading the rest of Part II as a calibration-restoring project may be useful. The NAI, AIT, AIC, and AIS assignment categories are not just AI-permission categories. They are calibration choices about which dimensions of student competence each assignment measures. The verification rubrics ask students to demonstrate the gap-closing practices from Part I. The process-evidence guidance is designed to recover what the polished final product no longer reveals on its own.

Before generative AI, a submitted paper, proof, program, lab report, or problem set was never perfect evidence of learning. Students could receive help from friends, tutors, solution manuals, prior students, websites, and commercial services. AI did not create the problem of attribution. It changed the scale, speed, and polish of outside assistance.

The core assessment question is now sharper:

⁵Quoted in McMurtrie (2024).

⁶Reported in Silva (2025).

What does this submitted work show that the student can understand, explain, verify, and do?

If an assignment can be completed well by a generic AI tool, the assignment may still have value as practice, but it is weak evidence of independent competence. This does not mean the assignment should be abandoned. It means the instructor should decide what the assignment is for. Is it meant to measure unaided knowledge? Is it meant to teach responsible use of tools? Is it meant to produce a polished final artifact? Is it meant to reveal process?

Recent learning-analytics research helps name the mechanisms by which a polished AI-assisted artifact can hide reduced learning. Fan et al. (2025) describe **metacognitive laziness**: students using generative AI on revision tasks performed measurably fewer self-monitoring activities, especially evaluation and orientation, even when their final artifact scored well. Chen et al. (2025) describe a related pattern they call cognitive offloading: in their comparative analysis of student help-seeking with human experts versus AI chatbots, students working with the chatbot tended to skip the diagnostic, evaluative, and iterative stages that ordinary help-seeking includes, jumping directly from request to implementation. Both mechanisms produce the same observable outcome: a polished submitted artifact accompanied by reduced student understanding of how it was produced. Designing assignments that make those skipped stages visible is the practical response to both mechanisms, and is the orientation behind the Plan-Use-Verify and Oral-Defense patterns in this Part.

Confusion arises when a course treats one kind of task as if it measured all of these things. In the AI era, instructors should separate them.

30 The Instructor’s Trilemma

The calibration problem in the previous section names what instructors *observe*. This section names what instructors *face* when they sit down to design a course, grade a stack of papers, or decide whether to raise a quiet conversation with a student about suspected AI use. Like students and institutions, instructors under AI pressure navigate a structural predicament that has not been clearly named in the discourse. We call it the **Instructor’s Trilemma**, and it has the same three-cornered shape as the Student’s Trilemma developed in Part I and the Institution’s Trilemma developed in Part III.

Before naming the three corners, a quieter question runs underneath all of them and deserves acknowledgment. Faculty are encountering AI not only as a pedagogical challenge but as a vocational one: *what am I uniquely for, in a classroom where AI can also explain, also tutor, also assess, also produce convincing prose?* The question is rarely asked aloud, partly because it is hard to ask, and partly because no answer is fully reassuring. It runs through the faculty discourse in a hundred quiet ways: the early retirement that was not planned, the senior colleague whose enthusiasm has gone flat, the conference panels titled “the future of higher education.” The Instructor’s Trilemma does not answer this question, but it names the three forces an instructor must navigate while the question hangs in the air.

The corners are:

Professional sustainability. Keeping the job and the conditions that allow you to keep doing it. This includes manageable workload, contract renewal, tenure case, course evaluations, time to engage deeply with students, and protection against the burnout that has become widespread in the profession. AI sharpens every one of these pressures: the cheating-detection arms race adds policing time without commensurate institutional support; the volume of AI-polished submissions complicates grading; the institutional inconsistency of AI policies leaves individual instructors to invent and defend their own.

Trustworthy professional judgment. What your assessments, grades, feedback, recommendations, and other professional acts of *vouching* actually signify about the students you teach. When you assign a B to a paper, write a letter of recommendation, or pass a thesis defense, you are making a claim that other people will rely on. The credential carries weight in part because individual faculty stand behind individual judgments. When AI assistance enters student work without being verified, those judgments lose their connection to what the student can actually do.

Real student learning. Students gain understanding, skills, and judgment as a result of taking your course. This corner has two faces, and the literature shows faculty are torn between them. The first face is *cognitive development*: students should leave with strengthened reasoning, deeper understanding, and the kind of thinking that holds up when the tools they used in your course are unavailable. The second face is *workforce preparation*: students should leave equipped for professional environments where AI is a permanent feature of the work. These are not opposed in principle, but in individual course-design decisions they often pull in opposite directions. An instructor who limits AI use to protect cognitive development may be failing the workforce-preparation mandate; one who teaches AI fluency to prepare students for professional work may be enabling the overuse that 95% of faculty in a recent national survey report worrying about (Elon University and American Association of Colleges and Universities, 2026).

Each pair of corners can be satisfied at the cost of the third. The failure modes echo the Student's and Institution's Trilemmas deliberately.

30.1 Three instructor failure modes

Sustainability and trustworthy judgment without real student learning. The instructor maintains a manageable workload and grades that can be defended in conversation, but the pedagogy itself has hollowed out. Assignments become defensive: in-class blue-book exams that limit AI use but also limit what can be assessed; multiple-choice tests that are easy to grade but reveal little; assignments rewritten until blind AI use “collapses under its own weight” (Samuel, 2025) but in the process becoming so constrained that students do not develop the kinds of thinking the course title promises. We call this the **instructor polished hollow**: the professional surface is intact, the integrity record is defensible, but the educational mission has hollowed out from inside.

Sustainability and real student learning without trustworthy judgment. The instructor protects their workload and produces actual student learning, but lets AI-assisted work into the grade book without demanding genuine ownership. AI grades the homework; students use AI to write the papers; the pedagogy may even be sophisticated and engaging; but the instructor can no longer defend in conversation what individual grades, recommendations, or evaluations actually signify. We call this the **instructor strategic compliance**: the instructor complies with the institutional appearance of credentialing function while quietly accepting that the verification was not done. This is the failure mode most common when the cost of demanding genuine ownership exceeds what the instructor can sustainably do.

Trustworthy judgment and real student learning without sustainability. The instructor holds high standards, refuses to lower expectations under AI pressure, redesigns assessments to surface genuine thinking, conducts the oral conversations and process audits the framework calls for, and produces real learning that students will carry with them. And burns out. The literature names this pattern directly in its observable forms: faculty reverting to blue books or opting for early retirement to avoid the perceived flood of machine-generated text (Manturuk, 2026). We call this the **instructor honest risk**: every commitment kept, no commitment surrendered, but the personal cost cannot be sustained.

The American Association of University Professors frames this pattern at the professional-labor level. They argue that the uncritical adoption of AI poses a threat to academic professions through work intensification and through implications for the faculty working conditions that affect student learning conditions (American Association of University Professors, 2025a). The intensification is real: more policing time, more course redesign cycles, more individual policy invention. The honest-risk failure mode is the price an instructor pays for refusing the other two.

AI did not create the Instructor's Trilemma. The trilemma names a structural tension between commitments faculty have always held in some balance. What AI did was sharpen each edge. Each pair of corners is harder to hold together when AI assistance lowers the apparent cost of giving up the third.

30.2 How the trilemma manifests differently by faculty type

The three corners are universal among teaching faculty, but the pull strength on each corner differs sharply by employment status. The dimension is unusually consequential in U.S. higher education: roughly two-thirds of the academic workforce is contingent, and roughly half is part-time adjunct. The trilemma looks different depending on which contractual category the instructor occupies. Four types are worth distinguishing.

Tenured faculty with reasonable course loads. These instructors have the most freedom across all three corners. Sustainability is relatively secure (the contract cannot be cancelled at the end of the semester), trustworthy judgment is supported by long-developed disciplinary reputation, and real student learning can be pursued through whichever pedagogical experiments the instructor finds compelling. The honest-risk failure mode is the failure they can most afford

and that they sometimes consciously choose: holding high standards even at significant personal workload cost. They may also be the most likely to romanticize the honest-risk posture in advice they give to junior colleagues for whom it is genuinely unsustainable.

Tenure-track faculty. Professional sustainability dominates because the tenure case is at stake, and almost every classroom decision feeds eventually into the tenure dossier. Course evaluations matter; student complaints carry weight; pedagogical experiments that backfire can be remembered. The polished-hollow failure mode is the most relevant risk: defensive pedagogy that protects the tenure case by avoiding controversy, at the cost of actual learning. The framework provides an alternative path through better-designed assignments and clearer rubrics, but the structural pressure toward defensiveness is real and should be acknowledged.

Full-time non-tenure-track faculty (lecturers, clinical faculty, teaching-stream professors, professors of practice). These positions have grown substantially in U.S. higher education and now constitute a meaningful share of teaching capacity. The trilemma is mixed: contract renewal pressure exists but is less acute than for adjuncts; teaching effectiveness is usually the explicit evaluation criterion (which protects the real-learning corner); but workload is often heavy and infrastructure for trustworthy judgment can be thin. Strategic compliance is a real risk when teaching loads exceed what allows for genuine verification of every student's work.

Adjunct and per-course contingent faculty. For this category, professional sustainability is dominant in a way that should not be moralized. An adjunct teaching four or five sections across two or three institutions cannot complete the time-intensive verification work the trustworthy-judgment corner requires, not because of any failure of values but because the structural conditions of the role make it impossible. The strategic-compliance failure mode is the failure they most often inhabit, often without having meaningfully chosen it. **Honest risk is a luxury good in the academy:** only faculty whose positions permit it can consistently afford to take it. Recommendations in this Part that assume an instructor can spend extra time on oral defenses, individualized feedback, and process audits should be read with this constraint in mind. Institutional change that does not address the underlying labor conditions will leave contingent faculty trapped in a failure mode that is structural, not personal.

These four types are not comprehensive. Graduate-student instructors, postdoctoral teaching fellows, and emeriti returning to teach face the trilemma with their own type-specific pull patterns. Discipline also matters: faculty in the humanities and writing-intensive fields face the trilemma in its most acute form because their disciplines are where AI-generated text is hardest to distinguish from student work, while faculty in disciplines with strong oral and laboratory traditions have more natural defenses against the strategic-compliance pattern. The framework's value is that the corners stay the same; what changes is the relative pull and the failure mode most likely to materialize given an individual instructor's structural conditions.

30.3 Why the Instructor’s Trilemma matters for the rest of this Part

The assessment models, design patterns, integrity procedures, and faculty-development approaches developed in the rest of this Part are tools for holding the three corners of the trilemma together under AI pressure. The three assessment models (Oral-Defense, Evidence-of-Thinking, Hybrid) are designed to make trustworthy judgment achievable at different course sizes without requiring the honest-risk personal cost. The CAT framework’s four assignment categories are designed to align with both faces of real student learning (cognitive development and workforce preparation) rather than forcing instructors to choose. The fair response sequence for suspected misuse is designed to protect against the policing exhaustion that drives the honest-risk failure mode.

Without the trilemma frame, the operational tools in this Part can read as a list of best practices. With it, they read as coordinated commitments that protect the instructor from each of the three failure modes simultaneously. The faculty discourse on AI is currently dominated by individual instructors describing whichever failure mode they are personally most worried about. The framework’s value is that it surfaces the structure beneath the individual stories and gives instructors a vocabulary for understanding why their colleagues are choosing different responses to the same forces.

31 Three design axes: class size, level, and discipline

The design choices an instructor makes about AI in a course are governed mostly by three roughly orthogonal dimensions: class size, course level, and discipline. Each drives a different family of decisions, and treating them as a single axis (often labelled informally as “undergraduate versus graduate”) tends to bury the design logic that this Part is meant to make visible. The three axes are summarized in table 4 and discussed in turn.

Class size determines what direct contact the instructor (or trained TAs) can have with each student. A 12-person seminar can support oral defenses on every major assignment; a 200-person lecture cannot. This axis is the primary driver of which of the three assessment models discussed below fits the course. Class size does not, by itself, determine *what* students should be able to do without AI.

Course level determines the substance of Core Competence, the assignment-category mix, and how much AI autonomy a student can responsibly handle. A first-year writing course and a graduate research seminar can both use Plan-Use-Verify assignments and oral defenses, but the first-year course will lean toward NAI and AIT categories with tightly scaffolded prompts, while the graduate seminar may lean toward AIC and AIS categories with broad student latitude. The reason is straightforward: verifying AI output requires the underlying competence the AI is being used to support, and that competence is what advancement levels are designed to build.

Table 4: The three design axes for AI-aware course design. Each axis drives a different family of decisions; treating them together produces clearer choices than treating “undergraduate versus graduate” as a single axis.

Axis	Drives	Sections in this Part
Class size	<i>How</i> you assess: which assessment model is feasible.	Oral-Defense Model (section 35); Evidence-of-Thinking Model (section 36); Hybrid Model (section 37).
Course level	<i>What</i> you assess and how much AI autonomy to grant: Core Competence definition, the typical assignment-category mix, expected verification capacity.	Four assignment categories (section 33); rubrics (section 43).
Discipline	The specifics of <i>how</i> AI fits in: what verification looks like, what AI tools are appropriate, what counts as acceptable use.	Discipline overlays (section 42).

Discipline determines the specifics of verification and AI use. A mathematics assignment is verified by checking limits, units, and edge cases; a humanities essay is verified by checking primary sources, quotation accuracy, and argument structure; a programming assignment is verified by writing test cases. The framework is uniform across disciplines, but the verification practices, AI tool choices, and category boundaries are not. Discipline overlays in section 42 address these specifics.

The three axes are largely independent. A small graduate seminar and a small first-year writing class share the assessment-model options that small classes support, but differ in category mix and AI autonomy. A large introductory mathematics class and a large introductory humanities class share the model constraints that large classes impose, but differ in what verification looks like. A large advanced seminar (rare but real) and a small introductory course share neither directly, and instructors in those settings can mix elements from different parts of this Part rather than picking one prepackaged path.

This Part is organized around all three axes. The CAT framework and the four assignment categories that follow are level-agnostic and discipline-agnostic at the framework level; their operational application varies along all three axes. The three assessment models that follow that vary primarily by class size. The discipline overlays vary by discipline. Where the prose in this Part says “a graduate course” or “an advanced undergraduate course,” the relevant claim is usually about course level (and sometimes about class size, when graduate courses are typically smaller); the prose attempts to be precise about which.

32 The CAT framework

This guide uses the CAT framework: **Core Competence**, **AI-assisted Practice**, and **Trustworthiness**. The framework generalizes a three-axis assessment model originally developed for graduate mathematics, where independent mathematical understanding, AI-augmented skill, and epistemic responsibility are assessed separately. The Foreword (page 15) explains the lineage; Appendices A and B present the original graduate-mathematics formulation in technical detail for instructors who want it.

Core Competence. What students must understand and be able to do without AI or with tightly bounded help.

AI-assisted Practice. What students can do when AI is used as a permitted tool for learning, revision, analysis, coding, feedback, or exploration.

Trustworthiness. How well students verify, disclose, critique, and take ownership of AI-assisted work.

A course should assess all three dimensions, but not always with equal weight. A foundational course may emphasize Core Competence. A capstone course may place more weight on AI-assisted Practice and Trustworthiness. A professional course may emphasize judgment, privacy, and defensible decisions.

The framework is consistent with UNESCO guidance that treats AI competence as part of a human-centered educational system rather than a substitute for teacher and learner agency (Miao and Cukurova, 2024; Miao et al., 2024). It has one non-negotiable feature: AI skill should not compensate for failure to meet the core learning goals of the course. A student who can produce strong AI-assisted work but cannot explain the central ideas has not met the same standard as a student who understands the material.

32.1 Core Competence

Core Competence is the part of the course that remains the student’s own. It may include definitions, methods, concepts, interpretations, calculations, source analysis, writing decisions, experimental reasoning, coding fundamentals, design principles, or ethical judgment.

Each instructor should identify the core explicitly. The useful question is not, “Can AI do this task?” The useful question is, “What must students still be able to do when AI is unavailable, inappropriate, or untrustworthy?”

Core Competence can be assessed through no-AI quizzes, in-class writing, oral explanations, lab practicals, live coding checks, source-analysis exercises, design critiques, short exams, or closed-resource problem solving. These assessments need not dominate the course, but they should be meaningful.

32.2 AI-assisted Practice

AI-assisted Practice is the student's ability to use AI productively without losing control of the work. This includes asking useful questions, constraining a model, comparing explanations, testing code, finding flaws, improving drafts, generating examples, and using AI feedback to improve human reasoning.

This is not a lesser skill. In many fields, responsible use of AI will become part of professional competence. Students should learn it in courses where faculty can set boundaries and provide feedback.

The instructor's task is to design assignments where AI use is visible and assessable. A strong AI-integrated assignment asks students not only for a final product, but also for evidence of choices, verification, revision, and judgment.

32.3 Trustworthiness

Trustworthiness includes verification, disclosure, privacy, source discipline, and intellectual ownership. It is the dimension that prevents AI use from becoming laundering: taking generated output and presenting it as independent understanding.

A student's responsible judgment is visible when they can answer questions such as these:

- What did AI contribute?
- What did you accept, reject, or revise?
- How did you check the output?
- What sources did you verify yourself?
- What part of the submitted work can you explain without AI?

These questions are useful across disciplines. They apply to a proof, a policy memo, a lab report, a design prototype, a literature review, a computer program, and a close reading of a text.

33 Four assignment categories

Every assignment should tell students what role AI may play. The guide uses four categories: NAI, AIT, AIC, and AIS. The categories are simple enough for a syllabus, but flexible enough for different disciplines (figure 8).

Before applying these categories, instructors may want to recommend that students complete the AI account configuration template in Part I (section 26.1) at the start of the course. The configuration step takes ten minutes, calibrates the AI to the student's level and discipline,

and creates a written record of the use boundaries the student has committed to. For a course with significant AIC-category work, requiring the configuration as a no-credit first assignment is a low-cost way to start the term with shared expectations.




Figure 8: **Choosing the assignment color.** The instructor decides what the assignment is for, then selects a color that matches that purpose. **NAI** protects independent competence; **AIT** allows AI for study support but not for submitted work; **AIC** permits responsible AI use in producing the submission; **AIS** makes AI itself the subject of analysis. The choice should appear on the assignment.

33.1 **NAI: No AI**

NAI assignments prohibit AI use. They are used to assess Core Competence.

Appropriate NAI assignments include in-class exams, short quizzes, oral explanations, live coding tasks, lab practicals, handwritten derivations, in-class writing, language fluency checks, and baseline skill demonstrations.

A NAI assignment should have a clear reason. It should not be a default expression of distrust. The reason is usually that the instructor needs evidence of what the student can do without outside help.

33.2 **AIT: AI as Tutor**

AIT assignments allow AI for study support, but not for producing submitted work. Students may use AI to ask for hints, examples, explanations, practice questions, or feedback before doing the work themselves.

AIT rules are useful when the instructor wants students to benefit from tutoring while still submitting independent work. A homework set might be AIT if students may ask AI to explain a concept but may not ask it to solve the assigned problems.

For AIT assignments, instructors may define a help ladder that structures the order in which AI assistance is permitted: concept explanation, then an analogous example, then one hint on the assigned problem, then a critique of the student’s attempt, and only then a full solution for ungraded practice. The ladder allows substantial AI support while protecting retrieval practice and productive struggle. Students who begin with a full solution to the assigned problem skip every learning move the assignment was designed to produce.

33.3 AIC: AI as Collaborator

AIC assignments allow AI as part of the working process. Students may use AI to brainstorm, revise, debug, critique, summarize, generate alternatives, or test ideas. AIC assignments require disclosure and verification.

AIC assignments work best when the process matters. For example, a student might submit an original plan, selected AI interactions, a revised product, and a verification note. The final product may be stronger because AI was used, but the student’s judgment remains visible.

33.4 AIS: AI as Object of Study

AIS assignments make AI itself part of the subject. Students might critique an AI-generated essay, identify hallucinated citations, repair flawed code, compare model responses, evaluate bias, or analyze how a tool handles a disciplinary problem.

AIS assignments are especially valuable because they teach students not to treat AI as an oracle. They make error detection and tool critique explicit learning goals.

33.5 A syllabus table for assignment categories

Category	AI role	Typical purpose
NAI	No AI	Measure independent competence.
AIT	Tutor only	Support study while preserving independent submitted work.
AIC	Collaborator	Teach responsible AI-assisted work with disclosure and verification.
AIS	Object of study	Teach critique of AI output, bias, errors, and limits.

33.6 Category mix shifts with course level

The four categories are level-agnostic at the framework level; any assignment in any course can fall into any of the four. The *typical mix*, however, shifts substantially with course level, and ignoring this shift is one of the more common ways instructors design AI policy poorly.

Introductory courses typically lean toward NAI and AIT. The reason is not distrust. The reason is that students at this level have not yet built the disciplinary competence required to verify AI output reliably; granting AIC-category latitude before that competence exists can amount to inviting students to accept work they cannot evaluate. A Calculus 1 student cannot distinguish a correct optimization solution from a confidently wrong one; a first-year writing student cannot reliably detect when an AI revision has shifted the meaning of a thesis. **NAI** protects the foundational skill; **AIT** lets the AI act as a study companion without entering the submitted work.

Intermediate undergraduate courses typically use a balanced mix. As students develop disciplinary competence, the verification side of AIC-category work becomes feasible, and AIS-category critique becomes possible. A junior-level course in any discipline can typically support meaningful AIC and AIS assignments alongside NAI and AIT ones.

Advanced undergraduate and graduate courses typically lean toward AIC and AIS. Students at this level can verify AI output, can defend submitted work in detail, and can engage AI tools as objects of disciplinary critique. NAI assignments still serve a purpose at this level (foundational-skills demonstrations, qualifying-exam preparation), but they are typically a smaller fraction of the course.

The shift is not strictly linear with level—a research-methods graduate seminar might sensibly require several NAI oral examinations, while an upper-division undergraduate elective might be almost entirely AIC. The shift is also not driven by class size: a 200-student introductory lecture and a 12-student introductory seminar share the typical NAI/AIT weighting; a small graduate seminar and a large graduate-equivalent course share the typical AIC/AIS weighting. Course level, not class size, is what shifts the typical mix.

A useful rule for an instructor designing a course: the Core Competence (NAI and AIT assessments) defines what the course is for; the AI-assisted Practice (AIC) and AI-as-object (AIS) work defines how that competence is exercised in a tools-rich environment. The proportion between these two halves is what shifts with course level.

34 Designing assignments that preserve thinking

The best AI-aware assignments make student thinking visible. They do not merely ask for a final answer. They ask for evidence of how the student framed the problem, what they tried first, what they asked AI, what they checked, and how the final work changed.

The pedagogical foundations for this design space are well established. A randomized

controlled trial of an AI-tutoring system in an introductory physics course at Harvard identified seven research-based practices that distinguish AI tools that support learning from AI tools that merely produce answers: facilitating active engagement, managing cognitive load, promoting a growth mindset, scaffolding content, ensuring accuracy of information and feedback, delivering targeted and timely feedback, and allowing for self-pacing (Kestin et al., 2025). The seven practices map onto principles long established in the learning-science literature (Ambrose et al., 2010; National Academies of Sciences, Engineering, and Medicine, 2018). Active engagement in particular rests on a large empirical base in undergraduate STEM: a meta-analysis of 225 studies found that active-learning approaches raised examination performance by 0.47 standard deviations, with students under traditional lecturing 1.5 times more likely to fail than those under active learning (Freeman et al., 2014). The cognitive-load principle is developed in detail in the synthesis by Sweller et al. (2011). The first five practices can be built into any assignment whether or not it uses AI; the last two are easier when AI is part of the design. In Kestin’s specific implementation (described in more detail in Part I section 24), the AI tutor’s system prompt handled the first three practices, the platform’s structured walkthrough handled scaffolding, and the remaining three are properties of AI as a medium; the assignment patterns recommended in this section are consistent with all seven, and instructors who want a deeper grounding for these patterns can read the Kestin study directly. A separate field experiment in high-school mathematics provides the converse evidence: an unrestricted AI interface that simply answered questions produced 48% practice gains but 17% exam losses, while a guardrailed variant designed to give hints rather than answers produced 127% practice gains and no exam loss (Bastani et al., 2025). The design choice that drove the difference was exactly what the AIT assignment category in this guide formalizes: AI is permitted as a hint-giver, not as a solver.

A separate concern is the social dimension of learning. A study of generative AI’s effect on peer interactions in computing courses found that students increasingly route their questions to AI rather than to each other, and that students seeking peer help are sometimes redirected to AI rather than receiving it (Hou et al., 2025). The pattern erodes the study groups, peer tutoring, and informal collaborative practices through which much undergraduate learning has historically happened, including the peer-instruction methods documented to improve conceptual understanding in undergraduate STEM (Mazur, 1997). Assignments designed to preserve thinking should preserve named peer interaction as one of the things they preserve.

A useful design pattern is the **plan–use–verify** structure (figure 9). The structure operationalizes the principle, well established in the learning sciences, that learners benefit from explicit self-monitoring before and after task engagement (National Academies of Sciences, Engineering, and Medicine, 2018).

Plan. Before using AI, the student writes a short plan, prediction, outline, proof sketch, lab strategy, or source note.

Use. The student uses AI within the assignment rules and records the relevant interaction.

Verify. The student checks the output and explains what was accepted, changed, rejected, or corrected.




Figure 9: **The Plan–Use–Verify assignment structure.** Each stage produces a small artifact the instructor can see. The three artifacts together show how the student thought, where AI helped, and what the student verified. Grading does not require reading every prompt; it requires seeing that all three stages happened.

This structure does not require grading every prompt. It asks students to show enough of the process that the instructor can see whether learning occurred.

34.1 Plan-Use-Verify in practice

The three artifacts work because each is short and gradable at a glance. The pattern is the same shape across disciplines, with only the lengths and content varying.

For a humanities essay, the Plan might be a one-paragraph thesis with the primary sources the student will use. The Use is a 3–5 line tool log naming what the student asked AI and what was returned in summary. The Verify is a 4–6 line note describing one place the student returned to a source to check an AI-suggested claim, plus what was accepted or rejected.

For a programming problem set, the Plan is a sketch of the algorithm or data structure the student intends to use, written before any code. The Use is the prompts and key responses from any AI debugging help, with the student’s own diagnostic questions surfaced. The Verify is a list of test cases run, with one or two showing edge cases that catch a typical bug.

For a mathematics problem, the Plan is the student’s setup of the problem (variables, constraints, choice of method) before AI use. The Use is a record of any AI hint or critique. The Verify is a unit check, a limiting-case check, or a comparison with a known result.

Three patterns make this work and three patterns break it.

Works when:

- The three artifacts together are short — one page or less for most assignments. Brevity forces selection of the meaningful moves.
- The artifacts are gradable at a glance, not parsed line-by-line. The instructor reads them as evidence of process, not as objects of grading.
- The student treats the structure as a thinking aid. The Plan forces commitment before

AI exposure; the Verify forces a check after.

Breaks when:

- The artifacts are graded on length or completeness. Process documentation drifts into theater (section 36.1).
- The Plan is written after the work is finished. The pattern collapses into retroactive rationalization.
- The instructor demands verbatim AI transcripts. The signal-to-noise ratio drops to near zero, and students learn to game the documentation rather than the work.

The pattern scales: same shape for a 1-page homework problem and for a semester project; only the lengths change. An instructor adopting Plan-Use-Verify across a course can keep the structural template constant and adjust expected lengths to the assignment.

34.2 Design features that improve assessment validity

Assignments are stronger when they include one or more of the following features. The pattern these features share is the one articulated by the backward-design approach to assessment (Wiggins and McTighe, 2005): working from the learning outcome to the evidence that would demonstrate it, then to the assignment that produces such evidence, rather than from familiar assignment forms outward.

- A first attempt before AI use.
- A short explanation of choices made during revision.
- A verification checklist.
- A source or evidence audit.
- A transfer question completed without AI.
- A short oral, written, or video explanation.
- A flawed AI response that students must critique.
- A comparison between the student's approach and an AI-generated approach.

These features move assessment away from polished output alone and toward reasoning, judgment, and process. The distinction between polished output and demonstrated learning (figure 10) is the conceptual heart of AI-aware assessment. This is the practical face of the long-established *learning versus performance* distinction in cognitive psychology (Soderstrom and Bjork, 2015; Bjork and Bjork, 2011): conditions that improve immediate performance can fail to produce durable learning, and the converse.




Figure 10: **Polished output vs. demonstrated learning.** Final polish and underlying competence are independent. The upper-left quadrant is the assessment-validity problem: high polish without learning. Strong AI-aware course design moves student work toward the right two quadrants and gives instructors evidence to distinguish all four cases.

34.3 Design features that weaken assessment validity

Some common assignment designs are now fragile.

- A take-home task graded only on final polish.
- A generic prompt that can be answered well by a public chatbot.
- A writing assignment with no draft history or author reflection.
- A coding assignment graded only on whether it runs.
- A problem set where students submit only final answers.
- A citation-heavy assignment with no source verification requirement.

These assignments may still be useful for practice, but they should not carry too much weight as evidence of independent understanding.

Table 5: **Three assessment models compared.** The Oral-Defense Model relies on live questioning to confirm understanding; the Evidence-of-Thinking Model relies on process artifacts that are difficult to outsource; the Hybrid Model combines both. The right model is the one that fits the course size, staffing, and goals.

	Oral-Defense Model	Evidence-of-Thinking Model	Hybrid Model
Best for	Small classes; seminars and honors courses; graduate courses	Large lectures; service courses; online courses	Medium courses; lab and project work; multi-TA staffing
Evidence sources	Oral conferences; live questioning; submitted work; in-class observation	Process artifacts; in-class quizzes; random audits; submitted work	Mix of oral and process; spot oral checks; process artifacts; in-class checks

35 The Oral-Defense Model

The Oral-Defense Model is best suited to small classes of any level: seminars, honors courses, writing workshops, independent studies, capstones, advanced labs, and many graduate courses. The model is class-size-driven; it works wherever the instructor (or trained TAs) can speak directly with each student about their submitted work. The assignment-category mix and the AI autonomy granted to students will shift with course level (section 33.6), but the assessment model itself is the same. Table 5 compares the three assessment models in this Part.

The model rests on a simple principle:

Students may use AI more extensively when they can personally defend the work.

An oral defense does not need to be adversarial. It can be a teaching conversation. Its purpose is to test whether the student’s submitted work has become part of the student’s understanding.

35.1 Recommended structure

A short oral defense can take 15 to 20 minutes. A fuller version can take 30 to 60 minutes. The following structure works across disciplines.

Segment		Purpose
Student summary	sum-	The student explains the submitted work in their own words.
Core question		The instructor probes a central concept, method, source, or decision.
Audit-and-repair		The student critiques a flawed AI response or a flawed piece of their own process.
Transfer question	ques-	The student applies the idea to a new but nearby situation.
Process defense		The student explains how AI was used, checked, and limited.

35.2 Examples across fields

In mathematics, the student may prove a result at the board, identify which hypotheses matter, and repair a flawed proof. In biology, the student may explain the design of an experiment and interpret a surprising result. In history, the student may defend the use of a primary source and distinguish evidence from background summary. In writing, the student may explain why a paragraph was revised and how the claim changed. In computer science, the student may walk through code and predict what happens on an edge case.

The format changes by discipline. The purpose remains the same: the student must show ownership.

35.3 Equity in oral assessment

Some students hear “oral audit” as threatening, even when the instructor intends a teaching conversation. A few small adjustments protect fairness without weakening the assessment. Share sample audit questions in advance so students can prepare. Grade for understanding, not verbal polish. Allow students to point to their work, write a formula, sketch a diagram, or explain in steps; oral does not have to mean spoken-only. The goal is ownership, not performance under pressure. These adjustments help multilingual students, shy students, neurodivergent students, and students who understand the material but need a moment to organize their explanation. The Oral-Defense Model is a calibration-restoring design; the calibration is of student understanding, not of presentation skill.

Two further considerations matter. First, students with documented accommodations through the institution’s disability services—including students with anxiety disorders, students who use augmentative communication, students who are deaf or hard of hearing—may need formats beyond the spoken-aloud audit. A short written explanation, a signed conversation, a typed exchange, or a recorded video with subtitles can satisfy the diagnostic purpose while accommodating the student. The framework is about ownership of work, not about fluency in any particular mode of demonstration. Second, instructors who can do so should

make alternative formats available to any student on request, not only those with formal accommodations on file. Some students whose understanding is sound will struggle in live questioning for reasons that do not appear on any accommodation letter; quiet flexibility about format protects the model's validity.

35.4 Suggested weighting for small classes

Component	Typical weight
Core competence assessments	30–40%
AI-integrated assignments or projects	25–35%
Oral defense or conference	20–30%
Disclosure, verification, and reflection	10–15%

These weights are only a starting point. A writing workshop may place more weight on conferences and revision. A lab course may emphasize practical demonstration. A proof-based course may require a higher independent competence threshold.

36 The Evidence-of-Thinking Model

The Evidence-of-Thinking Model is designed for large classes where one-on-one oral defense is not feasible for every student: large lectures, service courses, multi-section courses, general education courses, and online courses. Most large classes happen to be undergraduate simply because that is where enrollment concentrates, but the model is class-size-driven, not level-driven; a large advanced course or a large graduate-equivalent course uses the same evidence-of-thinking patterns, with the assignment-category mix shifting with course level (section 33.6).

The model rests on a different principle:

When individual oral examination is not feasible, the course should collect repeated, low-cost evidence of student thinking.

This model does not assume dishonesty. It assumes that polished products are no longer enough. In a large course, instructors need scalable ways to see whether students understand what they submit.

36.1 Scalable evidence sources

Useful evidence sources include short no-AI quizzes, homework-to-quiz transfer questions, in-class writing, process logs, source-verification sheets, draft histories, lab notebooks, code

test reports, AI-output critiques, video explanations, peer review, and random oral audits.

The goal is to distribute accountability across the term. A course with ten small checks is often more robust than a course with two high-stakes take-home assignments.

Process evidence should be short and targeted. Long AI transcripts and lengthy prompt logs should not be rewarded for length, since process documentation can drift into theater. The useful evidence is a decision, a verification step, a correction, or a change in understanding. Evidence-of-Thinking is not a paperwork model; it is a calibration model that reads small artifacts for signal.

36.2 Suggested weighting for large classes

Component	Typical weight
Short no-AI checks	25–35%
AI-assisted homework or projects	20–30%
Process evidence and verification notes	15–20%
AI critique or error-detection tasks	10–20%
Random oral, video, or TA-led checks	5–10%
Reflection and disclosure	5–10%

A large course does not need to use every component. The important point is balance. Students should have chances to practice with AI, but the course should also include direct evidence of independent understanding.

36.3 Random oral audits

A random oral audit is a short conversation about submitted work. It may last five to eight minutes. Students know in advance that any AI-assisted assignment may be selected.

A good audit asks ordinary questions.

- Explain the main idea of your submission.
- What did AI help with?
- What did you change after checking the output?
- Show me one place where you verified a claim.
- What would change if this assumption were different?

The audit should be framed as part of the course design, not as an accusation. It protects fairness by holding everyone to the same expectation: if students submit work, they should understand it.

36.4 TA-supported mini-conferences

In multi-section courses, trained teaching assistants can conduct short structured conferences. The instructor should provide a common script and rubric. TAs should grade a small sample together before conducting conferences independently.

This structure is most useful when assignments are AIC or AIS. It helps ensure that AI-assisted work remains connected to student understanding.

37 The Hybrid Model

Many courses fall between the small-class and large-class models. A class may be too large for full oral exams but small enough for some direct contact. A lab course may have many students but strong TA support. A writing course may have conferences for major projects but not for every assignment. The Hybrid Model is the realistic choice for most courses; it is the third of the three calibration-restoring designs presented in this Part. Process artifacts carry most of the load, while sampled live conversation provides direct verification at the rate the staffing allows.

Useful hybrid tools include group oral exams, sampled oral audits, short student conferences, video explanations, peer review with instructor spot-checking, TA-led recitations, and in-class transfer questions. The pattern is to use live explanation where it matters most and process evidence everywhere else. The remainder of this section develops the model with the same four-element treatment used for the Oral-Defense and Evidence-of-Thinking models: recommended structure, sampling strategy for live elements, a worked example, and suggested weighting.

37.1 Recommended structure

A hybrid course pairs one or more live elements with continuous process evidence. The pairing depends on the staffing and the course type. Table 6 sketches the mapping for four common course types.

The mapping is illustrative; many other pairings work. The instructor's task is to identify the live element the staffing supports, then design process evidence to cover the rest of the term. A common error is to over-design both elements, producing a course that asks students to document everything and defend everything, which collapses into paperwork. The hybrid

Table 6: Common hybrid pairings of live element and process element by course type. The live element provides direct verification of understanding; the process element distributes accountability across the term.

Course type	Live element	Process element
Lab course (intermediate enrollment, strong TA support)	Lab-time oral defense of submitted code or analysis	Lab notebook, code tests, in-class transfer questions
Medium lecture (40–80) with TAs	Sampled TA-led conferences once per term per student	Process logs, weekly transfer quizzes
Writing course with conferences (20–40)	Mid-term and final 20-minute conferences on major projects	Draft histories, author’s notes, in-class writing
Flipped-classroom STEM (any size)	In-class problem-solving with instructor circulation	AI-tutor logs from pre-class material, post-class transfer quizzes

model works when each element does its part: the live element verifies depth on a sampled basis; the process element provides continuous low-cost evidence.

37.2 Sampling strategy for live audits

When a course cannot conference every student on every assignment, the instructor must choose how to sample. Three sampling strategies are common, with different trade-offs.

Random sampling. Every student has the same probability of being audited on every assignment. Fairest in expectation, simplest to administer, hardest to argue with. Works well when the course has enough audits per term that most students will be sampled at least once. The drawback is statistical: a student with weak understanding may not be sampled on the assignment where the gap is largest.

Stratified sampling. The instructor stratifies the population (by submission quality, by stated AI use, by prior audit history, or by self-selection through optional sign-up) and samples within strata. Concentrates effort where it adds most assessment information. The risk is the appearance of targeting; the strata should be public and the rationale should be educational rather than punitive.

Threshold sampling. The instructor audits any submission that meets a threshold condition (a confident-sounding answer to a problem the student has not worked on in class; a major project crossing a stated complexity boundary; a flagged inconsistency between submitted work and prior work). Most efficient when the thresholds are known in advance. Vulnerable

to adverse-selection effects: students who learn the thresholds may submit work that just stays under them.

A hybrid course can mix strategies across the term. For example, a course might use random audits early (when the threshold for problematic submissions is still being calibrated) and stratified or threshold audits later (when the instructor has seen enough of each student's work to recognize patterns).

37.3 Worked example: a hybrid course across the semester

Consider a 60-student intermediate course with two TAs. Class meets twice weekly for 75 minutes; one TA-led recitation per week. The instructor designs the course as follows.

Week 1. Students complete the AI account configuration template (section 26.1) as a no-credit first assignment. The instructor explains the course's category mix (mostly AIT and AIC; one NAI mid-term; one AIS critique exercise late in the term).

Weeks 2–4. AIT and AIC homework assignments, each requiring a short Plan-Use-Verify artifact (one short paragraph each for plan, use, and verify). Weekly 5-minute in-class no-AI transfer quizzes on the prior week's material.

Week 5. Each student completes a 5-minute calibration conference with their assigned TA. The conversation is structured: the student summarizes one recent assignment, answers one transfer question, and describes their AI use pattern. The TA records a brief note. This early conference is for calibration, not grading.

Weeks 6–9. Continued homework + transfer quizzes. One AIS-category critique assignment around Week 8 (students critique an AI-generated answer to a course-relevant problem and identify two errors). Process logs are required for every assignment; the instructor reads only the verification notes, sampling the rest.

Week 10. Mid-term exam: paper-and-pencil, no-AI, weighted as the NAI component of the course.

Weeks 11–14. Larger AIC-category project. Students submit a project proposal at Week 11; receive AI-assisted draft feedback at Week 13; submit a final draft at Week 15. Each student's project culminates in a 10-minute oral defense during the final week.

Week 15. Project oral defenses (in-class, with the instructor present and TAs as second readers).

Final. Cumulative process review. The instructor scores process artifacts across the term, weighted as a single component.

The course distributes assessment across the term: weekly small no-AI checks, two oral conferences (early calibration + final defense), one paper-and-pencil mid-term, one AIS critique, one Plan-Use-Verify pattern on every AIC assignment. No single assessment carries excessive weight; no single modality is the only path to demonstrating learning.

37.4 Suggested weighting for hybrid classes

Component	Typical weight
Core competence assessments (no-AI quizzes, mid-term, NAI components)	25–35%
AI-integrated assignments or projects (AIC)	25–35%
Process evidence and verification notes (Plan-Use-Verify artifacts)	15–20%
Sampled oral, video, or conference checks	10–15%
AI critique tasks (AIS)	5–10%
Reflection and disclosure	5–10%

These weights are starting points. A lab course will weight the live lab-time defense higher; a writing workshop will weight conferences and revision higher; a foundational course will weight no-AI assessments higher. The principle is that no single modality should carry the whole assessment load: a student who is weak on one modality but strong on others should be visible to the instructor as the intermediate case they are, not as a failing student.

37.5 Two empirical anchors

Two recent studies provide empirical grounding for the patterns above. Kestin et al. (2025) showed in their RCT that an AI tutor designed around the seven pedagogical practices listed in section 34 could outperform in-class active learning on the introductory material it covered, when measured by post-test learning gains. The implication for course design is not that AI should replace classroom instruction. The implication is that, when AI tutoring is well-designed, it can effectively bring all students up to a level where in-person class time can be spent on advanced problem solving, project work, group discussion, and the kinds of higher-order tasks that benefit most from a faculty member’s presence. Used this way, AI tutoring functions as a flipped-classroom mechanism, with the AI handling first exposure and the instructor handling deepening, transfer, and oral assessment.

Kosar et al. (2024) report a more direct precedent for the Oral-Defense and Evidence-of-Thinking elements that hybrid courses combine. They redesigned a programming course for novice programmers by changing home assignments so they could not be directly completed by an AI chatbot, using lab time for oral defenses of submitted code with “understanding” questions, and converting the mid-term exam to paper and pencil with a focus on conceptual understanding. The redesign produced similar learning outcomes for student groups with and without GenAI access. The redesign described in that study is operationally close to the hybrid pattern recommended here. Their finding—that the with-AI and without-AI cohorts converged when the assessment design carried the load—is the most direct empirical evidence available that the assessment models in this guide work in practice.

38 Feedback and grading with AI

Instructors may also use AI. This can be legitimate and useful, but it requires boundaries.

Feedback design is one of the most-studied levers in the learning-science literature, building on the foundational case for formative assessment made by Black and Wiliam (1998). Effective feedback addresses three questions for the learner: where am I going, how am I going, and where to next; and it works best when directed at task and process rather than at the person (Hattie and Timperley, 2007).

AI may help instructors draft practice questions, generate examples, identify possible misconceptions, rewrite assignment instructions, produce alternate explanations, create low-stakes quizzes, and draft preliminary feedback. These uses can save time and improve student support.

But final evaluative judgment remains the instructor’s responsibility. AI should not assign final grades. It should not decide whether misconduct occurred. It should not be used in ways that violate student privacy or institutional data rules.

A good instructor-facing standard mirrors the student standard:

AI may assist judgment, but it must not replace judgment.

38.1 Safe uses of AI for feedback

Relatively safe uses include asking AI to suggest clearer wording for feedback already written by the instructor, generate additional practice questions for common errors, produce alternative explanations of a concept, or check whether a rubric covers the intended learning outcomes.

Even here, the instructor should review the output. AI-generated feedback can sound confident while missing the actual student difficulty. It can also become too generic to help.

When using AI to draft feedback on student errors, prompt the AI for scaffolded hints rather than complete corrections. An instructor who uses AI to fully rewrite a student’s flawed paragraph and hands it back has simply outsourced the cognitive offloading; the desirable difficulty (Bjork and Bjork, 2011) that produces learning is lost. Hints, leading questions, or partial models preserve the work the student still needs to do, which is the work that produces the learning.

A constructive recent example is the *Tutor CoPilot* system developed at Stanford, an AI tool finetuned on ethnographic observations of effective tutoring and integrated into an online tutoring platform. In a randomized study with 900 tutors working with 1,800 underserved students in the United States, the tool raised student pass rates by an average of 4 percentage points, with the largest gains (about 9 percentage points) for less-experienced tutors and smaller gains for already-experienced ones (Wang et al., 2024); the study is currently circulating as an Annenberg working paper pending journal publication, and the

specific effect sizes should be read with that status in mind. The pattern is informative even at preliminary status: AI assistance supporting human teachers produced measurable benefits without replacing the human teaching role. For instructors thinking about how AI might support their own feedback workflow, this study illustrates that the most useful design point is often AI that scaffolds the less-experienced rather than AI that replaces the experienced.

38.2 High-risk uses of AI for grading

High-risk uses include uploading identifiable student work to unapproved tools, using AI to assign scores without human review, using AI-generated accusations of misconduct, or using a model's explanation as the sole basis for a grade.

These practices can harm students and weaken trust. They should not be adopted without explicit institutional approval, privacy review, and human oversight.

39 Annotated dialogues: weak AI use as teaching cases

The following two dialogues show what weak AI use looks like and how an instructor might respond. They are paired with the strong-use dialogues in section 11 so that an instructor can show students both versions side by side. The contrasts are pedagogically the most useful part: the boundary between strong and weak use is rarely about whether AI was used, and almost always about how the student engaged with the work.

Dialogue 7: Weak undergraduate use (calculus, contrast with Dialogue 1)

Student. Solve this calculus optimization problem: a farmer wants to fence a rectangular field of 400 m^2 along a river, fencing only three sides. What dimensions minimize fencing cost?

→ Mistake: asked AI to solve, not for a hint.

AI. *Let length parallel to river be L and width be W . Constraint: $LW = 400$, so $L = 400/W$. Fencing cost is proportional to $f(W) = 400/W + 2W$. Setting $f'(W) = -400/W^2 + 2 = 0$ gives $W = \sqrt{200}$, so $L = 2\sqrt{200}$. Second derivative $f''(W) = 800/W^3 > 0$ confirms a minimum.*

Student. Thanks. I'll submit this.

→ Mistake: did no independent work, no verification, no reconstruction.

Student. [At the random oral audit a week later.] Instructor: "Why did you set up the function as $400/W + 2W$ rather than $L + 2W$?" Student: "I think because... I'm not sure. The AI did it that way."

→ The audit reveals what the submission concealed: no understanding of the constraint substitution that is the key insight. Course rule violated: "If you cannot explain it, it should not be in your work."

Student. Instructor follow-up: “What would change if the area constraint were 200 m² instead of 400?” Student: “I’d have to redo the problem with a new tool query.”

→ The student cannot transfer. This is the canonical failure mode the Learning Spiral (figure 6) is designed to prevent.

Instructor scoring. This case illustrates the upper-left quadrant of figure 10: high polish, low demonstrated learning. A reasonable response under most course policies is partial credit on the submission (the answer is correct), zero credit on the oral audit, and a teaching conversation about how to use AI in this course. The same submission paired with the engagement shown in Dialogue 1 would receive full credit on both.

Dialogue 8: Weak graduate use (proof, contrast with Dialogue 2)

Student. Prove that every continuous function on $[0, 1]$ attains its supremum.

→ Mistake: asked AI to produce the proof rather than to audit one.

AI. Let $M = \sup_{x \in [0, 1]} f(x)$. Since $[0, 1]$ is compact and f is continuous, $f([0, 1])$ is compact in $\mathbb{R}$, hence closed and bounded. Therefore M is attained.

Student. Looks fine. I’ll submit this.

→ Mistake: accepted a proof that uses a result (compactness implies attainment) that is essentially the conclusion. The AI proof is technically correct, but it bypasses the constructive Bolzano–Weierstrass argument the qualifying-exam syllabus expects.

Student. [Oral examination.] Examiner: “Walk me through the step where you concluded M is attained from compactness of $f([0, 1])$.” Student: “A compact subset of $\mathbb{R}$ is closed and bounded, so the supremum is attained.”

Student. Examiner: “Right. Now: prove that statement directly. Why is the supremum of a closed bounded subset of $\mathbb{R}$ attained?” Student: “...”

→ The student passed the surface but not the depth. GMIP D2 (proof construction *and repair*; see Appendix A) requires producing the underlying argument when probed, not just citing the named theorem.

Instructor scoring. The submission is correct mathematics. Under conventional grading it could pass. Under GMIP, the oral component reveals that the student has not built the underlying argument. Compare with Dialogue 2: the same student, given the same prompt, can engage with subtle critique, push back when AI is wrong, and produce a constructive proof. The framework distinguishes these two cases; the framework’s job is to make that distinction visible.

40 Academic integrity without detector dependence

AI detectors should not be the foundation of an academic integrity system. They produce both false positives and false negatives, and the magnitudes are not small. In a study evaluating seven widely used GPT detectors on TOEFL essays written by non-native English speakers, the average false-positive rate was 61.3%, all seven detectors unanimously misclassified 19.8% of the essays, and at least one detector flagged 97.8% of them as AI-generated (Liang et al., 2023). The mechanism is identifiable: most current detectors rely on text perplexity, a measure of how predictable a sequence of words is to a language model. Non-native writers exhibit lower lexical and syntactic variability, which produces lower perplexity that detectors then read as AI-generated. The same study found that detectors can be bypassed by prompting an AI to use more literary language: detection rates fell to near zero. Detectors therefore fail in two directions, and both directions disadvantage students who cannot afford to be on the wrong side of either error. The specific rates above were measured on the detector tools available in 2023; subsequent detector generations may reduce the false-positive rate, but the structural concern (that perplexity-based detection penalizes non-native writers) is intrinsic to the detection method and is unlikely to disappear without a fundamental change in how detection works.

Six in ten non-native English essays were flagged as AI-generated. The detector was wrong. —ed.

A citation is a claim, not a decoration. Verify it.

The institutional consequences extend beyond accuracy. Detector-based enforcement creates what the same authors term a presumption of guilt: students are flagged as suspect by an automated system, and the burden falls on them to prove innocence. Even when an accusation is later revoked, the student’s reputation and trust in the institution may not recover. The teaching-and-learning approach to academic integrity, developed long before AI tools were a concern, treats integrity as a pedagogical issue rather than a policing one (Bertram Gallant, 2017). The same orientation applies to AI-related concerns: when integrity is built into course design, fewer concerns arise, and those that do can be handled through evidence and conversation. When concerns arise, instructors should use multiple sources of evidence.

- Assignment policy and category.
- Student disclosure or lack of disclosure.
- Draft history or version record.
- Source verification.

- In-class performance.
- Oral explanation or targeted reassessment.
- Consistency with prior work, interpreted cautiously.
- A fair conversation with the student.

The goal is not to avoid accountability. The goal is to make accountability fair, educational, and evidence-based.

Spot-checking does not require reviewing every student's AI conversation on every assignment. Reviewing three or four threads per assignment, chosen at random or based on signal, creates probabilistic deterrence at sustainable cost. Students who know that any single thread might be reviewed engage differently from students who know nothing is ever checked. The probabilistic version preserves the framework's three-cornered balance: the instructor's trustworthy-judgment corner is honored without making sustainability impossible, and the student-side disclosure practice in section 14.4 becomes operationally useful for the instructor without requiring the workload of exhaustive review.

Faculty are not exempt from the integrity question. A senior at Northeastern University filed a formal complaint after she found her organizational behavior professor was using AI to produce his lecture materials, despite a syllabus that prohibited student AI use.⁷ The asymmetry is the issue: an integrity rule that applies to students but not to those teaching them is structurally fragile, and students notice.

The framework developed in Part I for student AI use applies symmetrically to instructor AI use. The Instructor's Trilemma developed in section 30 names the three forces faculty navigate under AI pressure: professional sustainability, trustworthy professional judgment, and real student learning. The Northeastern case can be read through that frame as a failure to hold trustworthy judgment together with the other two corners. The lecture materials presented to students implied a level of instructor engagement and expertise that the AI-assisted production did not actually back up; the syllabus claimed an integrity standard the instructor did not apply to himself.

The translation test gives both sides of the classroom a positive standard. A student who can translate their work into a different form owns it. A faculty member who can translate AI-assisted course material into their own language, with their own examples, in their own structure has integrated it. A faculty member who cannot has the analog of the polished hollow: course material the instructor cannot stand behind in conversation. The test is not punitive. It is a private check faculty can run on their own work before placing it in front of students, and a standard the institution can defend in public.

The translation test catches the *polished hollow* mode of the Student's Trilemma (section 5) directly, because the student cannot translate. The other two modes, *strategic compliance* and *honest risk*, require different instructor responses, developed in the rest of this Part. The parallel question for instructor misuse—what response sequence applies when faculty use

⁷Hill (2025); the case is discussed in full in Part III section 51.

AI without disclosure or in ways that contradict their own policies—is taken up in Part III section 51.

40.1 A fair response sequence

When AI misuse is suspected, a fair process can follow five steps.

1. Review the assignment policy and the student’s submitted disclosure.
2. Identify the specific concern, not merely a general suspicion.
3. Invite the student to explain the work and the process.
4. Use a short reassessment or oral explanation when appropriate.
5. Distinguish misunderstanding, poor documentation, and deliberate misconduct.

For minor first problems, an educational remedy may be more appropriate than a formal charge. Serious or repeated deception should be handled through the institution’s academic integrity process.

40.2 A constructive role for detectors

Although detector tools should not serve as the basis for misconduct decisions, the same study that documents their bias suggests a constructive alternative: detectors can be useful as student-facing self-check aids that flag clichéd phrases or repetitive patterns, encouraging the writer to vary structure or sharpen voice (Liang et al., 2023). Used this way, the detector is a writing tool the student applies to their own draft, not an enforcement tool the instructor applies to a submitted artifact. This use sits comfortably in the AIT assignment category: AI as feedback-giver on student-produced work, with the student responsible for accepting or rejecting the suggestions. Instructors who want to acknowledge the existence of detector tools without granting them evidentiary weight can recommend this self-check use directly.

41 Privacy, access, and equity

AI policy is not only about cheating. It is also about privacy, access, disability, language, and fairness. The four concerns interact: a policy that gets one right can quietly damage another. The subsections below treat them in turn.

41.1 Privacy

Students and instructors should not upload protected or confidential material into unapproved AI tools. This caution is consistent with European guidance that asks educators to engage with AI and data systems critically, ethically, and with attention to risk (European Commission, Directorate-General for Education, Youth, Sport and Culture, 2022). Material that should not be uploaded without explicit institutional approval includes identifiable student information, private correspondence, unpublished research data, human-subject data, patient or client information, restricted exam materials, recommendation letters, and copyrighted course materials when sharing is not permitted.

The instructor's privacy responsibility is not limited to declining to upload material themselves. It extends to course design: an assignment that requires students to upload classmates' work, group-project drafts containing other students' contributions, or interview transcripts collected under research-ethics protocols can create privacy exposures the students may not anticipate. When in doubt, consult institutional research-ethics or data-governance offices before requiring AI use that processes information about identifiable third parties.

41.2 Access

Some students may have paid AI tools, faster devices, better internet access, or more experience with AI. A course that requires AI should either provide equitable access or offer a no-cost alternative. A course that permits AI without requiring it should avoid giving hidden advantages to students who can afford premium tools.

Access has three layers. Tool access (free vs. paid tiers; rate limits; capability differences across models) is the most visible. Device access (a laptop with reliable browser support vs. a phone or shared library computer) is less visible but more durable. Internet access (campus wifi vs. home connections that may be intermittent or absent) determines whether AI is usable at all outside class.

The practical instructor response is a baseline-plus pattern: design assignments so the baseline AI capability available through institutional licensing or free tiers suffices to complete them, then accept that students with premium access may do additional exploration. Avoid assignments where premium tools confer a measurable assessment advantage that other students cannot match.

41.3 Equity in assessment design

Different assessment designs distribute their burdens differently across students. Process-evidence assessment can favor students with the time and home circumstances to document their work in detail. In-class no-AI assessment can favor students whose lives support uninterrupted focus. Oral defense can favor students whose language background, neurotype, or personality make spoken explanation comfortable. Each modality has its own equity profile,

and a course that uses only one modality concentrates the equity exposure on whichever students it disadvantages.

The hybrid approach to assessment (section 37) is partly a response to this. By distributing assessment across multiple modalities, the course gives most students at least one assessment modality where they perform near their best, rather than asking every student to perform at their best in the single modality the course happens to have chosen. This is not a guarantee of fairness, but it spreads the burden more evenly than monomodal assessment does.

Instructors who notice persistent gaps in performance across demographic groups should treat the gap as a design question first. Is one assessment modality contributing disproportionately to the gap? Is the documentation burden falling unevenly? Is the no-AI assessment running at a time of day or under conditions that disadvantage some students? These questions surface design adjustments more usefully than attribution to AI use itself.

41.4 Multilingual students and accessibility

Multilingual students deserve special care. Automated judgments about authorship, fluency, or style can penalize students whose writing differs from dominant patterns. Detectors specifically misclassify non-native English writing at high rates (section 40). Instructors should rely on process evidence and conversation rather than detector scores or vague impressions.

AI tools can be a genuine support for multilingual students, helping with idiomatic phrasing, vocabulary suggestions, and grammar checks. The Part I configuration template (section 26.1) explicitly addresses this use, and it sits comfortably in the AIT assignment category. Instructors should make this use legitimate rather than treating evidence of AI-assisted English revision as suspect.

Accessibility can cut both ways. AI may help some students brainstorm, read, summarize, translate, or organize work. It may also create new burdens if tools are inaccessible, expensive, or poorly integrated with assistive technology accommodations. Instructors should coordinate with campus accessibility offices when AI tools become part of required coursework, and should provide alternative access paths for students whose accommodations make standard AI interfaces inappropriate.

42 Discipline overlays

The framework is general, but implementation should be disciplinary. The same CAT structure can be translated into local practice.

42.1 STEM courses

In STEM courses, Core Competence may include conceptual understanding, problem solving, modeling, lab interpretation, coding fundamentals, and quantitative reasoning. AI-assisted Practice may include generating practice problems, debugging code, comparing derivations, checking edge cases, or critiquing a model. Trustworthiness includes unit checks, benchmark tests, reproducibility, source verification, and clear separation between evidence and speculation.

Strong STEM assignments often require a first attempt, a verification step, and a transfer question. For example, students might submit a solution, an AI critique of the solution, and a short no-AI explanation of the central method. Graduate-level mathematical training has additional structural requirements covered in Appendix A, including audit-and-repair tasks and the two-transcript grading scheme.

Mathematics courses divide along an axis that the broader STEM overlay does not always make explicit. In foundational proof-based courses (real analysis, abstract algebra, topology, transition-to-proof), NAI and AIT assessments often focus on definitions, short proofs, and proof-audit-and-repair tasks; AIC assignments may use AI to generate examples and non-examples or to critique a proof attempt, with oral defense for ownership. In computational and applied courses (numerical analysis, scientific computing, optimization, modeling), AIC assignments can integrate AI for code generation or method comparison under a verification rubric covering edge cases, stability, and reproducibility. The two patterns share the same loop, but they have different default category mixes and different assessment artifacts; the design choices follow from which kind of mathematical thinking the course is trying to build.

42.2 Writing and communication courses

In writing courses, Core Competence includes argument, structure, evidence, audience, revision, and voice. AI can help students ask revision questions, identify weak transitions, generate counterarguments, or compare possible openings. It should not erase the student's authorship. Disciplinary guidance from the Modern Language Association and the Conference on College Composition and Communication recommends teaching students to engage critically with AI-generated text rather than treating writing courses primarily as a site of detection (Modern Language Association and Conference on College Composition and Communication, 2023).

Strong writing assignments ask for draft history, an author's note, source verification, and a short explanation of revision choices. Students should be able to say why a sentence, paragraph, or claim changed.

42.3 Humanities courses

In humanities courses, Core Competence includes interpretation, close reading, source use, historical context, and argument from evidence. AI may help students generate background questions or compare interpretations, but primary sources and assigned texts should remain central.

Strong humanities assignments ask students to tie claims to specific passages, verify quotations, distinguish background summary from interpretation, and explain why an interpretation is plausible.

42.4 Social science courses

In social science courses, Core Competence includes research design, causal reasoning, data interpretation, theory comparison, and evidence evaluation. AI may help students generate possible confounders, critique survey questions, compare explanations, or summarize methodological options.

Strong assignments ask students to separate observed evidence from conjecture, justify methods, identify assumptions, and test alternative explanations.

42.5 Professional and applied courses

In professional courses, Core Competence often involves judgment under constraints. Students may need to apply standards, weigh risk, protect confidentiality, communicate with clients or patients, or justify decisions.

AI may support case simulation, draft critique, risk identification, or policy comparison. It should not replace professional responsibility. Assignments should ask students to explain their decisions and identify what information would change their recommendation.

The professional disciplines that have engaged most with AI in coursework so far are nursing, law, engineering, and business. Each raises a distinct version of the same pedagogical challenge: how to develop judgment when AI tools can produce confident-sounding answers that look professional but may be wrong.

Nursing. Core Competence in nursing courses includes clinical reasoning, patient assessment, ethical decision-making, and communication. AI may help students practice differential reasoning, generate case variations, or draft patient-education materials for critique. AI should not generate care recommendations that students submit without verification against course standards and clinical guidelines. A strong assignment asks students to compare an AI-generated triage suggestion with established clinical decision rules, identify where the two diverge, and explain which they would follow and why.

Law. Core Competence in law courses includes case analysis, statutory interpretation, legal writing, and reasoning under uncertainty. AI may help students explore alternative argument structures, generate counterarguments to a draft brief, or summarize unfamiliar areas of law for orientation. AI should not generate legal citations or substantive arguments students submit without verification. Hallucinated case citations are a documented risk; some appellate courts have already sanctioned attorneys for filings that included AI-generated fictitious cases. A strong assignment asks students to critique an AI-generated argument for hidden weaknesses, missing distinctions, and unverified citations.

Engineering ethics and applied engineering. Core Competence in engineering courses includes safety reasoning, standards compliance, ethical analysis, and professional judgment in design. AI may help students explore design alternatives, identify potential failure modes, or compare ethical frameworks. AI should not produce safety analyses or design decisions that students submit without independent verification. A strong assignment asks students to use AI to generate a list of potential failure modes for a design, then verify which are real engineering concerns versus surface-level pattern matches that do not apply.

In all three disciplines, the framework's core practices apply: independent attempt before AI use, verification against authoritative sources, disclosure of AI involvement, and the ability to defend decisions without AI. The professional context raises the stakes of failure, which is precisely why the disciplines have moved most quickly to develop explicit AI policies.

42.6 Creative and design courses

In creative and design courses, AI can support ideation, variation, critique, and iteration. The risk is that students may mistake generated novelty for artistic or design judgment.

Strong assignments ask for process portfolios, design rationales, iteration histories, constraints, and critique responses. Students should be able to explain why they chose one path rather than another.

The creative and design disciplines raise authorship and originality questions that other fields can largely sidestep. Where an engineering student is judged on whether a design works, a creative writing student is judged partly on whether the work is theirs.

Visual arts and design. Core Competence in visual arts includes observational skill, compositional judgment, technique, and the development of a personal visual vocabulary. AI image generation can produce visually polished output that bypasses each of these. A strong assignment requires sketches, drafts, iteration, and a written rationale explaining the artistic choices; AI may be a research tool or a comparison point, but the submitted work must show the student's hand. Some programs have moved to in-studio work as a foundation, with at-home assignments restricted to specific iteration tasks.

Music and composition. Core Competence in composition includes harmonic reasoning, voice leading, structural awareness, and the development of musical ideas over time. AI tools can generate plausible-sounding compositions that lack the structural reasoning a course is meant to build. Strong assignments ask for sketch development, theoretical annotations, and live performance or sight-reading components that confirm the student’s musical understanding.

Creative writing. Core Competence in creative writing includes voice, observation, structural choice, and revision. AI can produce text that is fluent but generic, or that imitates specific styles without the underlying observation that produced them. Strong assignments ask for drafts at multiple stages, marginalia explaining revision choices, and in-class writing under timed conditions that ground the work in the student’s own voice.

Design (industrial, graphic, UX). Core Competence in design includes user understanding, constraint analysis, iteration, and the rationale behind aesthetic choices. AI can generate variations quickly, which is valuable for ideation but dangerous if students mistake variation for design thinking. Strong assignments ask for user research, constraint documentation, iteration narratives, and presentation defenses where students explain why specific decisions serve specific users.

Across creative disciplines, the framework’s core insight applies with extra force: the goal of the course is to develop the student’s judgment, not to optimize the polish of their output. AI accelerates polish more than it accelerates judgment, so courses that grade on polish alone become particularly fragile in the AI era. The remedy is not to reject AI but to grade what the course is meant to teach.

43 Rubrics

A shared rubric helps students understand expectations and helps instructors grade consistently. The following five anchors work across many disciplines.

Correctness. The work is accurate, well supported, and appropriate to the course level.

Reasoning visibility. The work shows the path, not only the destination.

Verification discipline. Important claims, sources, calculations, code, or interpretations are checked.

Communication. The work is clear, organized, and suited to its audience.

Integrity. AI use is disclosed and consistent with the assignment policy.

For NAI assignments, the integrity anchor mainly concerns compliance with the no-AI rule.

For AIT assignments, it concerns whether AI was limited to tutoring. For AIC and AIS assignments, it concerns disclosure, verification, and ownership.

43.1 A simple four-level rubric

Level	Description
Excellent	Work is accurate, well explained, verified, and honestly documented. The student can defend the main decisions.
Satisfactory	Work meets the main goals. Some reasoning or verification is incomplete, but the student shows ownership.
Developing	Work contains gaps in understanding, verification, or documentation. The student needs more guided practice.
Unsatisfactory	Work cannot be explained, contains serious unsupported claims, violates assignment rules, or lacks ownership.

44 TA calibration and consistency

Large courses need consistent grading practices. A good rubric is not enough. TAs and instructors should practice applying it together.

A short calibration process can be used before major grading.

1. The instructor selects three to five sample submissions.
2. Each grader scores the samples independently.
3. The teaching team compares scores and discusses disagreements.
4. The rubric is clarified where needed.
5. Graders proceed in batches, with occasional spot checks.
6. The instructor reviews unusual or borderline cases.

The goal is not mechanical uniformity. The goal is shared interpretation of the course standards.

45 Implementation workflow for a course

For instructors redesigning under time constraints, three high-leverage moves are: assigning each task a NAI, AIT, AIC, or AIS label; adding one meaningful independent-competence

check per unit; and adding a brief AI-disclosure requirement to any AIC assignment. These three are the smallest set of moves that materially change how AI fits into a course.

The fuller workflow has seven steps.

1. Identify the core learning goals that must remain independent.
2. Assign each major task a NAI, AIT, AIC, or AIS category.
3. Choose the course model: Oral Defense, Evidence of Thinking, or Hybrid.
4. Add disclosure and verification requirements where AI is allowed.
5. Add at least one meaningful independent competence check.
6. Prepare a fair process for suspected misuse.
7. Tell students why the policy exists.

The last step matters. Students are more likely to follow a policy when they understand its educational purpose. A useful explanation is simple: the course allows AI where it supports learning and restricts AI where the instructor needs to see what students can do on their own.

46 Evidence base and limits

The instructor recommendations in this guide rest on the shared evidence base summarized in section 24 of Part I. Two findings matter most for course design. First, the AI-tutoring evidence is mixed: AI can help when designed around good pedagogy and students remain cognitively active; unrestricted access often improves apparent performance while weakening later independent performance. Second, cross-level meta-analyses find moderate positive overall effects but substantial heterogeneity by discipline, intervention duration, and assessment type (Liu et al., 2025; Ma and Zhong, 2025)—heterogeneity that the discipline overlays in this Part are designed to address.

The result is a design principle rather than a universal prescription. AI use should be allowed when it advances the learning goal and constrained when it hides the learning goal. A policy that is too permissive may reward polished output without understanding. A policy that is only prohibitive may fail to teach students how to use tools they will encounter outside the course. The instructor's task is to preserve independent competence while teaching responsible tool use.

47 Reference materials

The following templates and worksheets are intended for instructor use during course design and assessment. Many can be adapted directly into a syllabus, gradebook, or course-

management system. Department- and institution-level reference materials appear in Part III; the global technical appendices (GMIP, M/A/E framework, scientific writing) appear at the end of the volume.

47.1 Sample syllabus language

The following language may be adapted to fit a course. For instructors who prefer to start from a pre-written briefing rather than adapt the language below, the companion repository's `instructors/briefings.md` provides three ready-to-use syllabus paragraphs at different permissiveness levels (restrictive, default, and permissive); each is about 180 words and designed to drop into a syllabus without modification.

General course statement

AI tools are part of modern academic and professional work. In this course, you will learn when AI can support your learning and when it would interfere with the goals of an assignment. Each assignment will state whether AI use is prohibited, allowed for tutoring only, allowed as part of the working process, or required as an object of critique. When AI use is allowed, you are responsible for verifying the output, disclosing your use when asked, and submitting only work that you can explain.

No-AI assignment statement

This is a no-AI assignment. You may not use AI tools to generate, revise, translate, summarize, solve, check, or explain your answer while completing this task. The purpose is to assess what you can do independently.

AI-as-tutor assignment statement

You may use AI to study for this assignment by asking for explanations, examples, practice questions, or hints. You may not use AI to produce the submitted response. The submitted work should be written or completed by you without AI-generated text, code, or solutions.

AI-as-collaborator assignment statement

You may use AI as part of your working process. You must disclose the tools used, describe what they contributed, and explain how you verified the final work. You are responsible for all claims, sources, calculations, code, and interpretations in your submission.

47.2 Assignment label template

Field	Instructor entry
Assignment title	
AI category	NAI / AIT / AIC / AIS
Allowed uses	
Disallowed uses	
Required disclosure	
Required verification	
Independent component	
Privacy warning	
How this will be graded	

47.3 Student disclosure block

Tools used: Name the tools used, if any.

Tool role: Describe what the tools helped with.

Student role: Describe what you wrote, revised, checked, or decided.

Accepted, changed, or rejected: Describe the AI contributions you used, revised, or set aside.

Independent verification: Describe how you checked important claims, sources, data, code, or reasoning.

What I can explain without AI: Name the main parts of the work you can now explain on your own.

47.4 Oral defense rubric

Criterion	Evidence
Conceptual mand	Student explains central ideas in their own words.
Reasoning	Student can reconstruct the main argument, method, design, or interpretation.
Verification	Student identifies how important claims or outputs were checked.
Transfer	Student can adapt the idea to a nearby case.
AI process	Student explains what AI contributed and what was accepted, rejected, or revised.
Integrity	Student's explanation is consistent with the submitted disclosure and course policy.

47.5 Random oral audit script

A short audit may use the following sequence.

1. Please summarize your submission in two minutes.
2. What was the most important decision you made?
3. Did you use AI? If so, what did it contribute?
4. Show one place where you verified a claim, source, computation, code result, or interpretation.
5. What is one weakness or limitation in your submission?
6. Answer this nearby transfer question without AI.

The audit should be brief, structured, and consistent across students.

47.6 AI critique assignment template

Give students an AI-generated response to a course-relevant task. Ask them to complete four steps.

1. Identify what the response does well.
2. Identify at least two weaknesses, errors, unsupported claims, missing assumptions, or misleading statements.
3. Repair one part of the response.

4. Explain what this example teaches about using AI in the discipline.

This assignment works in many fields. A mathematics course can critique a proof. A writing course can critique an argument. A history course can critique a source summary. A computer science course can critique generated code. A lab course can critique an experimental explanation.

47.7 Course design checklist

- ☐ I identified the core abilities students must demonstrate independently.
- ☐ Each major assignment has an AI-use category.
- ☐ Students receive clear disclosure instructions where AI is allowed.
- ☐ AI-assisted assignments include verification or process evidence.
- ☐ The course includes meaningful no-AI or bounded-AI assessments.
- ☐ Large-class assessments include scalable evidence of thinking.
- ☐ Small-class assessments include oral defense, conferences, or live explanation where feasible.
- ☐ The syllabus explains privacy restrictions.
- ☐ The academic integrity process does not rely on AI detectors alone.
- ☐ TAs, if present, have a calibration process.

47.8 Instructor self-check

Before using AI in teaching or grading, ask:

- ☐ Am I using an approved tool for this kind of data?
- ☐ Does the AI output affect a grade or only help me prepare feedback?
- ☐ Have I reviewed the output before sharing it with students?
- ☐ Could this use expose private student information?
- ☐ Could this use create unfairness across students or sections?
- ☐ Am I still making the final instructional or evaluative judgment?

47.9 Short glossary

AI-assisted work. Work shaped by an AI tool through brainstorming, drafting, coding, revision, explanation, critique, or feedback.

Assessment validity. The degree to which an assessment measures the competence it claims to measure. Distinct from reliability (consistency of measurement) and fairness (procedural justice in the assessment process).

Core Competence. The knowledge and skills students must demonstrate without AI or with tightly bounded assistance.

Evidence of thinking. Process artifacts that reveal reasoning, choices, revision, verification, or explanation.

Oral defense. A live conversation in which a student explains submitted work and answers questions about it.

Process log. A short record of how a student approached the task, used AI, checked output, and revised the final work.

Trustworthiness. The ability to verify, disclose, critique, protect privacy, and take ownership of AI-assisted work.

Verification. Checking AI-assisted output against reliable sources, data, tests, methods, or reasoning.

Publication note

This manuscript is prepared as part of a three-guide package for possible publication and institutional adoption. The source is intentionally conservative in formatting: standard L^AT_EX structures, stable section labels, explicit templates, and minimal custom commands. Final production should use the current SIAM book macro package specified by SIAM Publishing at the time of submission.

Part III

A Framework for AI-Assisted Learning — A Departmental and Institutional Guide

*A workable policy is not one rule for every course.
It is a vocabulary every course can use.*

Part III is written primarily for the people who shape AI-assisted learning at the level where most readers will have direct influence: department chairs, curriculum committees, faculty

governance bodies, and the senior faculty within mathematics, science, engineering, and adjacent technical departments. It is also written for the smaller readership that engages with institutional-level policy: deans, provosts, teaching centers, academic-integrity offices, faculty senates, and university-wide administrative leaders. The framework gives both audiences practical tools for designing policy that protects learning while making AI-assisted practice an explicit part of academic life.

The central premise is simple: students will use AI, and academic units should teach them to use it without surrendering their own thinking. The framework separates Core Competence from AI-assisted Practice and Trustworthiness. It supports different course types, from small seminars to large multi-section courses, and recognizes that policy decisions sit at different levels: some belong to the department, some belong to the institution, and a few belong to both.

A note on scope and audience. The book's primary readers are faculty and graduate students in mathematics, science, engineering, and adjacent technical fields. These readers usually have substantial influence within their own department and limited influence over university-wide policy. The framework is written accordingly: department-level guidance is concrete and actionable, while institutional-level guidance is structured as what departments should expect from their institution and how to engage productively with institutional policy bodies. Readers in other roles (deans, provosts, central administrators) will find the institutional sections more directly addressed to their work.

Part III is organized in three logical movements rather than three rigid blocks. The first movement (Purpose and scope through Disclosure standard) establishes shared foundations both levels need: principles, vocabulary, the CAT framework, course-level AI-use categories, assessment models, the minimum-independent-competence requirement, and the disclosure standard. These topics are department-actionable and institution-relevant simultaneously, and the prose is calibrated for whichever level the reader is operating at. The second movement (Privacy through Risk register) addresses topics that are predominantly institutional in scope, with explicit notes on where departmental engagement is appropriate. The third movement (Looking forward through Reference materials) gathers the forward-looking discussion, evaluation metrics, model policy language, and reference tools that both levels share. Throughout, the prose distinguishes department-actionable practice from what departments should expect from their institution.

48 Purpose and scope

AI is now part of the ordinary academic environment. Students use it to explain concepts, draft prose, write code, summarize material, generate examples, prepare for exams, and test arguments. Faculty use it to design assignments, draft feedback, generate cases, refine explanations, and organize course materials. These uses will vary across tools and disciplines, but the policy problem is already clear at every level: departments and universities need a framework that protects learning without pretending that AI can be kept outside academic

life. A policy based only on prohibition will be brittle. A policy based only on permission will be too vague. A useful framework must help each course answer practical questions: What uses of AI are allowed? What must students still be able to do without AI? How should students disclose use? How should instructors grade AI-assisted work? How should suspected misuse be handled fairly?

Institutions have faced versions of this before. When calculators became widely available, schools that simply banned them protected the old form of arithmetic instruction without asking whether that form was still what it was for. Schools that simply allowed them, without rebuilding the surrounding judgment skills, watched students lose number sense without realizing they had lost it. Useful policy came later, after enough institutions worked out what calculator use required of the curriculum it was embedded in: which skills students still needed to demonstrate without the tool, which skills the tool enhanced rather than replaced, and how to teach the difference. AI is at the same point. The framework in this guide is an attempt to provide what most institutions are still missing: shared categories that distinguish AI use that supports learning from AI use that substitutes for it, and shared vocabulary for the policy decisions those categories enable.

This guide provides a framework for those decisions at the level where most decisions actually get made. It is not a single rule for every course. A first-year writing course, a physics lab, a history seminar, a programming course, a nursing simulation, and an engineering design studio need different practices. They should still share a common vocabulary, common ethical expectations, and common safeguards.

The framework is designed for the full range of courses taught at a university: courses of any class size, level, and discipline. The framework's design choices vary along these three axes (section 31 of Part II), and the structure here is intended to support that variation rather than impose a single uniform practice. Most of the practical work happens at the departmental level, and the prose is calibrated accordingly: where a section is primarily department-actionable, the framing speaks to faculty and chairs; where a section is primarily institutional, the framing speaks to deans, provosts, and central administrators; where both levels are involved, both audiences are addressed.

49 Institutional outcomes

A university that adopts this framework should be able to show progress in six areas.

1. Students receive clear guidance on how to use AI without surrendering their own thinking.
2. Faculty have shared categories for permitted, limited, and required AI use.
3. Courses preserve meaningful assessment of independent competence.
4. Departments use fair procedures for disclosure, privacy, access, and suspected misuse.
5. Teaching assistants and graders apply rubrics with reasonable consistency.

6. The institution reviews the framework regularly as tools, risks, and teaching practices change.

These outcomes give the framework a practical test. The policy is working only if it improves clarity, learning, fairness, and trust.

50 The Institution's Trilemma

Part I named the predicament students face under AI as the Student's Trilemma: a three-way pull, like a dilemma but with one more horn, among earning the grade, being honest about how the work happened, and demonstrating genuine understanding. Each pair pulls against the third. Almost every student now feels some version of it.

Institutions face the same structural problem in inverted form. Presidents, provosts, deans, and trustees are pulled between three goals that AI now makes harder to hold together: keeping the institution financially viable, defending what its credential actually certifies, and remaining honest about the educational mission the institution exists to serve. Each pair pulls against the third. Most administrators now feel some version of it, though they rarely name the pattern.

We call this the **Institution's Trilemma**. The corners are:

Institutional sustainability. Keeping the institution financially viable through enrollment quality and quantity, tuition, donor cultivation, research grant flow, and state appropriations, and keeping the public legitimacy that makes financial viability possible. Public skepticism erodes enrollment; falling enrollment erodes funding; falling funding erodes the institution's ability to make the case for its own value (Inside Higher Ed and Hanover Research, 2026). Administrators cannot work on one face of this without working on the other.

What the credential certifies. The degree as a meaningful claim about the institution's graduates: what they can do, what they understand, why they are worth hiring or admitting to graduate school. The credential's value depends on both the institution's testimony (what it says its graduates can do) and the underlying evidence (what they can actually do under direct examination). When AI assistance enters student work without being verified, testimony and evidence drift apart.

Educational mission. The deeper commitment to what students gain by attending. This varies by institutional type—Jesuit formation, liberal arts breadth, research apprenticeship, workforce readiness, professional preparation—but is always present in the institution's self-understanding and almost always present in the language it uses to describe itself publicly. The mission is what an administrator points to when asked why the institution exists beyond its operational and credentialing functions.

What unifies these institution-type-specific articulations is a single claim: that some learning is itself the payoff, not just a means to an end. A coder welcomes AI because

for them the drudgery is what they want eliminated and the working program at the end is the payoff. An artist resists AI because for them the drudgery—the patient hours of practice and revision—*is* the payoff and the finished piece is its trace. The educational-mission corner rests on the same distinction at the institutional scale. An institution that has accepted that all of its content can be delivered more efficiently by AI has tacitly accepted that nothing it does is itself the point. An institution that defends its educational mission is asserting the opposite: that some of what students pay for, and some of what graduates carry forward, is the intellectual practice itself, not just its end product.

Each pair of corners can be satisfied at the cost of the third. The failure modes echo the Student's Trilemma deliberately.

50.1 Three institutional failure modes

Sustainability and credential certification without mission. The institution maintains enrollment and credential reputation by becoming a sophisticated AI-gatekeeping operation. Students complete coursework, get verified through carefully designed oral defenses, receive degrees that mean what they say. But the institution has stopped trying to genuinely educate, in the sense the mission language implies, beyond the credentialing process itself. We call this the **institutional polished hollow**: the credential is intact, the public-facing story is intact, but the educational mission has hollowed out from inside. This is the failure mode of an institution that has outsourced the deep work of education to AI while retaining institutional gatekeeping for itself.

Sustainability and mission without credential certification. The institution maintains enrollment and stays committed to its educational mission, but lets AI-assisted work into the credential without demanding genuine ownership. Students graduate with degrees they cannot defend. The mission language remains genuine; the financial picture remains stable; the credential loses its meaning. We call this the **institutional strategic compliance**: the institution complies with the appearance of its credential function while quietly admitting that the verification was not done. This is the failure mode of an institution that wants to grow and stay true to its mission, but cannot bring itself to enforce the standards its credential implies.

Credential certification and mission without sustainability. The institution stays rigorous and mission-driven, refuses to lower standards in response to AI pressure, and runs out of students or money in the process. We call this the **institutional honest risk**: every commitment kept, no commitment surrendered, but the institution cannot keep the doors open. This is the failure mode most visible at the moment among small private tuition-dependent colleges that have closed in the past three years.

AI did not create the Institution's Trilemma any more than it created the Student's. The trilemma describes a structural tension that pre-existed AI and that any responsible institution always navigates. What AI did was sharpen each edge. Each pair of corners is harder to hold

together when AI assistance lowers the apparent cost of giving up the third.

50.2 How the trilemma manifests differently by institutional type

The three corners are universal among degree-granting institutions, but the pull strength on each corner differs by institutional type. Four types are worth distinguishing, with the understanding that real institutions often combine features of more than one.

R1 private research universities. At institutions like Stanford, Duke, and Tulane, sustainability is cushioned by endowments and tuition-pricing power, though research grant funding is now under direct AI strain: agentic-AI-assisted proposals are overwhelming funding agencies, which respond with submission caps and centralized review that reshape who gets funded (Rees and Wilsdon, 2026). The credential is the brand; protecting it is non-negotiable. Mission is articulated through the dual commitment to research training and undergraduate liberal education. The institutional polished-hollow failure mode is the most relevant risk: an institution that maintains its credential through sophisticated gatekeeping while quietly admitting that what it offers undergraduates has converged with what AI tutors offer for free.

Wealthy liberal arts colleges. At institutions like Williams, Pomona, and Kenyon, the trilemma takes its tightest form. Credential and mission corners are tightly coupled: the case for the institution *is* the case for the human-led residential experience. Sustainability is the threatening corner. With small enrollments and tuition dependence, the institutional honest-risk failure mode is the most visible and the most current. The closures of small private tuition-dependent colleges in recent years follow this pattern (Levine and Van Pelt, 2026). The case for these institutions in the AI era is not that they teach better content than larger universities, since AI tutors can deliver content well. It is that they teach in a setting where the human-led mentoring relationship is the educational point, not just the delivery mechanism.

Regional public universities. At most state-system institutions outside the research-flagship tier, sustainability dominates. The enrollment cliff hits hardest at this institutional type, and federal-plus-state funding cuts compound the pressure. Educational mission is increasingly framed in workforce-alignment terms rather than liberal-education terms, partly in response to state legislatures and partly in response to student-and-parent demand for clear post-graduation outcomes. Credential certification matters operationally (graduates need to perform when they leave) but is rarely the primary strategic conversation. The institutional strategic-compliance failure mode is the most relevant risk: an institution that maintains enrollment by accepting AI-assisted work without demanding genuine ownership, because the cost of demanding genuine ownership is enrollment the institution cannot afford to lose.

Community colleges. The trilemma is flatter here. Public trust in community colleges has remained strong even as trust in higher education broadly has eroded; one common explanation is that community colleges are seen as affordable, career-pathway-connected, and engaged with local employer needs (Pollard, 2026). Sustainability pressure exists but is less acute. The mission corner is more clearly articulated as access plus workforce alignment than

at other institutional types. The credential corner is dominated by a different question—credit transferability to four-year institutions—than at other types. Among the three failure modes the framework names, the one community colleges seem least vulnerable to is the honest-risk pattern; they are more vulnerable to local-mission drift if state policy shifts.

The four types are not comprehensive. Specialized institutions—HBCUs, religious-affiliated colleges, for-profit institutions, fully online providers—face the same trilemma but with type-specific pull strengths on each corner. The framework’s value is that the corners stay the same; what changes is the relative pull and the failure mode most likely to materialize.

50.3 Why the Institution’s Trilemma matters for the rest of this Part

The institutional policies developed in the rest of this Part—the assessment models, the integrity procedure, the faculty change-management approach, the implementation phases—are tools for holding the three corners of the trilemma together under AI pressure. Without the trilemma frame, these tools can read as a list of operational responses to a list of operational problems. With the trilemma frame, they read as coordinated commitments that protect the institution from each of the three failure modes simultaneously.

There is also a more concrete rhetorical purpose to this Part. Department chairs and program directors increasingly find themselves in conversations with deans, provosts, and boards about whether faculty roles can be replaced by AI tools at substantially lower cost. The Institution’s Trilemma supplies the structural answer to that question: the case for human-led education is the case that the educational-mission corner cannot be honored by a system that has hollowed out into pure credentialing. Throughout the rest of this Part, the operational tools developed—assessment models, integrity procedures, faculty change management, implementation phases—can be read as ammunition for the argument chairs need to make to deans when AI-replacement proposals arrive on the agenda. The argument is not nostalgic and it is not defensive. It is the argument that some of what students pay for, and some of what graduates carry forward, is the very thing that AI-as-replacement is supposed to eliminate.

The Northeastern case discussed in the next section can be read through this frame. It is not just a faculty integrity question. It is an early signal of the institutional polished-hollow pattern: an institution that maintained its credential and sustainability through procedural responses to a violation, but did not address the deeper mission question of what students were paying for.

51 The institutional problem

A senior at Northeastern University spent the final weeks of her undergraduate degree filing a formal complaint and a tuition refund request against her organizational behavior

professor. She had discovered that the professor was using ChatGPT, Perplexity, and an AI presentation generator to produce her lecture materials, with a stray AI prompt instruction visible in the slides. The syllabus for the course prohibited student AI use. The professor later acknowledged the AI assistance. The university declined to refund the tuition.⁸

The case names a pattern. The professor was wearing what the rest of this Part will call a **competence costume**: AI-polished material that signals more current and confident engagement than the underlying work supports. The student saw it. The institution declined to act on it. The asymmetry is the point. An integrity rule that applies to students but not to those teaching them is structurally fragile.

In the language of Part I (section 5), the Northeastern professor's costume was projected through *strategic compliance*: he understood his own course but hedged on disclosure. *Competence costume* names the outward projection; the taxonomy names the configuration that produces it.

The translation test introduced for students in Part I (section 6) gives institutions a positive standard to adopt symmetrically. Faculty who use AI assistance in course materials should be able to translate those materials into their own language and examples. Students using AI on coursework should be able to do the same. The same diagnostic, applied to both, is a policy the institution can defend in public.

AI creates a gap between polished performance and demonstrated learning. A student may submit fluent prose, working code, or a convincing solution without having developed the underlying competence. This problem is not limited to cheating. It can happen through ordinary over-reliance. A student may believe they understand because the final product looks good.

The institutional response should not be to treat every AI-assisted product with suspicion. The better response is to redesign expectations. Courses should assess both independent understanding and responsible tool use. Students should learn to verify, disclose, and defend AI-assisted work. Faculty should have clear models for small, medium, large, and online courses. Departments should not leave each instructor to invent policy alone.

The goal is a durable approach. AI systems will become stronger. Detection tools will remain imperfect. The safest assessments will be those that make student thinking visible and require students to explain, verify, apply, and revise their work.

Evidence from learning science supports this direction. Practice testing and distributed practice have strong support across learning contexts, and recent AI-tutoring studies suggest that carefully designed systems can help when they preserve active engagement rather than replace it (Dunlosky et al., 2013; Kestin et al., 2025). At the same time, field evidence warns that unrestricted generative AI can create short-term performance gains without durable learning (Bastani et al., 2025).

⁸Hill (2025).

52 Guiding principles

The following principles should guide institutional policy.

1. **Learning remains the center.** AI may support learning, but it should not replace the intellectual work students are meant to practice.
2. **Independent competence must be preserved.** Every course should identify what students must be able to do without AI or with tightly bounded support.
3. **AI literacy is a legitimate educational goal.** Students should learn how to use AI critically, productively, and ethically within their disciplines.
4. **Verification is required.** AI output should be treated as provisional until checked against evidence, data, sources, methods, or disciplinary standards.
5. **Disclosure should be normal, brief, and useful.** Students should know how to describe what AI contributed and what they verified.
6. **Assessment should reveal thinking.** Courses should gather evidence of reasoning, process, revision, judgment, and transfer, not only polished final products.
7. **Privacy and equity must be protected.** Students should not be required to use tools that compromise data, create avoidable costs, or disadvantage students by language background, disability, access, or prior exposure.
8. **Faculty judgment remains human.** AI may help instructors prepare feedback and materials, but instructors remain responsible for grades, academic-integrity judgments, and student evaluation.

These principles align with emerging international guidance that treats AI literacy, human oversight, ethics, privacy, and responsible educational design as central institutional responsibilities (Miao and Holmes, 2023; Miao et al., 2024; Miao and Cukurova, 2024; European Commission, Directorate-General for Education, Youth, Sport and Culture, 2022; OECD, 2026).

53 The CAT framework

University policy should distinguish three kinds of learning evidence: Core Competence, AI-assisted Practice, and Trustworthiness. The acronym CAT gives faculty and students a shared language.

Core Competence. What students must know and be able to do without AI or with strictly bounded assistance.

AI-assisted Practice. What students can do when they use AI responsibly as a tutor, critic, simulator, editor, coding aid, or research assistant.

Trustworthiness. How students verify, disclose, critique, and take ownership of AI-assisted work.

The distinction matters because these abilities can diverge. A student may produce strong AI-assisted work while having weak independent understanding. Another student may understand the material well but use AI carelessly. Good policy should not collapse these into one score.

Institutional guidance should therefore encourage courses to state which assignments measure Core Competence, which assignments develop AI-assisted Practice, and which assignments evaluate Trustworthiness. The student-facing development of the framework appears in Part I (especially section 7); the instructor-facing development appears in Part II (especially section 32). Graduate programs adopting the framework directly should consult Appendices A and B.

53.1 Alignment with UNESCO learning outcomes

UNESCO’s guidance (Miao and Holmes, 2023) proposes four broad categories of learning outcomes that universities should consider as they revise curricula for the AI era: *values* that ensure human-centred design and use of technology; *foundational knowledge and skills* (literacy, numeracy, basic scientific literacy) that remain prerequisites even when AI can perform specific tasks better than humans; *higher-order thinking skills* that ground critical evaluation of AI output and human-AI collaboration; and *vocational skills* required to work with AI systems in domains where AI is automating task units. The CAT framework above is consistent with all four. Core Competence corresponds to foundational knowledge and skills; Trustworthiness corresponds to values and to the critical-evaluation dimension of higher-order thinking; AI-assisted Practice corresponds to vocational skills and to the human-AI-collaboration dimension.

Together with the student-framework crosswalk (table 1) and the teacher-framework crosswalk (table 3) established earlier in this volume, the institutional framework here translates cleanly onto UNESCO’s structure for accreditation reporting, curriculum-revision proposals, and international-collaboration framing. National-policy environments outside UNESCO alignment can perform analogous mappings against their own framework documents (U.S. Department of Education, Office of Educational Technology, 2023); the CAT structure is intended to translate cleanly onto most national or regional frameworks rather than to compete with them. As universities revise curricula in the years ahead, this UNESCO categorization may prove a useful framework-independent reference point.

54 Shared vocabulary

A common vocabulary reduces confusion across courses. The following terms should be used consistently in university policy.

No-AI work. Work completed without AI assistance, except for accommodations or tools explicitly approved by the instructor.

Bounded-AI work. Work completed with limited, specified AI assistance, such as a single approved tool, a restricted prompt type, or a fixed time window.

Open-AI work. Work in which students may use AI broadly, subject to disclosure, verification, and course rules.

AI-assisted work. Student work improved with AI support but written, verified, and owned by the student.

AI-generated output. Text, code, images, explanations, summaries, solutions, data analyses, or other content produced by an AI system.

Assessment validity. The degree to which an assessment measures the competence it claims to measure. Distinct from reliability (consistency of measurement) and fairness (procedural justice in the assessment process).

Disclosure. A factual statement of which tools were used, what they contributed, and how the student verified the result.

Verification. The process of checking AI output against course materials, sources, data, code tests, experiments, professional standards, or independent reasoning.

Process evidence. Artifacts that show how work was produced, such as plans, drafts, prompt logs, version histories, verification notes, source checks, lab notebooks, or code tests.

Oral defense. A live explanation in which students answer questions about their work, reasoning, sources, methods, or AI use.

Random audit. A routine assessment check in which a student is asked to explain or verify submitted work. It is not, by itself, an accusation.

55 Course-level AI-use categories

Every assignment should state what kind of AI use is allowed. The following four categories are simple enough for students and flexible enough for departments.

55.1 **NAI:** No AI

AI use is not allowed. These assignments measure independent competence. Examples include closed-book exams, in-class writing, oral explanations, live coding checks, foundational calculations, language practice, and lab practicals.

NAI assignments should not dominate every course. They serve a specific purpose: they show what students can do without external scaffolding.

55.2 AIT: AI as Tutor

AI may be used for explanation, practice, review, hints, and study support, but not to produce submitted work. A student might ask AI to explain a concept, generate practice questions, or quiz them before an exam. The submitted response should still be the student's own unaided work.

55.3 AIC: AI as Collaborator

AI may help with brainstorming, drafting, revising, debugging, outlining, critique, source search, comparison, or testing. Students must disclose use and verify important claims. AIC assignments are useful when the course goal includes revision, professional practice, coding, research, modeling, or design.

55.4 AIS: AI as Object of Study

Students analyze AI output directly. They may identify hallucinations, critique bias, compare tools, repair flawed reasoning, evaluate generated code, examine source quality, or test how AI performs in a discipline. AIS assignments teach students to examine AI rather than merely use it.

56 Assessment models by course size

Class size changes what assessment can reasonably do. The same principles apply across courses, but the implementation must differ.

56.1 Oral-Defense Model

The Oral-Defense Model is best for small classes of any level: seminars, honors courses, capstones, writing workshops, advanced labs, research courses, and many graduate courses. The model is class-size-driven; it works wherever instructors can speak directly with each student about their work.

Students may use AI more extensively when they can be asked to defend what they submitted. A defense may be a board proof, source-analysis conversation, code walk-through, lab-method explanation, design critique, writing conference, or project presentation.

A course using this model should include at least one live explanation of AI-assisted work. The purpose is not to intimidate students. It is to confirm ownership and deepen learning. The operational treatment of this model, including recommended structure, suggested weighting, and equity considerations, appears in Part II (section 35).

56.2 Evidence-of-Thinking Model

The Evidence-of-Thinking Model is best for large lectures, service courses, online courses, and multi-section courses. In these settings, individual oral exams for every student may not be feasible.

Large courses should gather repeated, low-cost evidence of thinking. Useful tools include short no-AI quizzes, in-class writing, draft histories, source checks, process logs, AI-use disclosures, homework-to-quiz transfer questions, flawed-output critiques, code tests, lab notebooks, short video explanations, and random oral audits.

The goal is not surveillance. The goal is valid assessment. A polished final product is no longer enough evidence by itself. The operational treatment of this model, including scalable evidence sources, suggested weighting, and random oral audits, appears in Part II (section 36).

56.3 Hybrid Model

Many courses are too large for full oral exams but small enough for some live explanation. The Hybrid Model combines scalable process evidence with limited direct questioning. It may use group oral exams, sampled oral audits, TA-led conferences, short instructor meetings, video explanations, and written verification logs.

Hybrid designs are useful for medium-sized writing courses, upper-level STEM courses, lab courses, project-based courses, and professional programs. The operational treatment of this model, including a recommended structure with mapping table, sampling strategy for live audits, a worked example, and suggested weighting, appears in Part II (section 37).

57 Minimum independent-competence requirement

Every course should include some meaningful assessment of what students can do independently. The form should fit the discipline.

In STEM courses, this may be a quiz, exam, derivation, lab practical, coding check, or oral explanation. In writing courses, it may be in-class writing, a writing conference, or a revision defense. In humanities courses, it may be close reading, source analysis, or argument reconstruction. In professional programs, it may be a case explanation, simulation, or ethical decision defense.

Table 8: Recommended assessment model by course context.

Course context	Recommended model
Small seminar	Oral-Defense Model
Graduate course	Oral-Defense or Hybrid Model
Honors course	Oral-Defense or Hybrid Model
Medium undergraduate course	Hybrid Model
Large lecture	Evidence-of-Thinking Model
Multi-section course	Evidence-of-Thinking Model with TA calibration
Online course	Evidence-of-Thinking Model with stronger process documentation
Capstone or project course	Hybrid or Oral-Defense Model

The institution should not prescribe one method for every course. It should prescribe the principle: some course credit should reflect independent understanding, not only AI-assisted production.

A useful policy sentence is:

Every course should identify the core skills, concepts, or judgments students must demonstrate without AI or with tightly bounded assistance. The method of assessment should be appropriate to the discipline, level, enrollment, and learning goals of the course.

58 Disclosure standard

Disclosure should be normal and brief. It should not read like a confession. It should give instructors enough information to understand how AI shaped the work.

A recommended student disclosure block is:

Tools used: Name the tools used, if any.

Tool role: Describe what the tools helped with.

Student role: Describe what you wrote, revised, checked, or decided.

Accepted, changed, or rejected: Describe the AI contributions you used, revised, or set aside.

Independent verification: Describe how you checked important claims, sources, data, code, or reasoning.

What I can explain without AI: Name the main parts of the work you can now explain on your own.

For low-stakes assignments, a one-sentence disclosure may be enough. For substantial AI-assisted projects, instructors may require a process log, prompt excerpts, verification notes, or a short reflective memo.

59 Privacy, data governance, and tool approval

Students and faculty need clear rules about what may not be uploaded into AI systems. These rules should be plain enough to appear in syllabi and student orientation materials.

Unless a tool has been approved for the specific use, students and faculty should not upload:

- private information about students, classmates, patients, clients, research subjects, or employees;
- unpublished research data without permission;
- restricted exams, answer keys, or secure assessment materials;
- confidential institutional documents;
- private correspondence, recommendation letters, or evaluations;
- copyrighted course materials when sharing is not permitted;
- data governed by human-subjects, clinical, contractual, or legal restrictions.

The institution should maintain a list of approved tools and approved uses. A tool approved for general brainstorming may not be approved for student records, research data, grading, or protected course materials.

Data governance should also address retention, training use, third-party access, accessibility, cost, procurement, and support. These details should not be left to individual instructors alone.

60 Equity, access, and accessibility

AI policy should avoid creating a hidden paywall around learning. If a course requires AI use, students should have access to an approved tool at no additional cost or should be offered an equivalent alternative.

Equity also includes language background, disability, prior exposure, bandwidth, device access, and confidence using new tools. Some students will arrive with extensive AI experience. Others will be learning basic use for the first time. Courses should not mistake unequal access for unequal ability.

Accessibility can cut in both directions. AI tools may help some students summarize, rephrase, plan, or review material. They may also create barriers if tools are inaccessible, expensive, unstable, or poorly documented. Accommodation decisions should remain human and should follow institutional disability procedures.

The institution should also be careful with AI detectors. Evidence shows that some detectors misclassify writing by non-native English writers as AI-generated, raising concerns about fairness and reliability (Liang et al., 2023). For that reason, detection scores should not be treated as primary evidence of misconduct.

61 Academic integrity procedure

Academic integrity should be handled through evidence and due process, not through detector scores alone. A fair procedure should distinguish misunderstanding from deliberate misconduct. Early empirical work on ChatGPT and student writing emphasized that proactive course design and explicit policies are more effective than detection-driven enforcement (Cotton et al., 2024).

The case against detector-based enforcement is now quantitatively documented. In a study of seven widely used GPT detectors evaluated on TOEFL essays written by non-native English speakers, the average false-positive rate was 61.3%; all seven detectors unanimously misclassified 19.8% of the essays; and at least one detector flagged 97.8% of them as AI-generated (Liang et al., 2023). The mechanism is identifiable: most current detectors rely on text perplexity, a measure of how predictable a sequence of words is to a language model, and non-native writers exhibit lower lexical and syntactic variability that produces lower perplexity that detectors then read as AI-generated. The same study shows that detectors can be bypassed by prompting an AI to use more literary language, with detection rates falling to near zero. Detectors therefore fail in two directions, and both directions disadvantage students who can least afford the error: false positives concentrate on multilingual writers, and false negatives accrue to students with the resources and skill to refine prompts. Beyond accuracy, detector-based enforcement produces what those authors term a presumption of guilt, an institutional posture in which automated systems flag students as suspect and the burden falls on them to prove innocence; the reputational consequences of a flagged accusation may persist even when the accusation is later revoked.

The accuracy concerns translate into harm at the case level. A senior at the University of California Davis whose paper was flagged by Turnitin as AI-generated spent the final months of her senior year defending against the accusation; her grades reportedly began to slide as she fought it.⁹ The arms race that detector reliance creates has produced its own absurdity: students now use “humanizer” tools that adjust AI-generated text to defeat detection, and some students who do not use AI at all subscribe to the same tools to reduce the chance of being falsely flagged. As one professor quoted in recent reporting put it, the better the writer

⁹Klee (2023).

you are, the more AI thinks you're AI.¹⁰ Detector dependence does not create the integrity problem. It produces a parallel problem of its own.

A recommended process is:

1. Review the assignment rules and the student's disclosure.
2. Examine process evidence, such as drafts, version history, source notes, lab notebooks, code history, or prompt logs.
3. Compare submitted work with in-class work when relevant.
4. Invite the student to explain the work in a non-accusatory meeting or written response.
5. Use a short reassessment or oral explanation when needed.
6. Document the concern, evidence, student response, and decision.
7. Use educational remedies for minor first violations when appropriate.
8. Use formal academic-integrity procedures for serious, deceptive, or repeated misconduct.

The institutional standard should be clear:

AI-detector output may be considered, if allowed by institutional policy, but it should not serve as the sole or primary evidence of academic misconduct. Decisions should rest on assignment rules, process evidence, student explanation, and professional judgment.

62 Faculty use of AI

The framework applies to faculty as well as students. Faculty may use AI to support teaching, but they remain responsible for professional judgment.

The Northeastern case discussed in section 51 is one of the early reported instances where this responsibility was tested in public. The translation test offers an internal standard faculty can apply to themselves: a faculty member using AI assistance on a piece of course material should be able to translate the material into their own language, with their own examples, in their own structure, before placing it in front of students. If the translation is awkward, the integration is incomplete.

Appropriate uses may include drafting examples, generating practice questions, identifying possible misconceptions, improving assignment wording, drafting feedback for human review, creating case variations, and preparing discussion prompts.

¹⁰Kingkade (2026); the line is attributed in that article to Erin Ramirez of California State University, Monterey Bay.

Faculty should not delegate final grades, academic-integrity decisions, accommodation decisions, or sensitive student evaluation to AI. They should not upload protected student work, student records, or confidential materials into unapproved systems.

A concise policy sentence is:

Instructors may use AI to assist teaching preparation and feedback drafting, but the instructor remains responsible for all grades, evaluative judgments, academic-integrity decisions, and protection of student information.

63 Required syllabus elements

Every syllabus should answer five questions.

1. What AI uses are allowed or prohibited in this course?
2. Which assignments are NAI, AIT, AIC, or AIS?
3. What must students be able to do without AI?
4. How should students disclose AI use?
5. What materials may not be uploaded into AI systems?

For large or multi-section courses, syllabi should also explain how the course will check understanding. Examples include short no-AI quizzes, process logs, random oral audits, in-class writing, draft histories, or TA-led conferences.

64 Departmental responsibilities

Departments should not require every instructor to use the same AI policy. They should require every course to make its policy clear.

At the department level, useful responsibilities include:

- adopting shared definitions and assignment categories;
- identifying discipline-specific standards for verification;
- developing sample assignments and rubrics;
- supporting faculty who teach large or multi-section courses;
- calibrating TA grading when process evidence is assessed;
- reviewing course policies for consistency and fairness;

- updating policies as tools and student practices change.

Departments should also identify courses where AI literacy is itself a learning outcome. In those courses, students should practice critique, verification, prompt design, error detection, and responsible disclosure.

65 Teaching-center responsibilities

Teaching centers can turn policy into practice. They should provide templates, examples, workshops, and consultation for faculty.

High-value resources include:

- a syllabus-statement bank;
- sample NAI, AIT, AIC, and AIS assignments;
- discipline-specific verification examples;
- rubrics for process evidence and disclosure;
- guides for oral defense and random audits;
- TA calibration materials;
- privacy and data-use checklists;
- academic-integrity conversation scripts;
- student-facing learning modules.

The most useful teaching-center materials are short, editable, and course-ready. Faculty adoption improves when instructors can adapt templates rather than build from scratch.

66 Training plan

An institutional framework should include training for students, faculty, TAs, and academic leaders.

66.1 Student training

Student orientation should teach the Learning Spiral (figure 6): try first, ask for help rather than replacement, work independently, verify, and disclose. Students should also learn what not to upload, how to check sources, and how to prepare for no-AI assessments. Institutions

can reduce friction further by providing a recommended baseline AI configuration—a starter set of personal preferences students can adopt during orientation—since most students do not configure their AI accounts without prompting and benefit substantially from doing so.

A short module can be completed in the learning-management system before the first AI-permitted assignment. Departments may add discipline-specific examples.

66.2 Faculty training

Faculty training should focus on course design rather than tool novelty. Useful topics include assignment redesign, assessment validity, process-based grading, responsible feedback, privacy, equity, suspected misuse, and discipline-specific AI literacy.

Faculty should leave training with a revised syllabus statement, at least one redesigned assignment, and a plan for assessing independent competence.

66.3 TA training

TAs need practical calibration. They should learn how to grade disclosure blocks, process logs, audit-and-repair tasks, and short explanations. They should also know what to do if they suspect misuse.

A strong TA calibration session uses sample submissions. The instructor and TAs grade the same examples, compare judgments, revise the rubric, and document standards before grading the full class.

66.4 Administrative training

Administrators and academic-integrity officers need a shared understanding of the institution's AI categories, detector limits, privacy rules, and due-process expectations. This reduces inconsistent handling across departments.

67 Implementation plan

The four-phase implementation described below sits within a broader international policy context. UNESCO's guidance on generative AI in education proposes seven steps for regulating GenAI in education systems: endorsing or developing data-protection regulations; adopting whole-of-government AI strategies; solidifying ethics regulations; adjusting copyright laws to address AI-generated content; elaborating regulatory frameworks specific to generative AI; building capacity for proper use among teachers, learners, and researchers; and reflecting on the long-term implications of these tools (Miao and Holmes, 2023). National and regional

governments are responsible for most of these steps, and the framework in this Part is designed to be compatible with whichever regulatory environment a particular institution operates in. The four-phase plan that follows is the institutional-level operationalization within whatever national regulatory environment applies; the policy structure here is meant to slot into national guidance rather than substitute for it.

A university should adopt AI-assisted learning policy in phases (figure 12).

67.1 Phase 1: Pilot

Select a small set of courses across disciplines and sizes. Include at least one large undergraduate course, one writing-intensive course, one STEM course, and one seminar or capstone. Test the shared vocabulary, assignment labels, disclosure form, and assessment models.

Collect feedback from students, instructors, TAs, and academic-integrity staff. Revise the materials before wider adoption.

67.2 Phase 2: Department adoption

Departments adapt the framework to local needs. They identify core competence standards, discipline-specific verification practices, sample assignments, and course-size models. They also decide which courses should include explicit AI literacy outcomes.

67.3 Phase 3: University adoption

The university integrates the framework into syllabus guidance, teaching-center resources, student orientation, academic-integrity procedures, and data-governance policy. The policy should remain flexible enough for local course design.

67.4 Phase 4: Review and revision

AI policy should be reviewed annually. The review should examine student learning, faculty workload, integrity cases, equity concerns, privacy incidents, and student feedback. The goal is not permanent policy language. The goal is stable principles with adaptable implementation.

67.5 Cost and resource considerations

The four-phase plan above describes what to do but not what it costs. Honest implementation discussion has to acknowledge that AI-aware teaching is a real institutional investment. Disclosure-protocol design and faculty workshops in Phase 1 typically require teaching-center

staff time and a modest stipend pool for pilot instructors. Process-evidence rubrics in Phase 2 increase per-assignment grading time, which means departments with high enrollment-to-instructor ratios may need to reallocate teaching assistant support or accept somewhat coarser rubrics. Oral assessment, the most future-proof element of the framework, is the most expensive to scale; small classes can absorb it within existing contact hours, while large classes typically cannot without TA support or a hybrid in-class/recorded model. Calibration sessions across instructors are the lowest-cost element on a per-session basis but require persistent scheduling investment that departments often underestimate.

For resource-constrained departments, a useful low-cost adoption path is to begin with three minimal-cost elements that do not require new staffing: explicit per-assignment AI policies (faculty time only); the disclosure-block standard (one-time setup); and a single calibration session per term (faculty time only). The full framework can then be phased in over multiple years as resources allow. The goal of the implementation plan is not to require simultaneous adoption of every element, but to give departments a coherent target state to grow toward.

For institutions in UNESCO-aligned policy environments, the CAT-to-UNESCO crosswalks established earlier in this volume (Tables 1 and 3, and the learning-outcomes mapping in section 53.1) can be used directly during accreditation reporting, curriculum-revision proposals, and international-collaboration framing. Where the institution's accreditation language references UNESCO student or teacher AI competence frameworks, the crosswalks make the alignment explicit and reduce the translation work for adopting departments. National-policy environments outside UNESCO alignment can perform analogous mappings against their own framework documents; the CAT structure is intended to translate cleanly onto most national or regional frameworks rather than to compete with them (U.S. Department of Education, Office of Educational Technology, 2023).

68 Faculty change management

The framework laid out in the previous sections describes what an institution should do. It does not describe how a university actually moves from current practice to that target state. The gap between the two is faculty change management, and it is where most well-designed AI policies fail.

68.1 The change-management problem

Faculty are not, in general, opposed to AI policy in principle. Most instructors recognize that AI tools have changed what students bring to class and what assignments measure. The resistance an institution encounters is rarely ideological. It is practical. New disclosure conventions add a small amount of work to every assignment. Process-evidence rubrics take longer to grade. Oral assessments require time many faculty do not have. Calibration sessions across instructors require coordination that departments are not staffed to provide.

A faculty member who is already teaching three courses, advising graduate students, serving on two committees, and writing a tenure file does not need a new burden. The honest framing is that AI-aware teaching is real additional work, at least at first, and institutional change management has to start from that recognition.

There is a second source of resistance worth naming. Some faculty believe their existing practice already addresses AI use adequately. They may be right; some courses are naturally AI-resistant by design. They may also be wrong, and the wrong cases tend to surface only when a student fails an oral audit or when a graduate program discovers a thesis chapter cannot be defended. The change-management challenge is to provide enough scaffolding that even faculty who do not see the problem can adopt practices that protect them and their students if it appears.

68.2 Stages of faculty adoption

Faculty distribute across at least four broad groups, each requiring a different kind of support. The grouping draws loosely on the diffusion-of-innovations literature, which has consistently identified distinct adopter categories with characteristic motivations and barriers (Rogers, 2003). The labels matter less than the recognition that one approach will not reach all groups.

Skeptics doubt that AI policy is needed at all in their course. They may believe their existing safeguards (in-person exams, oral components, structured drafts) already prevent the failure modes the policy targets. The most useful intervention is not advocacy but evidence: invite skeptics to participate in calibration sessions where they grade student submissions blind alongside other instructors. The conversations that follow are usually more persuasive than any presentation.

Pragmatists will adopt new practices if shown clear benefit and if the marginal workload is manageable. They want concrete templates, sample rubrics, and ready-made assignment redesigns rather than principles. The most useful intervention is shared resources: a department-level library of AI-aware assignments organized by course type, a small set of validated rubrics, and one or two paragraphs of standard syllabus language. Pragmatists are the largest group on most faculties; the framework's adoption depends on them.

Early adopters are already running ahead of policy. Some have been redesigning assignments for two years; some have published their experiments. Their needs are different: they want recognition, the chance to influence institutional direction, and protection against being penalized for trying. The most useful intervention is to involve them in policy design and to let them serve as resources for the pragmatist group rather than asking them to attend mandatory training built for skeptics.

Overwhelmed faculty want to comply but lack the bandwidth to redesign assignments, learn new tools, or attend additional meetings. They are often contingent faculty, lecturers teaching multiple sections, or faculty in caregiving roles. The most useful intervention is reducing the cost of compliance: providing a pre-made syllabus paragraph, a default rubric, and a single accessible workshop rather than a sequence of trainings.

A successful adoption strategy reaches all four groups with calibrated support, not all with the same training.

68.3 Barriers and how to address them

Several specific barriers recur across institutions adopting AI-aware teaching.

Time cost of process-evidence grading. Reading a draft history, an annotation, and a verification statement takes longer than reading a final submission. The remedy is sampling rather than universal review. A reasonable practice is to grade the final work for everyone and to read process artifacts for a randomly selected subset, plus any submission that raises a concern. Students should be told sampling is in place; this preserves the deterrent effect without the workload of universal audit.

Lack of confidence in oral assessment. Many faculty have not conducted structured oral assessments since their own qualifying exams. They worry about consistency, fairness, and student anxiety. The remedy is training that uses real student work, scripts a small set of standard questions, and provides a simple rubric. Most faculty develop confidence after running three or four oral check-ins; the first one is the hardest.

Concern about fairness to multilingual writers. Detector-based enforcement is documented to misclassify non-native English writing as AI-generated (Liang et al., 2023). Faculty who teach diverse student populations are right to worry about this. The remedy is to abandon detector-only enforcement entirely (section 40) and to rely on process evidence and student explanation. Faculty who understand this shift in evidentiary standard tend to become advocates for the framework rather than skeptics of it.

Workload disparities across course sizes. A 200-student lecture and a 12-person seminar cannot use the same model. Faculty in large courses sometimes feel the framework is designed for small classes and does not address their reality. The remedy is to make the Evidence-of-Thinking Model (section 36) explicit in change-management materials, with concrete examples drawn from large courses.

Tenure-track and lecturer differences. Tenure-track faculty have research obligations that make additional teaching workload genuinely costly to their careers. Contingent faculty and lecturers carry heavier teaching loads and have less time for redesign. Both groups need accommodation, and an institutional approach that ignores these differences will fall hardest on contingent faculty. The remedy is differentiated support: course-release time for redesign work where possible, shared resources that scale across many sections, and recognition that adoption will not look the same across employment categories.

68.4 When the standard playbook is not enough

The four-category adopter model and the barrier remedies above describe what works for most faculty. Departments should also be honest that a small number of faculty will not

move regardless of support level. A tenured professor who refuses to use disclosure standards, declines calibration sessions, or actively grades against process-evidence rubrics in their courses can stall department-wide adoption if their courses are gateway requirements for the major. There is no graceful technical solution to this problem; it is a governance problem and should be handled through department governance rather than through additional faculty-development effort.

What works in these cases is usually some combination of the following: making the policy a department-level requirement rather than an instructor-level recommendation (so noncompliance becomes a governance issue, not a personal-style issue); ensuring that course evaluations and curricular reviews include questions about AI-policy implementation; assigning the noncompliant faculty member to courses where the policy footprint is smaller (electives rather than gateways); and accepting that uniform adoption may take longer than the four-phase plan suggests. Department chairs should not personalize these situations into individual conflict; the goal is institutional adoption, not unanimous instructor agreement, and a small number of holdouts is compatible with effective department-wide practice as long as gateway courses and qualifying milestones are covered.

Institutions should also recognize that a faculty member who appears to be a persistent skeptic may instead be a thoughtful dissenter whose objection is substantive. A serious objection to a particular framework element—for example, that oral assessment as currently implemented produces a particular fairness concern in their student population—may be evidence that the framework needs adaptation rather than that the dissenter needs persuading. Department leadership should distinguish between these cases before treating dissent as obstruction.

68.5 What works

Several practices recur across institutions that have made meaningful progress.

Faculty learning communities that meet regularly to discuss AI in their courses produce more change than any number of workshops. The conversations expose what individual instructors are trying, what is working, and what is not. They also create peer accountability that pure top-down requirements cannot. Five-to-eight faculty meeting monthly is a typical scale.

Paired assignment exchanges between faculty in similar courses help calibrate expectations. Two instructors teaching different sections of calculus exchange one assignment, grade independently, and compare scores. Differences reveal hidden assumptions about what counts as AI-assisted Practice; conversations resolve the differences. This kind of calibration is rare in higher education but produces large effects when adopted.

Calibration sessions where multiple instructors grade the same student submission anonymously, then discuss, build shared standards more effectively than written rubrics alone. The sessions also surface the assessment-validity problems the framework is designed to address (figure 10); faculty often see the polished-output-versus-demonstrated-learning gap clearly when grading a peer's submission.

TA training oriented toward process evidence rather than detector use. Graduate teaching assistants are often the first line of grading in large undergraduate courses; their habits propagate widely. Training should cover the assignment categories (figure 8), the assessment models (table 5), and the verification standards (figure 7).

Leadership modeling by deans, department chairs, and senior faculty matters more than is usually acknowledged. When a department chair publicly redesigns one of their own courses and discusses what worked and what did not, faculty notice. When senior faculty refuse to engage and continue teaching as before, junior faculty hesitate to invest the time in change.

68.6 What does not work

A short list of patterns that institutions try, and that consistently disappoint, deserves explicit naming.

Mandatory training without local input. Centrally designed training that does not reflect what nursing instructors, mathematics instructors, and writing instructors actually face produces compliance without change. Faculty attend the session, sign the attendance sheet, and continue teaching as before.

One-size-fits-all rubrics. A rubric designed for first-year writing will not work for a graduate proof course. A rubric designed for upper-division engineering will not work for studio art. The framework is intentionally flexible across disciplines (figure 7); change-management materials should preserve that flexibility rather than collapsing it.

Detector-only enforcement. Using AI-detection scores as the primary basis for academic-integrity action produces false positives, fairness problems for multilingual writers, and erodes trust in faculty judgment (Cotton et al., 2024). The teaching-and-learning approach to academic integrity, well-developed before AI tools became prevalent, is more robust (Bertram Gallant, 2017). Institutions that rely on detectors create both injustices and litigation risk.

Discipline-blind requirements. Mandates that ignore disciplinary differences place a heavier burden on some faculty than others. A mandate to require draft histories on every assignment is sensible for a writing course and unworkable for a calculus course where the assignment is solving twenty problems. Change-management materials should distinguish requirements that apply universally (disclosure, no-detector-only enforcement) from practices that should be adapted by discipline (specific verification protocols, choice of assessment model).

68.7 Department leadership

The department chair is the most consequential change agent for AI-aware teaching. A dean's policy memo can require new syllabus language; a chair's example shapes daily practice. Three concrete moves are within most chairs' authority.

A one-page departmental compact stating what the department expects of every course (clear AI-use rules in syllabi, a defensible Core Competence assessment, disclosure expectations) and what faculty can expect in return (shared rubric library, calibration meeting time, TA training support). The compact should be brief enough to fit on one page and should be revisited annually.

A one-page departmental compact outlasts a fifty-page policy manual. Faculty read the compact.

A calibration meeting protocol held once a semester, where faculty bring one anonymized submission and the group grades it together. The first such meeting feels awkward; subsequent meetings tend to be the most pedagogically valuable hour of the term. Department chairs who model participation in these meetings make them sustainable.

A teaching-portfolio question for tenure and promotion review. Adding “how does the candidate’s teaching address responsible AI use” as one item among many in a teaching dossier signals that the institution treats this as part of teaching rather than as an extra burden. The signal matters more than the assessment; faculty who know this question will be asked tend to take the work seriously.

68.8 Faculty support resources

Institutional resources that support faculty change should be visible, accessible, and curated. A short, well-organized set is more useful than a long list of links.

Teaching center workshops on AI-aware assignment design, oral assessment, and process-evidence grading. Workshops should be offered both synchronously (with live discussion) and asynchronously (recorded for those who cannot attend). Department-specific sessions, where the workshop draws examples from the participants’ own discipline, are consistently more useful than generic sessions.

Discipline-specific examples curated from the disciplinary overlays in Part II. Faculty in nursing want to see nursing examples; faculty in philosophy want philosophy examples. A small, well-chosen library of example assignments organized by discipline does more work than a large undifferentiated repository.

Assignment redesign clinics where a small group of faculty bring an existing assignment, work through one round of redesign with a facilitator, and leave with a revised draft. The format respects faculty time (one session, one assignment) and produces tangible output (a revised assignment they will use). Clinics scale better than one-on-one consultations.

Shared rubric libraries organized by assignment category (figure 8) and by assessment model (table 5). Rubrics should be lightly curated for quality but not over-standardized. The goal is to give faculty a starting point they can adapt, not to impose uniformity.

These resources become more useful as they accumulate use. A library that contains five well-tested rubrics is more valuable than a library that contains fifty rubrics of unknown quality. Curation is part of the work.

69 From department to institution: how departmental practice can shape institutional policy

Most readers of this book will not personally drive university-wide AI policy. They will, however, shape what their own department does, and that departmental practice can become the template for what eventually becomes institutional policy. This short section describes how the bridge from department to institution typically works in mathematics, science, and engineering settings, and what departmental faculty can do to make that bridge productive rather than frustrating.

69.1 The pattern in practice

Universities rarely adopt comprehensive AI policy in a single top-down decision. Far more often, one or two departments develop coherent practice first, accumulate concrete examples of what works, and then become reference cases when university-wide bodies begin formulating institutional policy. A mathematics department that has six semesters of experience with the Oral-Defense Model, three years of disclosure-block practice across all courses, and a documented academic-integrity procedure is, in effect, providing the institution with a tested case study. When the faculty senate or the academic-integrity committee finally takes up AI policy, they look for examples of what is already working. Departments that have done the work first are positioned to inform that conversation; departments that have done nothing become subject to whatever the institution decides without their input.

The pattern is consistent with how academic policy typically evolves. Curriculum innovations, assessment innovations, and integrity-policy innovations almost always begin in a single department and spread through demonstration. AI-aware practice is no exception.

69.2 What departmental faculty can do to support the bridge

Faculty whose primary influence is at the department level can still make their work institutionally legible. The most useful actions are: (1) document departmental practice in writing, including what was tried, what was kept, and what was discarded, so that the institutional record can refer to it rather than having to reconstruct it; (2) participate in any institutional working groups, faculty senate committees, or teaching-center conversations that touch AI policy, even briefly, because departmental experience is the missing input these bodies usually lack; (3) accept invitations to share departmental practice with other departments, which both helps adjacent departments avoid reinventing the wheel and creates

the cross-departmental relationships that institutional policy formation depends on; (4) name the institutional support that would help departmental practice scale, including the Cost and resource considerations from section 67.5, so that institutional policy can be informed by what departments need rather than only by what institutions can mandate.

69.3 What departments should ask of their institution

Departments doing AI-aware work productively need a small number of things from their institution that institutions are usually willing to provide once they know to provide them: clear privacy and data-governance policy for AI tool use (so departments do not each have to negotiate licensing separately); a centrally supported pool of AI-tool licenses that gives equity-of-access guarantees so departments do not depend on student personal accounts; an academic-integrity procedure at university level that is compatible with department-level evidence standards (specifically, that does not require detector-tool output as the basis for misconduct decisions, given the documented unreliability described in section 61); recognition in the merit-review process of the time faculty spend on AI-aware course redesign; and a teaching center that can offer the workshop and consultation services described in section 65. Departments that ask for these clearly, in writing, with reference to the framework principles in this volume, are far more likely to receive them than departments that ask for general support without specifics.

69.4 The honest limit

Some institutional decisions are out of departmental hands. Most universities will not adopt the framework here in its entirety; some will adopt fragments; some will adopt nothing and eventually develop their own framework that overlaps imperfectly with this one. Departments should not wait for institutional adoption to do good work in their own courses. The departmental framework is designed to be useful regardless of what the institution does, and the institutional framework that follows is offered as guidance for the institutions that wish to use it, not as a precondition for departmental practice. Even in institutions that adopt nothing centrally, a department running coherent practice is providing real value to its students; that value does not depend on whether the policy environment around the department keeps up.

70 Looking forward

A framework written in the present moment will encounter capabilities that did not exist when it was written. The volume should acknowledge this honestly without becoming dated. Some predictions about the next several years will be wrong; the framework's structure is designed to remain useful even when specific implementations need refresh.

70.1 What is likely to change

Several trends seem stable enough to plan around.

AI capability will continue to improve. Recent work on mathematical reasoning suggests that current limitations on proof construction and on multi-step computational reasoning will narrow over the next several years (Ahn et al., 2024). Code generation, multimodal evaluation, and agentic workflows that operate beyond a single chat are all developing faster than university policy cycles.

Tool ubiquity will increase. AI capabilities are moving from separate applications into operating systems, productivity software, and integrated development environments. The distinction between “using AI” and “using a word processor” will blur, and assignment policies that try to specify which tools count will become harder to maintain. Disclosure-and-verification standards generalize better than tool-by-tool restrictions.

Detector reliability will decrease. As AI-generated text becomes harder to distinguish from competent human writing, the underlying signal that detection tools rely on will weaken. Detectors are already failing in two directions: they produce documented false positives against multilingual writers (Liang et al., 2023), and they produce false negatives against students who refine prompts to evoke more literary language. The first direction will increase as detector vendors tune for false-positive control, and the second will increase as student prompting skill becomes ubiquitous. Institutions that built integrity processes around detectors will face mounting reliability problems on both fronts and will need to migrate to the process-based evidence the framework already emphasizes.

Student baseline familiarity will increase. The undergraduates of 2030 will arrive on campus having used AI tools throughout secondary school. Orientation curricula that explain what AI is will become unnecessary; orientation curricula that teach responsible use, verification habits, and disclosure norms will become more important.

70.2 What is likely to stay the same

Other elements appear durable across capability shifts.

The need for independent competence will not disappear. A graduate student who cannot reason without tools is not a researcher; a physician who cannot diagnose without AI is not a physician. The specific topics that count as foundational competence may shift, but the underlying requirement that human expertise must extend beyond tool-mediated procedure remains.

Oral assessment will continue to discriminate well between independent understanding and tool-mediated performance, regardless of tool capability. A student who has not built an underlying argument cannot reconstruct it under questioning, and AI capability does not change that. Oral assessment is expensive but durable.

Process evidence will continue to provide signal. The asymmetry between polishing output and developing judgment (figure 10) does not disappear as AI improves. Polish becomes easier; judgment requires the same effort it always did. Assignments that grade process retain validity longer than assignments that grade only final polish.

Disclosure and verification will become more important, not less. As AI involvement becomes ubiquitous and harder to detect, the burden shifts from external detection to internal disclosure. Verification standards (figure 7) and disclosure norms (Appendix C) protect both the integrity of student work and the trustworthiness of the academic record.

70.3 What this means for the framework

The framework's core elements appear stable across the anticipated changes. The CAT distinction (figure 2) will remain useful regardless of capability shifts because it identifies what an academic course is for, not what a tool can do. The four assignment categories (figure 8) describe the role AI plays in an assignment, which is a curricular choice rather than a tool fact. The three assessment models (table 5) reflect course size and staffing constraints that change slowly. The verification standard generalizes across disciplines and across tool generations.

What will need refresh is the specific content within those structures. Tool lists will need annual updates. Prompt examples will date quickly; the prompt examples in Part I are illustrative rather than canonical. Disclosure templates will evolve as institutional norms develop. Detector references should be removed from policy documents wherever they appear; this volume already takes that position, but many existing institutional policies have not.

The recommendation for departments adopting the framework is to plan annual reviews of implementation specifics (tools, prompts, templates) on a stable foundation of structural commitments (categories, models, verification, disclosure). Annual reviews are part of Phase 4 (section 67.4); the change-management chapter argues that those reviews are also faculty change-management opportunities.

70.4 Modesty about prediction

Specific predictions will be wrong. Some capabilities anticipated above will arrive faster than expected; others will plateau. Some failure modes will be solved; others will appear that no one currently anticipates. The framework is designed to be more durable than any specific tool list, but it is not designed to be permanent. Universities adopting it should expect to revise the implementation details every year and the structural commitments every five years or so.

A useful posture is to treat this volume as a working document rather than a settled answer. Departments that find the framework helpful are encouraged to adapt it, document their adaptations, and share what they learn. The framework improves by use, not by being right

in advance.

71 Evaluation metrics

A university should evaluate whether the framework improves learning and fairness. Useful metrics include:

- percentage of syllabi with clear AI-use rules;
- percentage of courses identifying independent competence assessments;
- student understanding of AI disclosure expectations;
- faculty confidence in AI-aware assignment design;
- TA grading consistency on process-evidence rubrics;
- student performance on no-AI transfer checks;
- quality of student verification statements;
- number and type of academic-integrity cases;
- frequency of detector-only misconduct claims;
- student reports of fairness, access, and privacy clarity;
- adoption of approved tools rather than unsupported tools;
- faculty workload effects.

The purpose of measurement is improvement. If students disclose AI use but cannot verify claims, the student training needs revision. If faculty avoid AI assignments because grading process evidence is too burdensome, the institution should provide better rubrics or TA support. If large courses rely mainly on polished take-home work, assessment validity may need attention.

72 Risk register

Institutional adoption should include a simple risk register. Figure 13 positions the entries below by likelihood and impact.

Table 9: Common risks and institutional responses.

Risk	Why it matters	Recommended response
Over-reliance on AI	Students may produce polished work without durable understanding.	Require independent competence checks and evidence-of-thinking assignments.
Detector dependence	Detection tools may be unreliable and unfair.	Use process evidence, student explanation, and due process.
Privacy breaches	Students or faculty may upload protected data.	Maintain approved-tool lists and clear no-upload rules.
Unequal access	Paid tools can create academic advantage.	Provide institutionally supported tools or equivalent alternatives.
Faculty workload	Process grading and audits can add burden.	Provide templates, TA calibration, sampling methods, and reusable rubrics.
Policy fragmentation	Students receive conflicting rules across courses.	Use shared vocabulary and assignment categories while preserving local flexibility.
False confidence	AI output can make weak understanding look strong.	Include transfer checks, oral explanations, and audit-and-repair tasks.
Underuse of AI literacy	Students may be told what not to do without being taught what to do.	Include explicit student training and AIS assignments.

73 Model institutional policy statement

The following language can be adapted for a university policy page.

The university recognizes that artificial intelligence tools are now part of academic, professional, and creative work. Students should learn to use these tools responsibly while continuing to develop the independent knowledge, skills, and judgment that each course requires.

Each course must state how AI may or may not be used. Instructors should identify which assignments prohibit AI, which allow AI for tutoring or study

support, which allow AI as a collaborator, and which ask students to analyze AI directly. Students are responsible for following course rules, verifying AI-assisted work, protecting private or restricted information, and disclosing AI use when required.

Instructors remain responsible for grading, feedback, and academic-integrity judgments. AI tools may assist instruction, but they should not replace human evaluation. Suspected misuse should be handled through process evidence, student explanation, and established academic-integrity procedures. AI-detector scores should not be used as the sole or primary evidence of misconduct.

74 Evidence base and limits

Institutional AI policy is developing faster than the research base. The shared evidence summarized in section 24 of Part I establishes the learning-science foundation that AI-aware policy must preserve. For institutions specifically, two additional findings are stable enough to guide practice: automated text detectors are not reliable enough to serve as the sole basis for misconduct decisions, and AI tools raise privacy, access, and equity concerns that universities must address directly.

A third stable finding worth naming at the institutional level is the *performance-versus-learning gap*. The OECD's 2026 Digital Education Outlook frames the central evidentiary problem as follows: generative AI tools may enhance the apparent quality of student work—its performance at educational tasks—without improving actual learning, the durable knowledge and skill acquisition the work was supposed to demonstrate (OECD, 2026). Multiple studies converge on this distinction. Institutions that evaluate the success of their AI policies primarily by submitted-artifact quality will systematically miss this gap; institutions that evaluate success by post-AI independent performance, oral-defense outcomes, or transfer-question accuracy will catch it. The implementation plan in this Part is designed to surface the gap rather than mask it; the assessment models and rubrics in Part II are the operational tools.

Other questions remain unsettled. The best balance between no-AI, bounded-AI, and open-AI assessment will vary by field, level, and class size. AI tools will also change. For that reason, this framework should be adopted as a living policy. It gives institutions a shared starting point, not a final answer.

75 Reference materials

The following worksheets, templates, and checklists support institutional adoption of the framework. Each may be adapted for local use; the items are designed to be lifted into faculty-handbook material, syllabus-design guidance, or chair-level operations documents. Global technical references (GMIP, M/A/E framework, scientific writing) appear at the end of the volume.

75.1 Course classification worksheet

Instructors or departments can use this worksheet before the semester begins.

1. What are the three to five core abilities students must develop in this course?
2. Which of those abilities must students demonstrate without AI?
3. Which assignments are NAI, AIT, AIC, or AIS?
4. What evidence will show that students understand submitted work?
5. Is the course best suited to the Oral-Defense Model, Evidence-of-Thinking Model, or Hybrid Model?
6. What privacy restrictions are most relevant to this course?
7. What disclosure format will students use?
8. What student training is needed before the first AI-permitted assignment?
9. What grading support or TA calibration is needed?
10. How will the course be evaluated and revised after the term?

75.2 Syllabus template

AI use in this course

This course uses the following AI-use categories.

NAI assignments. AI use is not allowed. These assignments assess what you can do independently.

AIT assignments. AI may be used for study support, explanation, practice, and hints, but not to produce submitted work.

AIC assignments. AI may be used as a collaborator for brainstorming, revision, debugging, critique, or testing. You must verify the work and disclose your use.

AIS assignments. AI is part of the object of study. You may be asked to analyze, critique, compare, or repair AI output.

For any assignment that permits AI, your submission must follow the disclosure instructions given with the assignment. You are responsible for all submitted work. You should not submit AI output that you cannot explain, verify, or defend.

Do not upload private, restricted, copyrighted, confidential, clinical, human-subject, or secure assessment material into AI systems unless the instructor has explicitly approved that use.

75.3 Disclosure block

Tools used: Name the tools used, if any.

Tool role: Describe what the tools helped with.

Student role: Describe what you wrote, revised, checked, or decided.

Accepted, changed, or rejected: Describe the AI contributions you used, revised, or set aside.

Independent verification: Describe how you checked important claims, sources, data, code, or reasoning.

What I can explain without AI: Name the main parts of the work you can now explain on your own.

75.4 Privacy checklist

Before using AI in a course, ask:

1. Is the tool approved for this type of use?
2. Does the task involve student records, grades, private communications, clinical data, research-subject data, or confidential material?
3. Could student work be stored, reused, or used for training by the vendor?
4. Are students required to create personal accounts?
5. Is there a no-cost option or equivalent alternative?
6. Is the tool accessible to students with disabilities?
7. Are students told what not to upload?
8. Has the instructor provided a privacy-safe alternative when needed?

75.5 Academic-integrity workflow

When AI misuse is suspected, instructors should avoid beginning with accusation. A fair workflow is:

1. Re-read the assignment's AI-use rule.
2. Gather available evidence, including process evidence and comparison to other course work.
3. Invite the student to explain the work.
4. Ask targeted questions about sources, reasoning, method, code, data, or revision choices.
5. If needed, assign a short reassessment tied to the submitted work.

6. Decide whether the issue is misunderstanding, poor documentation, prohibited assistance, or deliberate misrepresentation.
7. Apply the appropriate educational or formal response under institutional policy.

75.6 Department adoption checklist

A department is ready for adoption when it has:

- approved shared AI-use vocabulary;
- provided sample syllabus language;
- identified discipline-specific verification practices;
- developed examples of NAI, AIT, AIC, and AIS assignments;
- stated expectations for independent competence;
- provided a disclosure template;
- created a process for large-course TA calibration;
- clarified privacy and tool-use rules;
- aligned academic-integrity procedures with the institutional framework;
- scheduled periodic review.

75.7 Program evaluation rubric

A university can evaluate implementation using a five-level rubric.

Table 10: Institutional implementation rubric.

Level	Description
1. Fragmented	AI policies vary widely, many syllabi are unclear, and faculty have few shared resources.
2. Emerging	Some departments use clear policies, but vocabulary, privacy rules, and assessment models are inconsistent.
3. Functional	Most courses state AI-use rules, and the institution provides basic templates, training, and academic-integrity guidance.
4. Mature	Departments use shared vocabulary, course-size models, disclosure standards, privacy safeguards, and calibrated assessment practices.

Level	Description
5. Exemplary	AI literacy, independent competence, responsible judgment, equity, privacy, and assessment validity are integrated into regular course design and program review.

75.8 One-page adoption summary

The institutional framework can be summarized in ten commitments.

1. Teach students to use AI without surrendering their own thinking.
2. Require every course to state its AI-use rules.
3. Preserve meaningful assessment of independent competence.
4. Treat AI literacy as an educational outcome.
5. Require verification of AI-assisted work.
6. Normalize brief, factual disclosure.
7. Protect privacy, access, and equity.
8. Avoid detector-only misconduct decisions.
9. Support faculty with templates, training, and TA calibration.
10. Review policies regularly as tools and practices change.

Acknowledgment of source framework

This institutional guide was drafted to accompany the Student Guide and Instructor Guide in the *Learning with AI* package. It generalizes a graduate mathematics assessment framework into a broader university model. The underlying ideas include Core Competence, AI-assisted Practice, Trustworthiness, verification, process evidence, oral grounding, and transparent disclosure.

Publication note

This manuscript is prepared as part of a three-guide package for possible publication and institutional adoption. The source is intentionally conservative in formatting: standard $\text{\LaTeX}$ structures, stable section labels, explicit templates, and minimal custom commands. Final production should use the current SIAM book macro package specified by SIAM Publishing at the time of submission.

A Appendix A: Graduate Mathematical Intelligence Profile

*Two transcripts.
One records what the student can prove.
The other records what they can responsibly verify.*

This appendix presents the *Graduate Mathematical Intelligence Profile* (GMIP), the source framework from which the CAT framework in this book is generalized. GMIP was developed for graduate mathematics departments that expect students to use AI tools as part of serious mathematical work. It is presented here in technical-reference form for departments and instructors who want to adopt or adapt it directly.

GMIP separates what departments want students to become from what tools can produce on their behalf. It treats graduate mathematical ability as a profile of stable competencies, each of which can be assessed under different tool conditions. The framework is intentionally conservative about claims and grading. It assumes that AI models will continue to improve, and it therefore prioritizes assessments that require mathematical judgment, verification habits, coherent reasoning under questioning, and ethical disclosure of tool use.

A.1 Two-transcript design

GMIP organizes evaluation into two linked transcripts.

Mathematics Transcript. What the student can do with no assistance or with limited assistance. Primary instruments are no-AI written exams, oral examinations at the blackboard, and short proof or concept tasks under closed-resource conditions.

AI-Integrated Transcript. What the student can do when using tools responsibly to accelerate and deepen mathematical work. Primary instruments are open-tool homework with process artifacts, AI-critique and error-detection tasks, comparative human–AI problem solving, and AI-assisted research projects.

A grade can be computed from these transcripts, but the larger goal is diagnostic: GMIP is designed to tell a student and their advisor what kind of mathematician the student is becoming and what to strengthen next.

A.2 Five core design principles

GMIP rests on five design principles that distinguish it from conventional mathematical assessment.

1. **Tool-conditional assessment.** Evaluate the same competency in multiple tool regimes (no-AI, bounded-AI, open-AI). The same student may show strength in one regime and weakness in another, and that pattern is itself diagnostic.
2. **Evidence-first grading.** Require artifacts that make reasoning inspectable: scratch work, proof sketches, checklists, verification logs. A correct final answer does not establish mathematical understanding if no reasoning evidence is visible.
3. **Verification over fluency.** Reward students who can detect and repair errors, including errors produced by tools. AI models can generate fluent mathematics that contains subtle flaws. The student who notices and fixes those flaws is demonstrating exactly the skill the discipline values.
4. **Process integrity.** Require transparent disclosure of tool usage and enforce clear boundaries on what counts as the student’s own contribution. The boundaries should be stated in writing, applied consistently across courses, and discussed openly with students.
5. **Oral grounding.** Reserve a meaningful portion of the grade for live reasoning, where understanding cannot be outsourced. Oral examinations remain the most future-proof assessment because they probe the present-tense thinking of the student in front of the examiner.

A.3 The eight competency domains

GMIP evaluates graduate mathematical intelligence through eight competency domains (figure 14). Departments can adopt all eight or start with a subset. Each domain includes characteristic abilities, observable evidence, and suggested assessments.

D1. Conceptual Understanding and Definitions. Students can state definitions precisely, explain motivations, give nontrivial examples and counterexamples, and identify what changes when hypotheses change. Typical evidence includes definition recall and reformulation in plain language, minimal counterexample construction under relaxed assumptions, and dependency tracing showing which steps use which hypotheses.

D2. Proof Construction and Proof Repair. Students can write complete proofs, choose strategies, and repair broken arguments. They can also certify whether a tool-produced proof is correct. Typical evidence includes an independent proof outline written before tool use, an annotated proof audit that identifies gaps and fixes them with cited lemmas, and error-localization tasks that ask the student to find the first invalid inference in a flawed argument.

D3. Problem Solving and Technique Selection. Students can decide what to try next, not just execute a known routine. They can explain why a technique is appropriate and when

it fails. Typical evidence includes strategy memos describing a short plan with alternatives and failure modes, timed in-class problems that require branching choices, and post-solution reflection on what generalizes and what does not.

D4. Mathematical Communication. Students can communicate with precision and audience awareness, including writing that is readable and honest about uncertainty. Typical evidence includes short expository notes with a guiding narrative, board proofs with clean structure and appropriate signposting, and peer-review work that gives useful mathematical feedback on others' arguments.

D5. Modeling and Interpretation. Where the program emphasizes applied or computational mathematics, students can translate between assumptions, mathematical form, and interpretation. They can state what a model proves, what it suggests, and what it cannot justify. Typical evidence includes assumption-to-consequence maps, sensitivity and robustness reasoning, and clear separation between theorem, computation, and interpretation in written work.

D6. Computational Mathematics and Experimentation. Students can use computation to explore, test conjectures, and validate intermediate steps. They can reason about numerical error and reproducibility. Typical evidence includes reproducible computational notebooks or scripts, sanity checks on edge cases with alternative implementations, and explicit complexity and stability discussions.

D7. AI Tool Proficiency for Mathematical Work. Students can use AI tools to accelerate work while maintaining control over correctness. They can prompt well, constrain outputs, and integrate tool output into rigorous work. Typical evidence includes prompt-and-response logs with rationale, iterative refinement that steers a tool toward a valid proof or counterexample, and tool triangulation that compares results across methods and resolves conflicts.

D8. Research Integrity, Ethics, and Professional Judgment. Students disclose tool use honestly, respect attribution norms, and maintain rigorous standards. They avoid laundering tool output as personal work. Typical evidence includes transparent tool-disclosure statements, attribution and citation discipline for external sources, and consistent integrity under questioning, where the student's account of how the work was produced does not change between contexts.

A.4 Assessment architecture

GMIP uses a small number of recurring assessment types. Each type is designed to remain informative even as AI tools improve.

- A. In-class exams in two modes.** Use short, frequent in-class exams in two modes. *No-AI mode*: closed network, no tools, focused on core reasoning and definitions. *Bounded-AI mode*: a controlled tool setting, such as one approved model, limited time, and required disclosure, focused on verification and judgment. Problems should be chosen so that a correct solution depends on one or two key decisions, which makes it possible to grade reasoning quality rather than output volume.
- B. Take-home work with strict process.** Allow open tool use on homework and take-home exams, but require process artifacts. A strong default is to grade each submission in two layers: mathematical correctness and exposition, and process integrity and verification. Required artifacts include a short “plan first” section written before tool use, a tool log of prompts and outputs with notes on what was accepted or rejected, a verification checklist documenting independent checks, and a brief disclosure statement.
- C. Individual oral examinations.** Oral exams are the most future-proof component because they probe live understanding. A recommended structure is a 30 to 60 minute session with two segments: blackboard reasoning on definitions, short proofs, and probing questions that test hypothesis-awareness; and audit-and-repair on a short argument that contains subtle flaws which the student must identify and fix.
- D. Research-style performance tasks.** At least once per term, assign a task that resembles real mathematical work: exploring a conjecture, building a small theory, or producing a carefully written note. These tasks reward disciplined verification, good exposition, and honest handling of uncertainty. Specific task templates include a conjecture laboratory (computational exploration leading to a conjecture, a proof attempt, and a failure analysis), a counterexample clinic (generating families of examples that stress-test a statement), and a mini-survey (synthesizing a small literature slice with correct attribution and clean exposition).

A.5 Grading scheme

GMIP supports a grade that is both legible to students and useful to faculty. A simple adoption-friendly scheme is a weighted sum of two scores.

Mathematical Mastery (MM). Performance in no-AI and oral components. Emphasis on reasoning without external scaffolding.

AI-Integrated Mathematical Practice (AIMP). Performance in bounded-AI exams, take-home work, and research-style tasks. Emphasis on verification and responsible tool use.

A typical department-level choice is to weight MM slightly higher than AIMP early in the program and to converge toward parity by the end of the second year. Departments can adjust weights to match program goals.

The same five rubric anchors apply to all components and make grades comparable across courses: *correctness* (the mathematics is correct, complete, and appropriately justified); *reasoning visibility* (the work shows the path, not just the destination); *verification discipline* (the student checks claims and catches errors, including tool errors); *communication* (structure, definitions, and reader guidance are strong); and *integrity* (tool use is disclosed and consistent with course rules).

A.6 Future-proofing stress tests

Before adopting a new assessment, GMIP suggests asking whether it still works if tools become much stronger. A recent survey of LLM mathematical reasoning capabilities shows steady improvements on structured tasks alongside persistent weaknesses on proof construction, judgment, and adaptation to new contexts (Ahn et al., 2024). The four stress tests below are designed to remain informative against this moving target.

Outperform test. If a model can produce a polished solution, does the task still measure judgment, verification, or understanding? If not, the task may need redesign.

Explanation test. Can a student defend the work at the board under probing questions? If the assessment cannot accommodate live questioning at all, it should be paired with another that can.

Error test. Does the assessment include opportunities to detect and repair subtle mistakes? Pure-output assessments lose validity faster than audit-and-repair assessments.

Transfer test. Does success require adapting ideas to a new but nearby context, not just reproducing patterns? Tasks that ask only for pattern reproduction become weaker as model capabilities grow.

An assessment that passes all four tests is likely to remain informative for some time. An assessment that fails one or more should be reviewed and either redesigned or weighted less heavily.

A.7 Operational policy

Tool disclosure standard. Every graded artifact that allows tool use should include a short disclosure block. The disclosure should be factual and brief. It should state the tools used (model and version when known), what the tool contributed (ideas, proof sketches, code, editing), and what the student verified independently. The detailed disclosure block template appears in Part I of this book; the same template is appropriate for graduate work.

Minimal integrity boundaries. GMIP works best when departments enforce a few bright lines.

- Students may not submit tool output as a proof without an audit that identifies why each step is valid.
- If a student cannot explain an argument orally, it does not count as mastery.
- All external sources, whether human or tool-mediated, must be cited or acknowledged.
- Instructors should provide a short list of allowed and disallowed tool uses for each assessment type.

Calibration and consistency. To keep grading consistent across courses, departments can run a short calibration each term: instructors grade a small sample using the shared anchors and discuss discrepancies. The goal is not identical grading, but stable interpretation of what mastery and responsible practice mean. The same calibration practice is recommended for the broader CAT framework in Part II.

A.8 Templates

Template A: Tool disclosure block.

Tools used: (model name and version)

Tool role: (what the tool helped with)

Independent verification: (how the student checked the result)

Template B: Verification checklist (short form).

- ☐ I can restate the problem in my own words and identify the main obstacle.
- ☐ I checked each nontrivial inference, or I cited a lemma or theorem.
- ☐ I tested at least one edge case or limiting case.
- ☐ I produced a second derivation or an alternative argument for one key step.
- ☐ I can explain the solution aloud without reading.

Template C: Oral examination question bank outline.

- *Definitions and examples.* Core objects of the course and their motivating examples.
- *Short proofs.* Five-to-ten minute arguments built from standard tools.
- *Hypothesis stress tests.* Modify assumptions and ask the student to predict what breaks.

- *Audit-and-repair.* Locate and fix a flaw in a short proof.
- *Synthesis.* Connect two results from the course in a new way.

GMIP is offered as a starting point. Departments are expected to adapt it to their own program structures, comprehensive examination practices, and research training requirements. The companion document in Appendix B provides a more compact assessment-policy formulation that some departments may find easier to adopt directly.

B Appendix B: AI-Integrated Graduate Mathematics Assessment Framework

*Three orthogonal axes.
A weak score on one cannot be paid off by a strong score on another.*

This appendix presents the *AI-Integrated Graduate Mathematics Assessment Framework*, a department-level policy formulation that complements the GMIP technical reference in Appendix A. Where GMIP describes a profile of competencies and assessments, this framework provides a more compact policy structure suitable for adoption by a department or graduate program. The two are designed to be used together. The original M/A/E terminology of this framework is preserved here. The Foreword explains how M/A/E maps onto the CAT framework used in the main text of this book.

B.1 Purpose and rationale

This framework offers a unified departmental approach to integrating AI into graduate mathematics instruction and assessment. The goal is to ensure that students develop deep mathematical understanding, advanced AI-assisted research skills, and strong ethical and epistemic judgment.

The framework recognizes that AI is now a permanent component of mathematical practice. Assessment must therefore distinguish between human mathematical competence and AI-augmented competence while rewarding responsible and transparent use of AI tools. A department that does not make this distinction risks two failures simultaneously: it cannot reliably assess what its students understand, and it does not prepare them for professional and research environments in which AI tools are routine.

B.2 Core principles

The framework rests on five principles.

1. **Mathematical depth.** Students must demonstrate genuine conceptual understanding independent of AI.
2. **AI literacy.** Students must learn to use AI as a mathematical instrument, not a substitute for reasoning.
3. **Epistemic responsibility.** Students must verify, critique, and transparently document AI contributions.
4. **Process-oriented assessment.** Evaluation emphasizes reasoning and methodology, not only final answers.

5. **Equity and clarity.** Expectations for AI use must be explicit, consistent, and fair across courses.

B.3 The M/A/E competency model

Graduate mathematical proficiency is evaluated along three orthogonal axes.

M — Mathematical Understanding. What students can do without AI: definitions, proof construction, conceptual reasoning, problem-solving fluency.

A — AI-Augmented Mathematical Skill. What students can do with AI as a tool: prompting, verification, integration of AI output with human reasoning, error detection in AI-generated mathematics.

E — Epistemic and Ethical Responsibility. How students disclose, attribute, verify, and take ownership of AI-assisted work.

Each course should assess all three dimensions using multiple instruments. A student must meet minimum standards in M to pass, regardless of performance in A or E. This minimum-M requirement is what prevents AI proficiency from compensating for weak mathematical foundations.

B.4 Assessment structure

Recommended weighting, adjustable by course type, is as follows.

Component	Weight
Mathematical Understanding (M)	40–50%
AI-Augmented Skill (A)	30–40%
Epistemic and Ethical Responsibility (E)	15–25%

Departments may adapt these weights to their program goals, but the underlying principle is non-negotiable: AI proficiency cannot compensate for weak mathematical foundations. A student who excels in A and E but fails to meet the minimum M standard has not met the program’s mathematical requirements.

B.5 Assessment instruments

No-AI assessments (M).

- In-class written exams without AI tools.
- Short proof and concept-explanation tasks.
- Oral examinations at the blackboard.

AI-integrated assessments (A).

- AI-assisted homework and take-home exams.
- AI-critique and error-detection tasks.
- Human–AI comparative problem solving.
- AI-assisted mini research projects.

Epistemic and ethical assessments (E).

- AI-usage documentation and prompt logs.
- Reflective methodological reports.
- Error analysis and verification tasks.
- Attribution and reproducibility statements.

B.6 Oral examination framework

The oral exam is the keystone assessment in this framework and should include three components.

1. **Pure reasoning.** The student proves or explains a mathematical result without AI.
2. **AI critique.** The student analyzes an AI-generated solution, identifying its strengths and weaknesses.
3. **Problem reframing.** The student formalizes an ambiguous problem and discusses how AI could assist or mislead in addressing it.

This three-component structure distinguishes deep understanding from surface-level AI reliance. A student who passes the first component but stumbles on the second has demonstrated mathematical knowledge without AI literacy. A student who handles the second smoothly but cannot complete the first has demonstrated tool fluency without mathematical depth. The framework requires both.

B.7 Documentation requirements

For AI-integrated assignments, students must submit the following.

- Description of AI tools used (name and version).
- Prompt evolution history.
- AI-generated outputs that influenced the submission.
- Independent verification and corrections.
- Final mathematical reasoning, in the student's own words.

This documentation is itself an evaluable artifact. Sloppy or incomplete documentation lowers the E score even when the mathematics is correct. Clean documentation that traces the student's reasoning through tool interactions and back to verified mathematical conclusions is what professional research practice requires.

B.8 Grading rubrics

Departments should adopt explicit rubrics for each competency axis. The framework recommends the following criteria.

Mathematical Understanding (M).

- Conceptual clarity.
- Logical rigor.
- Proof structure.
- Independent reasoning.

AI-Augmented Skill (A).

- Quality of prompts and decomposition.
- Ability to detect and correct AI errors.
- Integration of AI output with human reasoning.

Epistemic and Ethical Responsibility (E).

- Transparency of AI use.
- Verification rigor.

- Intellectual honesty.
- Awareness of AI limitations.

B.9 Safeguards against AI substitution

To prevent AI from replacing genuine learning, the framework recommends four institutional safeguards.

- A minimum M-threshold required to pass courses.
- A significant proportion of no-AI assessments throughout the program.
- Oral examinations for all graduate students at appropriate program milestones.
- Process-based grading rather than output-only grading.

The threshold rule is the most important. Without it, students can in principle compensate for weak mathematical understanding through fluent AI use, which defeats the purpose of graduate training. The threshold need not be high to be effective; what matters is that it exists and is enforced consistently.

Although direct empirical evidence on graduate-mathematics adoption of these specific safeguards is not yet available, course-redesign experience in adjacent STEM contexts is broadly consistent with the recommended structure. Kosar et al. (2024) report a computer-science course redesigned around paper-and-pencil mid-term exams focused on conceptual understanding, oral defenses of submitted code during lab time, and home assignments that could not be directly completed by an AI chatbot; in that setting, student groups with and without GenAI access produced similar learning outcomes. The redesigned course operationalizes three of the four safeguards above (no-AI assessment, oral examination, process-resistant assignment design); its results provide the most direct external evidence available that the safeguard structure works in practice. Departments adopting the framework here may find the Kosar redesign a useful operational reference, with appropriate adaptation for graduate-mathematics content and assessment instruments.

B.10 Implementation phases

A department adopting this framework should plan a phased rollout.

Phase 1 — Pilot. Adopt the framework in selected graduate courses. Test the assessment instruments, calibrate grading among the instructors involved, and collect feedback from students.

Phase 2 — Faculty workshops. Run workshops on AI-integrated assessment for the broader graduate faculty. Use sample student work from Phase 1 as calibration material.

Phase 3 — Department-wide adoption. Roll out the framework across all graduate courses and qualifying examinations. Establish a periodic review cycle, since AI tools and student practices will continue to evolve.

Faculty are encouraged to adapt the framework to local context while preserving its core epistemic commitments. The minimum M-threshold and the requirement for transparent AI disclosure should not be relaxed even as other elements of the framework are localized.

B.11 Expected outcomes

A graduate program that adopts this framework should expect to produce students who possess deep mathematical understanding, who can use AI as a sophisticated research tool rather than a crutch, who demonstrate strong epistemic judgment when AI output is involved, and who uphold ethical standards in AI-assisted mathematics. The framework is not a guarantee of these outcomes; it is a structure that supports them.

The framework is not designed to produce uniformity across departments. Different graduate programs will reasonably weight M, A, and E differently based on their research culture and student goals. The framework provides shared vocabulary and a common assessment architecture; the specifics of weighting, rubric stringency, and assessment frequency are local decisions.

C Appendix C: Best Practices for Ethical AI-Assisted Scientific Writing

If a reviewer asks, the author should have an answer that does not begin with the tool.

This appendix gives practical guidance for scientists, including graduate students and faculty researchers, who use AI tools when preparing research manuscripts. The goal is to maintain scientific integrity, ensure transparency, and preserve human accountability while benefiting from AI assistance.

The guidance here applies to manuscript preparation, code development associated with research, literature review, and grant writing. The emphasis is on STEM and computational research where equations, code, and numerical results dominate verification practice; humanities and qualitative-research writers will find the underlying principles (verification, disclosure, responsibility) transferable but should also consult the discipline-specific guidance for humanities, writing, and creative disciplines in Part II (section 42). It complements the verification and disclosure practices described in Part I for student work and the GMIP framework in Appendix A for graduate-level mathematical work. Researchers preparing manuscripts for journal submission should also consult their target journal’s specific AI policy, which may impose requirements beyond those described here.

C.1 Core principles

Human responsibility is non-negotiable. Every scientific result, analysis, interpretation, and conclusion in a research manuscript must be something the author can defend without reference to AI assistance. The author understands every equation and derivation. The author can explain the methodology to reviewers. The author recognizes the relevant literature and scientific context. The author made the intellectual choices, even if AI helped explore options.

A practical test makes this concrete: if a reviewer asks, “Why did you choose this approach?” the author should have a genuine answer rooted in scientific understanding rather than a procedural answer about what the tool suggested. If the only available answer is a tool-mediated one, the work is not ready for submission.

AI systems are tools, not collaborators. AI tools may assist with drafting and refining prose, exploring analytical approaches that the author then verifies, code prototyping that the author then tests and validates, and literature searching that the author then reads and evaluates. AI tools cannot generate original scientific results, make intellectual decisions about methodology, take responsibility for correctness, or qualify for authorship.

This distinction matters when journals ask whether AI tools should be listed as authors. The professional consensus across publishing bodies is that they should not, because authorship

implies responsibility, and tools cannot bear responsibility.

Transparency proportional to impact. The level of AI disclosure should match the level of AI involvement. Minor uses, such as grammar checking and formatting, warrant a brief acknowledgment. Moderate uses, such as drafting text or assisting with code, warrant a clearer description in the acknowledgments. Substantial uses, such as significant drafting, analysis, or exploration, warrant a detailed description in both the acknowledgments and the methods section.

C.2 The verification cornerstone

The requirement that all AI-assisted content be independently verified is what transforms AI assistance into a contribution the author can defend before peer review and disciplinary critique. Without verification, AI-assisted text is at best a draft and at worst a hazard.

Verification works better when the AI's output is calibrated to the author's actual technical level and goals. Researchers who configure their AI accounts with persistent preferences specifying their field, current projects, and uncertainty-flagging requirements receive output that is easier to verify because the AI has been told what kind of confidence claims to make and how to mark unfamiliar territory. The student-facing version of this configuration practice is described in section 2 of Part I (with the full configuration template at section 26.1); the principles transfer directly to research use, with the technical-work accuracy guard (showing work, labeling claims as confident or tentative) becoming particularly valuable for research-grade writing.

A practical verification workflow has five steps.

1. AI assists with draft content or analysis.
2. The author reviews the output critically: does this make sense?
3. The author verifies the output technically: is this correct?
4. The author confirms meaning: does this say what I intend?
5. Once verified, the content represents the author's verified work.

Each step is necessary. Reviewing for plausibility (step 2) is not the same as verifying technical correctness (step 3). A passage that reads smoothly and reaches a familiar conclusion can still contain errors that only careful technical inspection will surface. A passage that is technically correct can still misrepresent the author's intended meaning if the author has not examined it for fidelity to the underlying scientific argument.

C.3 Disclosure guidelines

Where to disclose.

- Acknowledgments (the primary location for AI disclosure).
- Cover letter, when AI use was substantial enough to merit editorial attention.
- Methods section, when AI assisted with analysis or modeling.
- Data availability or supplementary materials, when AI assisted with code that affects reproducibility.

Sample disclosure language. *Minimal* (grammar checking, prose refinement only):

The authors used *[AI tool]* for grammar checking and prose refinement. All scientific content was developed independently by the authors.

Standard (drafting and revising portions of exposition):

During the preparation of this article, the authors used *[AI tool and version]* for drafting and revising portions of the exposition. All scientific results and interpretations were independently verified and are the sole responsibility of the authors.

Comprehensive (substantial AI involvement in drafting, code prototyping, or analysis exploration):

This manuscript was prepared with AI assistance using *[tool and version]*. Initial drafts of *[sections]* were generated with AI assistance and substantially revised by the authors. Code was prototyped with AI assistance and independently tested. All scientific results were developed by the authors and independently verified. The authors take full responsibility for all content.

The right disclosure level is the one that accurately represents what AI contributed to the work. Under-disclosure is unethical; over-disclosure can imply more AI involvement than actually occurred and may misrepresent the author’s contribution in the other direction.

Fairness considerations in detector-based screening. A growing number of journals use AI-text detectors during submission screening. These tools have documented limitations that disproportionately affect non-native-English writers. In an analysis of ICLR 2023 papers, those whose first authors were from non-English-native countries showed systematically lower

text perplexity than papers by English-native first authors (Liang et al., 2023); the same study reported a 61.3% average false-positive rate when those same detectors were evaluated on TOEFL essays by non-native English speakers. The mechanism is intrinsic to detector design: most current detectors rely on text perplexity, and non-native writers exhibit lower lexical and syntactic variability, which detectors then misread as AI-generation. Authors writing in English as a second language should expect higher false-positive rates from journal screening tools and should be prepared to provide process evidence—drafts, revision history, or prompt logs—if questioned. The detailed-disclosure standards in this appendix support exactly that kind of evidence-based response. The broader implication is that the research-publishing community should treat detector output as one signal among several rather than as a basis for editorial decisions, and authors who can document their writing process protect themselves against detector-based misclassification.

C.4 Common pitfalls

Citation hallucination. AI tools sometimes suggest citations that do not exist or that exist but do not support the claim being made. Empirical work documents substantial fabrication rates in AI-generated bibliographies, often in the range of tens of percent depending on the model version and topic specificity, with errors common even in citations that point to real papers (Walters and Wilder, 2023). The risk is high enough that researchers should treat every AI-suggested citation as unverified until checked. The remedy is straightforward: verify that every cited source exists, read every cited paper, and never cite an unexamined source. This is not new advice for scholarly work; AI simply makes it more important by increasing the volume of plausible-looking but incorrect citations a researcher may encounter.

Accepting unverified analysis. AI presents analysis confidently. The temptation to accept it without verification is real, especially when the analysis appears to confirm a hypothesis the researcher already favors. This pattern is documented in the educational research literature on AI use in learning, where it has been characterized as metacognitive laziness: individuals who use generative AI to complete cognitive tasks tend to disengage from monitoring whether the task has actually been completed correctly (Fan et al., 2025). The empirical work focuses on student learning, but the cognitive mechanism applies to research writing as well; confident output exerts a pull toward acceptance even when verification is the explicit norm. The remedy is to work through every step independently, to test code with known inputs, and to refuse to include any analysis the researcher cannot follow on their own.

Meaning drift in prose. AI revision can subtly change the meaning of scientific claims. A sentence that originally said a result “suggests” a conclusion may be revised to say it “shows” or “proves” that conclusion, an inflation of certainty that the underlying evidence may not support. The remedy is to review every paragraph after AI revision, check that hedging language matches the author’s actual confidence, and rewrite any claim the author does not

recognize as their own.

Invisible code errors. AI-generated code that runs successfully can still produce incorrect results. Compilation success and execution success are not correctness. The remedy is to test with known inputs whose correct outputs the researcher can verify independently, to check boundary conditions, and to compare results against analytical solutions or established benchmarks where they exist.

C.5 Pre-submission checklist

The following checklist, completed before manuscript submission, encodes the standards in this appendix.

- ☐ I understand every equation and analysis in this paper.
- ☐ I can defend every claim without reference to AI.
- ☐ I have verified all AI-assisted content independently.
- ☐ I have read all cited papers.
- ☐ I have tested all code and can reproduce key results.
- ☐ My disclosure accurately reflects AI use in this work.
- ☐ I have checked the target journal’s AI policy and complied with its specific requirements.

A manuscript that passes this checklist is not necessarily ready for submission, but a manuscript that fails any item is not ready. The failures should be addressed before any further submission preparation.

C.6 Verification standards by AI contribution type

The following table summarizes what verification looks like for each common type of AI contribution to a research manuscript.

AI contribution	Verification standard	What the author must do
Prose drafting	Read critically	Ensure the text says what the author means; rewrite anything that drifts in meaning or certainty.
Analytical approach	Work through independently	Verify every step of the analysis on the author's own; do not include analysis the author cannot reconstruct.
Code generation	Test extensively	Validate correctness through known-input testing, boundary checks, and benchmark comparisons; do not assume execution implies correctness.
Citation suggestions	Read the sources	Confirm each cited source exists and supports the claim; treat unverified citations as unsuitable for inclusion.
Numerical results	Reproduce independently	Check against known values, analytical solutions, or independent reimplementations.

C.7 Summary

AI-assisted scientific writing is ethically sound when four conditions hold. The researcher's intellectual contribution is genuine: AI assists, the researcher creates. Verification is thorough: every AI contribution is independently checked. Disclosure is proportional: transparency matches AI involvement. Responsibility is the researcher's: the researcher can defend everything in the manuscript without reference to AI.

These four conditions are easier to satisfy when the researcher has already developed strong verification habits in graduate training. The framework in Appendices A and B is one path to those habits. Researchers whose graduate training predates the AI era can develop the same habits through deliberate practice and through the verification workflow described above.

The standards in this appendix are general. Specific journals, funding agencies, and professional societies may impose additional requirements. Researchers should treat this appendix as a baseline rather than a ceiling.

D Appendix D: Math-Specific Quick Reference

*In mathematics, the question is rarely whether the answer is right.
It is whether the student understands why.*

This appendix is a one-stop reference for using the framework in mathematics courses. It consolidates the practical patterns scattered across Parts I and II: student-side AI study habits for proof-based and computational mathematics, instructor-side category mixes and assessment artifacts for typical math course types, and a few fast design rules for instructors under time pressure. The student-side material is in second person; the instructor-side material is in third-person summary.

D.1 Student use in mathematics

The student-side principle is simple. AI should support mathematical thinking, not replace it. In the language of the book, students should preserve Core Competence, use AI for AI-assisted Practice, and show Trustworthiness through verification, disclosure, and explainability.

Proof-based mathematics. Proof-based courses include discrete mathematics, transition-to-proof, linear algebra with proofs, abstract algebra, real analysis, topology, and most graduate theory courses. In these courses, the main danger is accepting a polished proof that the student cannot reconstruct, defend, or adapt under questioning.

A compact proof-based workflow is:

1. State the theorem and rewrite the relevant definitions from memory.
2. Sketch a plan before consulting AI.
3. Use AI only for hints, audits, or comparison of approaches, not for full proofs.
4. Rebuild the proof independently in your own words, ideally from memory or at the board.
5. Check whether the argument survives a small variation of the statement.

Recommended habits during the cycle: try the problem first, even briefly; ask AI for a hint, a simpler analogous problem, or an audit of a proof attempt rather than a complete proof; verify every major step by checking where each hypothesis is used; use AI to generate examples, non-examples, and borderline cases for definitions, especially in analysis and algebra where understanding often depends on seeing what fails when a hypothesis is dropped.

Computational and applied mathematics. Computational and applied courses include calculus, differential equations, numerical analysis, optimization, modeling, scientific computing, statistics, and data-oriented math courses. In these courses, the main danger is

accepting output that looks plausible because the algebra works or the code runs, even when the method, assumptions, or interpretation are wrong.

A compact computational workflow is:

1. Write down the model, formula, or plan first.
2. Use AI for hints, diagnostics, or method comparison.
3. Solve or code independently.
4. Run verification checks on cases where the answer is known or constrained, using the verification kit in section 13.1.
5. Be prepared to explain the method, not just the output.

Recommended habits: identify variables, assumptions, equations, or algorithmic goals before asking AI for help; use AI to suggest solution strategies, debug code by diagnostic questions, or generate test cases; verify quantitative work with the math verification kit (special case, limiting case, units, known formula, independent route); test edge cases and reproducibility for code; ask AI to explain why a method is appropriate, what it assumes, and when it can fail.

Strong AI prompt patterns for mathematics. The strongest AI prompts in mathematics preserve the student’s active role. Table 11 lists patterns that work across course types.

Table 11: Strong AI prompt patterns for mathematical work. The pattern is consistent across goals: ask AI to help you do the thinking, not to replace it.

Goal	Strong prompt pattern
Start a proof	Give me one hint about the main idea, but do not prove it.
Check a proof	Audit my proof attempt and identify the first step that needs more justification.
Learn a definition	Give me an example, a non-example, and a borderline case of this definition.
Debug code	Ask me diagnostic questions until I find the bug.
Prepare for oral defense	Ask me to explain each step of this argument until you find what I do not understand.
Verify a method	What assumptions does this method use, and how would I test whether they hold here?

D.2 Instructor use in mathematics

For instructors, the main design problem is to separate what students must demonstrate independently from what they may do responsibly with AI assistance. The NAI, AIT, AIC, and AIS assignment categories from section 33 and the Oral-Defense, Evidence-of-Thinking, and Hybrid assessment models already provide that structure.

Typical course patterns. Table 12 gives practical default mixes for common mathematics course types. These are starting points, not rigid prescriptions.

Assessment artifacts that work especially well in mathematics. The strongest mathematics assessments make reasoning visible rather than grading only final polish. Across course types, the following artifacts are especially effective:

- *First-attempt plans*: a 2–3 line proof idea, setup, or strategy written before AI use.
- *Homework-to-quiz checks*: AI-allowed homework paired with a short no-AI in-class question on the same underlying idea.
- *Audit-and-repair tasks*: students critique and repair flawed AI-generated proofs, derivations, or code.
- *Verification notes*: short statements showing how a solution, derivation, or computation was checked.
- *Oral defenses*: brief conversations in which students explain a proof, method, or design choice in their own words.
- *Transfer questions*: nearby variants of a problem that reveal whether students learned the structure rather than memorized a polished answer.

Fast design rules for mathematics instructors. When redesign time is limited, three moves usually provide the largest benefit. The same three appear in the implementation workflow in section 45; they are repeated here for the math reader.

1. Label each assignment clearly as NAI, AIT, AIC, or AIS.
2. Add at least one meaningful independent-competence check to each unit.
3. Require one small process artifact when AI is allowed: a plan, one representative AI exchange, and a verification note are often enough.

For small proof-based classes, the Oral-Defense Model is often the most robust design because it allows substantial AI use while preserving standards of proof and explanation. For large service courses, the Evidence-of-Thinking Model is usually more feasible: homework-to-quiz

checks, random oral audits, brief process logs, and AI-critique items scale better than full oral examinations.

In mathematics, AI is most helpful when it acts as a tutor, critic, or diagnostic partner rather than a substitute author. Students should use it to get unstuck, test understanding, generate examples, audit arguments, and debug computations. Instructors should design assessments that preserve independent competence while making verification, process, and judgment visible.

References

- Janice Ahn, Rishu Verma, Renze Lou, Di Liu, Rui Zhang, and Wenpeng Yin. Large language models for mathematical reasoning: Progresses and challenges. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics: Student Research Workshop*, pages 225–237, St. Julian’s, Malta, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.eacl-srw.17.
- Susan A. Ambrose, Michael W. Bridges, Michele DiPietro, Marsha C. Lovett, and Marie K. Norman. *How Learning Works: Seven Research-Based Principles for Smart Teaching*. Jossey-Bass, San Francisco, 2010. ISBN 9780470484104.
- American Association of University Professors. Artificial intelligence and academic professions. AAUP Topical Report, 2025a. URL <https://www.aaup.org/reports-publications/aaup-policies-reports/topical-reports/artificial-intelligence-and-academic>. AAUP topical report on the labor and professional implications of AI in higher education. References AAUP Policy Documents and Reports 12th edition (Johns Hopkins University Press, 2025). Wayback Machine snapshot to be added during VER pass.
- American Association of University Professors. Artificial intelligence and academic professions. Technical report, American Association of University Professors, Washington, DC, 2025b. URL <https://www.aaup.org/sites/default/files/2025-07/Artificial-Intelligence-and-Academic-Professions-Bulletin.pdf>. Report of the ad hoc Committee on Artificial Intelligence and Academic Professions.
- Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakcı, and Rei Mariman. Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*, 122(26):e2422633122, 2025. doi: 10.1073/pnas.2422633122.
- Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT ’21)*, pages 610–623, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 978-1-4503-8309-7. doi: 10.1145/3442188.3445922.
- Tricia Bertram Gallant. *Academic Integrity in the Twenty-First Century: A Teaching and Learning Imperative*. ASHE Higher Education Report 33(5). Jossey-Bass, San Francisco, 2008. ISBN 9780470373668.
- Tricia Bertram Gallant. Academic integrity as a teaching and learning issue: From theory to practice. *Theory Into Practice*, 56(2):88–94, 2017. doi: 10.1080/00405841.2017.1308173.
- Elizabeth Ligon Bjork and Robert A. Bjork. Making things hard on yourself, but in a good way: Creating desirable difficulties to enhance learning. In Morton Ann Gernsbacher, Richard W. Pew, Leaetta M. Hough, and James R. Pomerantz, editors, *Psychology and*

- the Real World: Essays Illustrating Fundamental Contributions to Society*, pages 56–64. Worth Publishers, New York, 2011. ISBN 9781429230438.
- Paul Black and Dylan Wiliam. Assessment and classroom learning. *Assessment in Education: Principles, Policy & Practice*, 5(1):7–74, 1998. doi: 10.1080/0969595980050102.
- Angxuan Chen, Mengtong Xiang, Junyi Zhou, Jinjin Jia, Junyu Shang, Xinyu Li, Dragan Gašević, and Yizhou Fan. Unpacking help-seeking process through multimodal learning analytics: A comparative study of ChatGPT vs human expert. *Computers & Education*, 226:105198, 2025. doi: 10.1016/j.compedu.2024.105198.
- Debby R. E. Cotton, Peter A. Cotton, and J. Reuben Shipway. Chatting and cheating: Ensuring academic integrity in the era of ChatGPT. *Innovations in Education and Teaching International*, 61(2):228–239, 2024. doi: 10.1080/14703297.2023.2190148.
- John Dunlosky and Janet Metcalfe. *Metacognition*. SAGE Publications, Thousand Oaks, CA, 2009. ISBN 9781412939720.
- John Dunlosky, Katherine A. Rawson, Elizabeth J. Marsh, Mitchell J. Nathan, and Daniel T. Willingham. Improving students’ learning with effective learning techniques: Promising directions from cognitive and educational psychology. *Psychological Science in the Public Interest*, 14(1):4–58, 2013. doi: 10.1177/1529100612453266.
- Elon University and American Association of Colleges and Universities. National survey: 95% of college faculty fear student overreliance on AI and diminished critical thinking among learners who use generative AI tools. Elon University / AAC&U, January 2026. URL [https://www.aacu.org/newsroom/national-survey-95-of-college-faculty-fear-student-overreliance-on-ai-and-diminished-](https://www.aacu.org/newsroom/national-survey-95-of-college-faculty-fear-student-overreliance-on-ai-and-diminished-critical-thinking-among-learners-who-use-generative-ai-tools)Non-scientific survey conducted between October 29 and November 26, 2025, using a list of college and university faculty members developed by AAC&U and Elon University. N=1,057 respondents diverse across academic disciplines, school sizes, job titles, and student populations. Co-authors include Lee Rainie (Elon University’s Imagining the Digital Future Center) and Eddie Watson (AAC&U VP for Digital Innovation). Wayback Machine snapshot to be added during VER pass.
- European Commission, Directorate-General for Education, Youth, Sport and Culture. Ethical guidelines on the use of artificial intelligence and data in teaching and learning for educators. Technical report, Publications Office of the European Union, Luxembourg, 2022. URL <https://op.europa.eu/en/publication-detail/-/publication/d81a0d54-5348-11ed-92ed-01aa75ed71a1/language-en>.
- Yizhou Fan, Linxuan Tang, Huixiao Le, Kaiwen Shen, Sitan Tan, Yueying Zhao, Yingjie Shen, Xinyu Li, and Dragan Gašević. Beware of metacognitive laziness: Effects of generative artificial intelligence on learning motivation, processes, and performance. *British Journal of Educational Technology*, 56(2):489–530, 2025. doi: 10.1111/bjet.13544.
- Scott Freeman, Sarah L. Eddy, Miles McDonough, Michelle K. Smith, Nnadozie Okoroafor, Hannah Jordt, and Mary Pat Wenderoth. Active learning increases student performance in

- science, engineering, and mathematics. *Proceedings of the National Academy of Sciences*, 111(23):8410–8415, 2014. doi: 10.1073/pnas.1319030111.
- Chahna Gonsalves. Addressing student non-compliance in AI use declarations: Implications for academic integrity and assessment in higher education. *Assessment & Evaluation in Higher Education*, 50(4):592–606, 2024. doi: 10.1080/02602938.2024.2415654. URL <https://www.tandfonline.com/doi/full/10.1080/02602938.2024.2415654>. Study finding that 74% of King’s Business School students failed to declare AI usage despite being instructed to do so; cited reasons include fear of negative perceptions and potential lower marks. Wayback Machine snapshot to be added during VER pass.
- John Hattie and Helen Timperley. The power of feedback. *Review of Educational Research*, 77(1):81–112, 2007. doi: 10.3102/003465430298487.
- Higher Education Policy Institute (UK). Student generative AI survey 2025. HEPI / Kortext, February 2025. URL <https://www.hepi.ac.uk/2025/02/26/student-generative-ai-survey-2025/>. UK survey reporting that 92% of students use AI in some form in 2025, up from 66% in 2024. Reported widely including by Fortune (July 2025). Wayback Machine snapshot to be added during VER pass.
- Kashmir Hill. The professors are using ChatGPT, and some students aren’t happy about it. *The New York Times*, May 2025. URL <https://www.nytimes.com/2025/05/14/technology/chatgpt-college-professors.html>. Wayback Machine snapshots: https://web.archive.org/web/*/https://www.nytimes.com/2025/05/14/technology/chatgpt-college-professors.html.
- Irene Hou, Owen Man, Kate Hamilton, Srishty Muthusekaran, Jeffin Johnykutty, Leili Zadeh, and Stephen MacNeil. “All roads lead to ChatGPT”: How generative AI is eroding social interactions and student learning communities. In *Proceedings of the 30th ACM Conference on Innovation and Technology in Computer Science Education V. 1 (ITiCSE 2025)*, pages 1–8, Nijmegen, Netherlands, June 2025. Association for Computing Machinery. ISBN 979-8-4007-1567-9. doi: 10.1145/3724363.3729024. Preprint available as arXiv:2504.09779.
- James M. Hyman. AI and the next layer of human work: A position paper for students, faculty, and administrators. In preparation. Companion paper to *Learning with AI*; will be available at the book’s GitHub repository., 2026a.
- James M. Hyman. Teaching AI responsibly to principled skeptics: A framework for students, faculty, and administrators. In preparation. Companion paper to *Learning with AI*; will be available at the book’s GitHub repository., 2026b.
- Inside Higher Ed and Generation Lab. Survey: College students’ views on AI. Inside Higher Ed Student Voice survey, August 2025. URL <https://www.insidehighered.com/news/students/academics/2025/08/29/survey-college-students-views-ai>. Survey of 1,047 students from 166 two- and four-year institutions, public and private nonprofit, conducted in July 2025 in partnership with Generation Lab. Margin of error plus or minus three percentage points. Wayback Machine snapshot to be added during VER pass.

- Inside Higher Ed and Hanover Research. Presidents puzzled on rebuilding public trust in higher ed: 2026 survey of college and university presidents. Inside Higher Ed, April 2026. URL <https://www.insidehighered.com/news/governance/executive-leadership/2026/04/09/presidents-puzzled-rebuilding-public-trust-higher>. N=430 presidents from public and mostly private nonprofit institutions, two-year and four-year. Margin of error 4 percentage points. Wayback Machine snapshot to be added during VER pass.
- Daniel Kahneman. *Thinking, Fast and Slow*. Farrar, Straus and Giroux, New York, 2011. ISBN 9780374275631.
- Enkelejda Kasneci, Kathrin Seßler, Stefan Küchemann, Maria Bannert, Daryna Dementieva, Frank Fischer, Urs Gasser, Georg Groh, Stephan Günemann, Eyke Hüllermeier, Stephan Krusche, Gitta Kutyniok, Tilman Michaeli, Claudia Nerdel, Jürgen Pfeffer, Oleksandra Poquet, Michael Sailer, Albrecht Schmidt, Tina Seidel, Matthias Stadler, Jochen Weller, Jochen Kuhn, and Gjergji Kasneci. ChatGPT for good? On opportunities and challenges of large language models for education. *Learning and Individual Differences*, 103:102274, 2023. doi: 10.1016/j.lindif.2023.102274.
- Greg Kestin, Kelly Miller, Anna Klaes, Timothy Milbourne, and Gregorio Ponti. AI tutoring outperforms in-class active learning: An RCT introducing a novel research-based design in an authentic educational setting. *Scientific Reports*, 15(1):17458, 2025. doi: 10.1038/s41598-025-97652-6.
- Tyler Kingkade. To avoid accusations of AI cheating, college students are turning to AI. NBC News, January 2026. URL <https://www.nbcnews.com/tech/internet/college-students-ai-cheating-detectors-humanizers-rcna253878>. Wayback Machine snapshots: https://web.archive.org/web/*/https://www.nbcnews.com/tech/internet/college-students-ai-cheating-detectors-humanizers-rcna253878.
- Miles Klee. She was falsely accused of cheating with AI — and she won’t be the last. *Rolling Stone*, June 2023. URL <https://www.rollingstone.com/culture/culture-features/student-accused-ai-cheating-turnitin-1234747351/>. Wayback Machine snapshots: https://web.archive.org/web/*/https://www.rollingstone.com/culture/culture-features/student-accused-ai-cheating-turnitin-1234747351/.
- Tomaž Kosar, Dragana Ostojić, Yu David Liu, and Marjan Mernik. Computer science education in ChatGPT era: Experiences from an experiment in a programming course for novice programmers. *Mathematics*, 12(5):629, 2024. doi: 10.3390/math12050629.
- Nataliya Kosmyna, Eugene Hauptmann, Ye Tong Yuan, Jessica Situ, Xian-Hao Liao, Ashly Vivian Beresnitzky, Iris Braunstein, and Pattie Maes. Your brain on ChatGPT: Accumulation of cognitive debt when using an AI assistant for essay writing task, 2025. URL <https://arxiv.org/abs/2506.08872>.
- Arthur Levine and Scott Van Pelt. *From Upheaval to Action: What Works in Changing Higher Ed*. Johns Hopkins University Press, Baltimore, March 2026. ISBN 978-1421454030.

- Weixin Liang, Mert Yuksekgonul, Yining Mao, Eric Wu, and James Zou. GPT detectors are biased against non-native English writers. *Patterns*, 4(7):100779, 2023. doi: 10.1016/j.patter.2023.100779.
- Xiaohong Liu, Baoxin Guo, Wei He, and Xiaoyong Hu. Effects of generative artificial intelligence on K-12 and higher education students' learning outcomes: A meta-analysis. *Journal of Educational Computing Research*, 2025. doi: 10.1177/07356331251329185.
- Nan Ma and Zhiyong Zhong. A meta-analysis of the impact of generative artificial intelligence on learning outcomes. *Journal of Computer Assisted Learning*, 41(5):e70117, 2025. doi: 10.1111/jcal.70117.
- Kim Manturuk. The best defense against AI cheating (opinion). Inside Higher Ed, April 2026. URL <https://www.insidehighered.com/opinion/career-advice/teaching/2026/04/13/best-defense-against-ai-cheating-opinion>. IHE opinion essay describing faculty postures of surveillance vs. supplication in response to AI cheating. Wayback Machine snapshot to be added during VER pass. Author identity to be confirmed during VER (not visible in search snippet).
- Eric Mazur. *Peer Instruction: A User's Manual*. Prentice Hall Series in Educational Innovation. Prentice Hall, Upper Saddle River, NJ, 1997. ISBN 9780135654415.
- Donald L. McCabe and Linda Klebe Treviño. Academic dishonesty: Honor codes and other contextual influences. *The Journal of Higher Education*, 64(5):522–538, 1993. doi: 10.1080/00221546.1993.11778446.
- Beth McMurtrie. Professors ask: Are we just grading robots? *The Chronicle of Higher Education*, June 2024. URL <https://www.chronicle.com/article/professors-ask-are-we-just-grading-robots>. Print version appeared in the July 5, 2024 issue. Wayback Machine snapshots: https://web.archive.org/web/*/https://www.chronicle.com/article/professors-ask-are-we-just-grading-robots.
- Fengchun Miao and Mutlu Cukurova. AI competency framework for teachers. Technical report, UNESCO, Paris, 2024. URL <https://www.unesco.org/en/articles/ai-competency-framework-teachers>.
- Fengchun Miao and Wayne Holmes. Guidance for generative AI in education and research. Technical report, UNESCO, Paris, 2023. URL <https://www.unesco.org/en/articles/guidance-generative-ai-education-and-research>.
- Fengchun Miao, Kelly Shiohira, and Natalie Lao. AI competency framework for students. Technical report, UNESCO, Paris, 2024. URL <https://www.unesco.org/en/articles/ai-competency-framework-students>.
- Modern Language Association and Conference on College Composition and Communication. MLA-CCCC joint task force on writing and AI working paper: Overview of the issues, statement of principles, and recommendations. Technical report, MLA-CCCC Joint Task Force on Writing and AI, 2023. URL <https://hcommons.org/app/uploads/sites/1003160/2023/07/MLA-CCCC-Joint-Task-Force-on-Writing-and-AI-Working-Paper-1.pdf>.

- Ethan Mollick. *Co-Intelligence: Living and Working with AI*. Portfolio/Penguin, New York, 2024. ISBN 9780593716717.
- Ethan R. Mollick and Lilach Mollick. Assigning AI: Seven approaches for students, with prompts. Working paper, The Wharton School, University of Pennsylvania, 2023. URL <https://arxiv.org/abs/2306.10052>. arXiv:2306.10052.
- National Academies of Sciences, Engineering, and Medicine. *How People Learn II: Learners, Contexts, and Cultures*. National Academies Press, Washington, DC, 2018. ISBN 9780309459648. doi: 10.17226/24783.
- OECD. *OECD Digital Education Outlook 2026: Exploring Effective Uses of Generative AI in Education*. OECD Publishing, Paris, 2026. ISBN 978-92-64-74128-7. doi: 10.1787/062a7394-en.
- DeRionne Pollard. On Public Trust in Community Colleges. Quoted in Colleen Flaherty, “Presidents Puzzled on Rebuilding Public Trust in Higher Ed,” *Inside Higher Ed*, April 2026. URL <https://www.insidehighered.com/news/governance/executive-leadership/2026/04/09/presidents-puzzled-rebuilding-public-trust-higher>. Quote attributed to DeRionne Pollard, president and CEO of the American Association of Community Colleges: “Public trust in higher education has not eroded at community colleges. For eight decades we have been the most reliable source of affordable, accredited, career-connected higher education in most communities. If anything, people are turning to community colleges more frequently because they know that we are affordable, connected to career pathways and deeply engaged with the business needs of local communities.” Published 9 April 2026.
- Geraint Rees and James Wilsdon. Could agentic AI topple grant-funding systems? *Nature*, 652:1119–1121, April 2026. doi: 10.1038/d41586-026-01297-y. URL <https://www.nature.com/articles/d41586-026-01297-y>. Lead author affiliation: Geraint Rees, UCL Vice-Provost of Research, Innovation & Global Engagement. Wayback Machine snapshot to be added during VER pass.
- Henry L. Roediger and Jeffrey D. Karpicke. Test-enhanced learning: Taking memory tests improves long-term retention. *Psychological Science*, 17(3):249–255, 2006. doi: 10.1111/j.1467-9280.2006.01693.x.
- Everett M. Rogers. *Diffusion of Innovations*. Free Press, New York, 5 edition, 2003. ISBN 9780743222099.
- Doug Rohrer and Kelli M. Taylor. The shuffling of mathematics problems improves learning. *Instructional Science*, 35(6):481–498, 2007. doi: 10.1007/s11251-007-9015-8.
- Anita Samuel. Faculty + AI: A retrospective of 2025. Tech-Savvy Teaching (Substack), December 2025. URL <https://anitasamuel.substack.com/p/faculty-ai-a-retrospective-of-2025>. Substack essay describing four faculty postures toward AI: Hardliners, Redesigners, Resigned, and Disengaged. Wayback Machine snapshot to be added during VER pass.

- Elise Silva. University students feel ‘anxious, confused and distrustful’ about AI in the classroom and among their peers. The Conversation, July 2025. URL <https://theconversation.com/university-students-feel-anxious-confused-and-distrustful-about-ai-in-the-classroom-a> Reports findings from focus groups with 95 University of Pittsburgh students (Spring 2025). Acknowledged co-investigators: Annette Vee, Patrick Manning, Jessica FitzPatrick, Jessica Ghilani, Catherine Kula, Patty Wharton-Michael, Jialei Jiang, Sean DiLeonardi, Birney Young, Mark DiMauro, Jeff Aziz, Gayle Rogers. Wayback Machine snapshots: https://web.archive.org/web/*/https://theconversation.com/university-students-feel-anxious-confused-and-distrustful-about-ai-in-the-classroom-a
- Nicholas C. Soderstrom and Robert A. Bjork. Learning versus performance: An integrative review. *Perspectives on Psychological Science*, 10(2):176–199, 2015. doi: 10.1177/1745691615569000.
- Lihui Sun and Liang Zhou. Does generative artificial intelligence improve the academic achievement of college students? A meta-analysis. *Journal of Educational Computing Research*, 62(7):1896–1933, 2024. doi: 10.1177/07356331241277937.
- John Sweller, Paul Ayres, and Slava Kalyuga. *Cognitive Load Theory*, volume 1 of *Explorations in the Learning Sciences, Instructional Systems and Performance Technologies*. Springer, New York, 2011. ISBN 9781441981257. doi: 10.1007/978-1-4419-8126-4.
- U.S. Department of Education, Office of Educational Technology. Artificial intelligence and the future of teaching and learning: Insights and recommendations. Technical report, U.S. Department of Education, Washington, DC, 2023. URL <https://www.ed.gov/sites/ed/files/documents/ai-report/ai-report.pdf>.
- James D. Walsh. Everyone is cheating their way through college. *New York Magazine / Intelligencer*, May 2025. URL <https://nymag.com/intelligencer/article/openai-chatgpt-ai-cheating-education-college-students-school.html>. Wayback Machine snapshots: https://web.archive.org/web/*/https://nymag.com/intelligencer/article/openai-chatgpt-ai-cheating-education-college-students-school.html.
- William H. Walters and Esther Isabelle Wilder. Fabrication and errors in the bibliographic citations generated by ChatGPT. *Scientific Reports*, 13(1):14045, 2023. doi: 10.1038/s41598-023-41032-5.
- Rose E. Wang, Ana T. Ribeiro, Carly D. Robinson, Susanna Loeb, and Dorottya Demszky. Tutor CoPilot: A human-AI approach for scaling real-time expertise. EdWorkingPaper 24-1056, Annenberg Institute at Brown University, 2024.
- WGU Labs. AI as the new digital divide for first-generation students. WGU Labs blog (EdTech Survey), 2024. URL <https://www.wgulabs.org/posts/ai-as-the-new-digital-divide-for-first-generation-students>. WGU Labs EdTech Survey reporting that 48% of continuing-generation students were aware of ChatGPT compared with 34% of first-generation students; 11% of continuing-generation students

had used ChatGPT for coursework compared with 7% of first-generation students. Survey covered nine diverse colleges and universities. Wayback Machine snapshot to be added during VER pass.

Grant Wiggins and Jay McTighe. *Understanding by Design*. Association for Supervision and Curriculum Development (ASCD), Alexandria, VA, 2 edition, 2005. ISBN 9781416600350.

Yan Zhao, Yuhe Yue, Zhonghua Sun, Qiang Jiang, and Gangsheng Li. Does generative artificial intelligence improve students' higher-order thinking? A meta-analysis based on 29 experiments and quasi-experiments. *Journal of Intelligence*, 13(12):160, 2025. doi: 10.3390/jintelligence13120160.




Figure 11: *
The costume was convincing. The question was not.




Figure 12: **Four-phase implementation timeline.** Universities adopt the framework progressively rather than at once. Phase 1 builds confidence in a small set of courses; Phase 2 lets departments localize; Phase 3 creates institutional consistency; Phase 4 maintains and refines. Each phase produces concrete outputs that feed the next.

Impact	high	Unequal access	Privacy breaches	Detector dependence
	medium	Policy fragmentation	Faculty workload	Over-reliance on AI
	low		Underuse of AI literacy	False confidence
		low	medium	high
		Likelihood		

Figure 13: **Risk register heat map.** Each risk from Table 9 positioned by likelihood and impact. Risks in the upper-right zone deserve the strongest controls; risks in the lower-left zone deserve monitoring rather than immediate intervention. Local conditions will shift particular risks; institutions should re-assess annually.




Figure 14: **The eight competency domains.** Four mathematical reasoning domains (D1–D4) are assessed primarily under no-AI conditions; two tool-mediated domains (D6, D7) are assessed under bounded- or open-AI conditions; two cross-cutting domains (D5, D8) span both. A graduate program that uses GMIP weighs and assesses all eight, though the relative emphasis may vary by area of mathematics.

Table 12: Default AI-policy mixes for common mathematics course types. The mixes are starting points; individual courses should adjust to local conditions.

Course type	Typical category mix	Strong assessment artifacts
Introductory procedural (Calculus I, College Algebra)	Mostly NAI and AIT, selective AIC for study support or critique tasks.	Short no-AI quizzes, homework-to-quiz transfer checks, brief explanation of one step, AI-generated error critique, formula or concept checks.
Intermediate proof transition (Discrete Math, Intro to Proof)	Balanced NAI and AIT with selective AIC and AIS.	In-class proof sketches, definition recall, proof audit-and-repair tasks, short oral explanations, annotated correction of flawed AI proofs.
Upper-division proof (Real Analysis, Abstract Algebra, Topology)	Mixed NAI, AIC, and AIS; AIT remains useful for study support.	Board proofs, oral defense of homework, plan-use-verify submissions, proof comparison tasks, hypothesis-stress tests, AI-proof audit tasks.
Computational or numerical (Numerical Analysis, Scientific Computing)	Balanced AIC with meaningful NAI and AIS components.	Code test reports, reproducible notebooks, boundary-case checks, short no-AI conceptual quizzes, model comparison memos, critique of AI-generated code.
Applied modeling (ODE/PDE modeling, mathematical biology, operations research)	AIC and AIS with targeted NAI checks on core concepts.	Assumption-to-model notes, verification of outputs against benchmarks, sensitivity checks, short oral or video defense, process logs, AI-assisted model critique.
Graduate proof-intensive seminar or core course	Substantial AIC and AIS with NAI oral or written competence checks.	Blackboard examinations, oral defense, proof repair, comparative AI/human solution analysis, disclosure blocks, verification notes, mini research projects.